

# You Might Not Want Async

# Quick Questions

- Concurrency with Python
- Threads
- Multi-processing
- Single-thread asynchrony
- asyncio

# Me

- Call me TP
- Follow @uranusjr
- <https://uranusjr.com>



djangogirls

TAIPEI



<http://macdown.uranusjr.com>



*www.* **Bimetek** .com

[NEWS](#)[SPEAKERS](#)[SCHEDULE](#)[SPONSORS](#)

# OSDC

ACADEMIA SINICA TAIPEI, TAIWAN

**04/11 - 04/12**

**OPEN SOURCE DEVELOPERS'  
CONFERENCE 2014**



中央研究院人文社會科學館

INTERNATIONAL CONFERENCE HALL,  
RESEARCH CENTER FOR HUMAN AND  
SOCIAL SCIENCE, ACADEMIA SINICA

» **REGISTER**

**TWITTER**

推文關於 "#osdc.tw"



**NEWS**

Apr 30

OSDC 2014 照片出爐



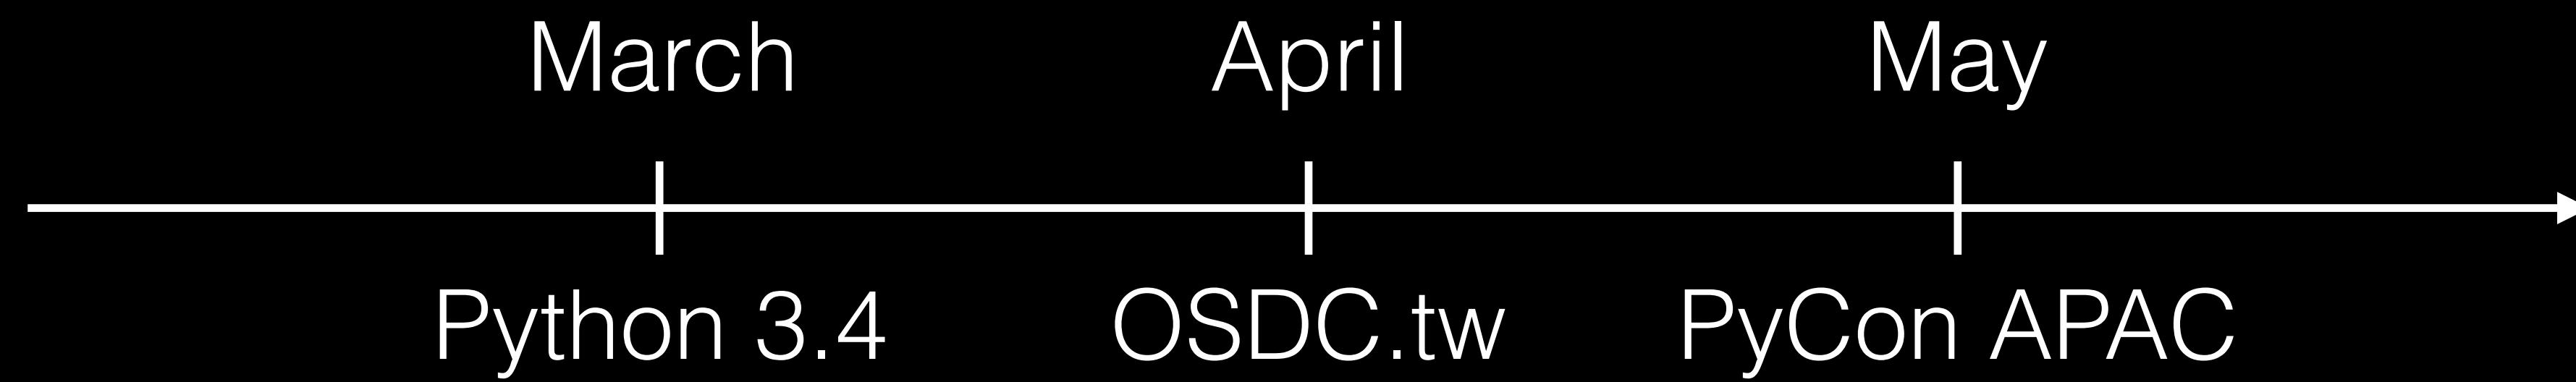
MAY 17-18 2014 in [Taipei](#)

**from everything import future**

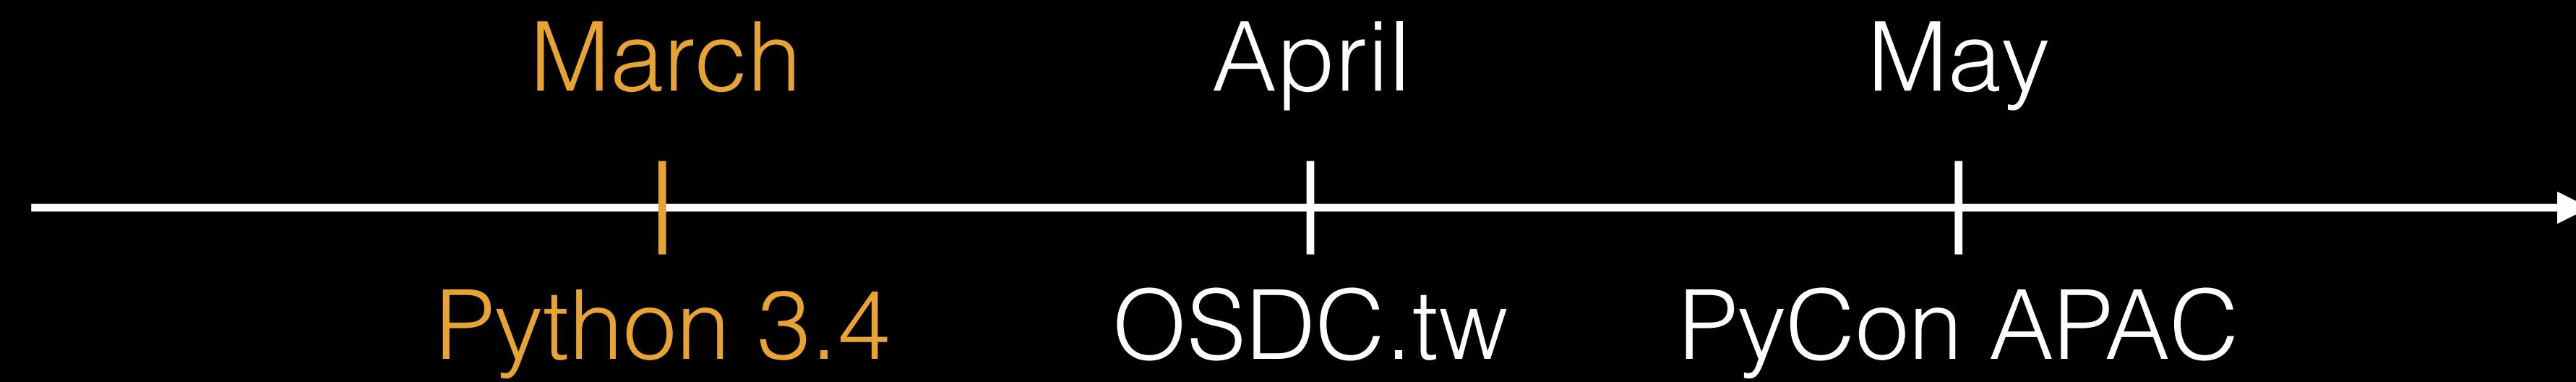
A photograph of four young men standing side-by-side in a dark setting. They are all looking towards the left of the frame with neutral, slightly weary expressions. The man on the far left is wearing a dark zip-up jacket over a grey t-shirt. The second man from the left is wearing a grey t-shirt. The third man is wearing a light-colored striped polo shirt. The man on the far right is wearing a dark zip-up jacket over a grey t-shirt.

meh

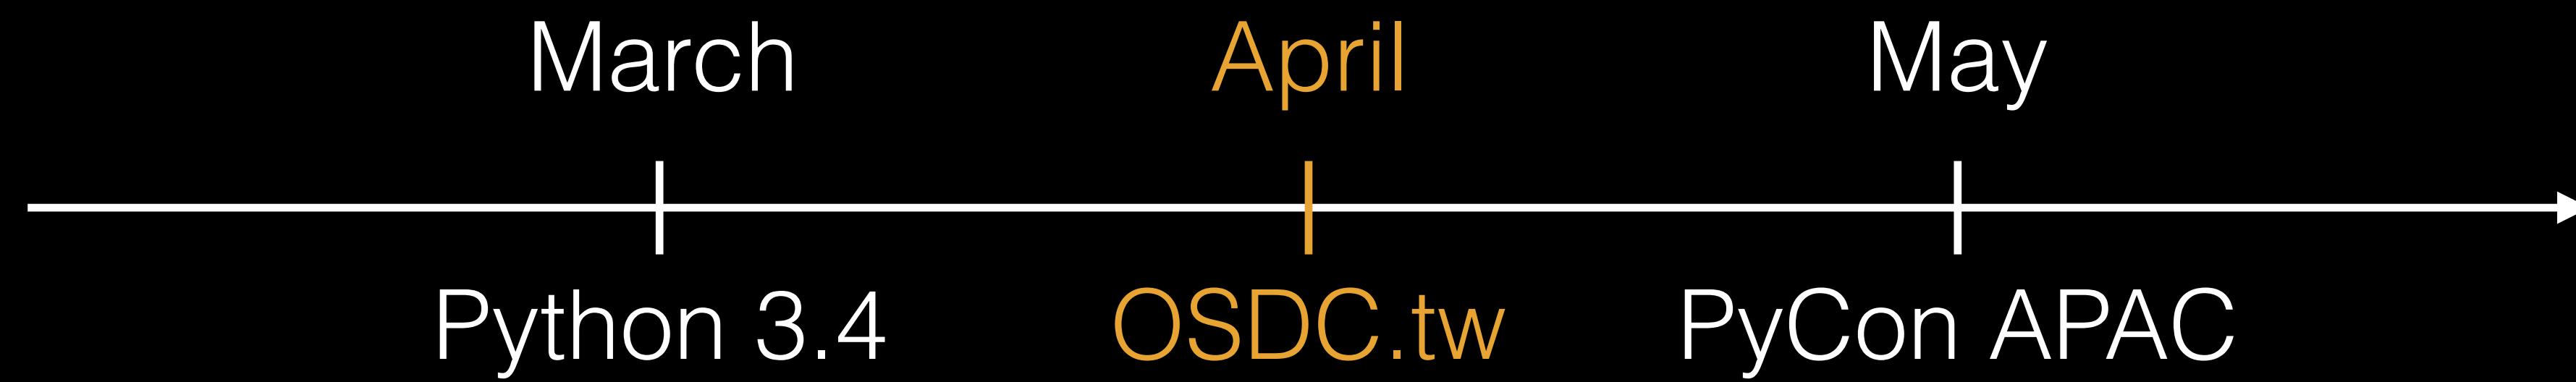
# 2014



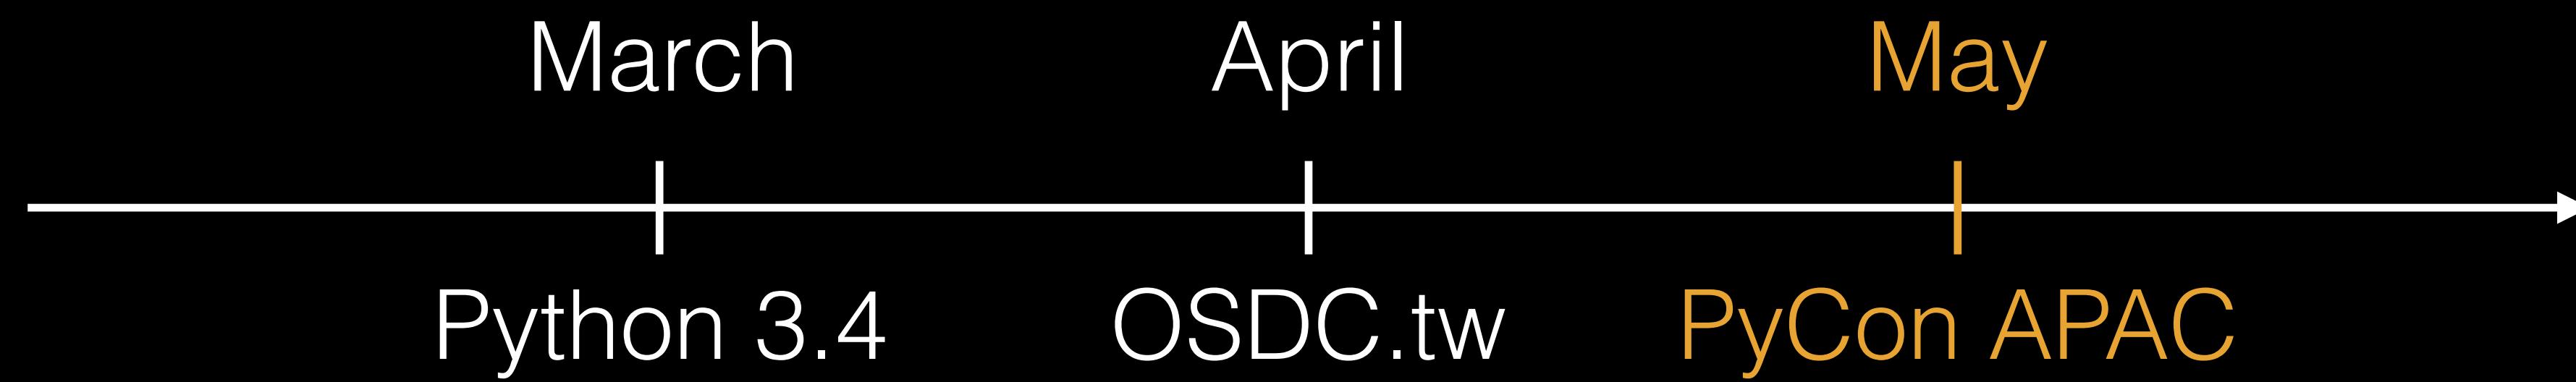
# 2014



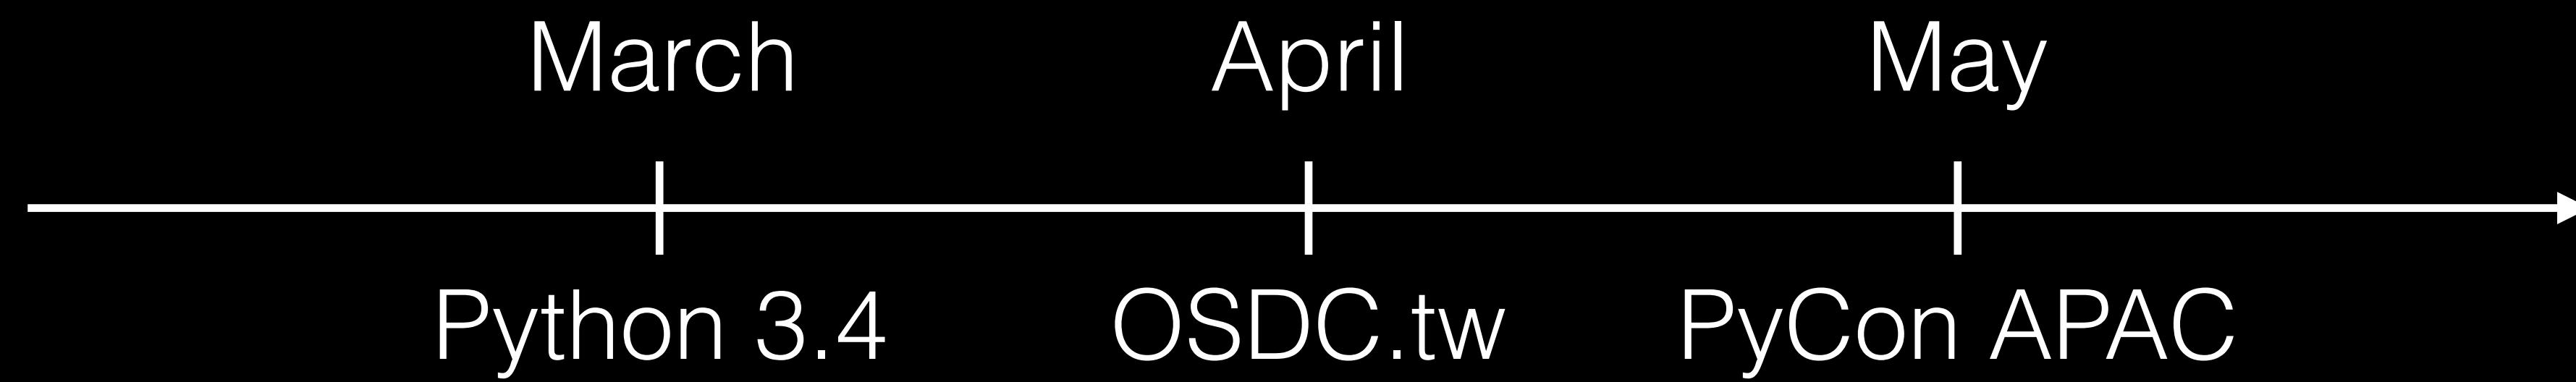
# 2014



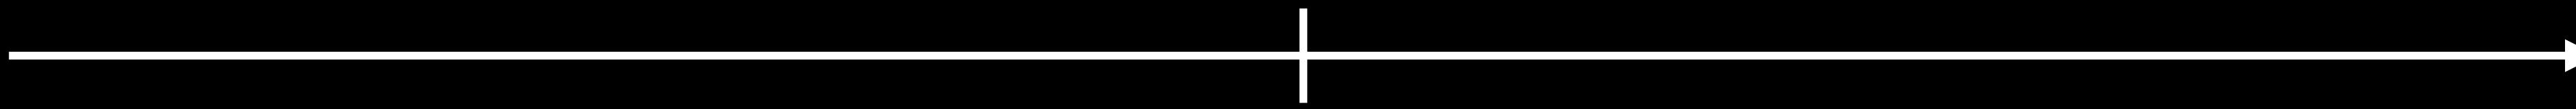
# 2014

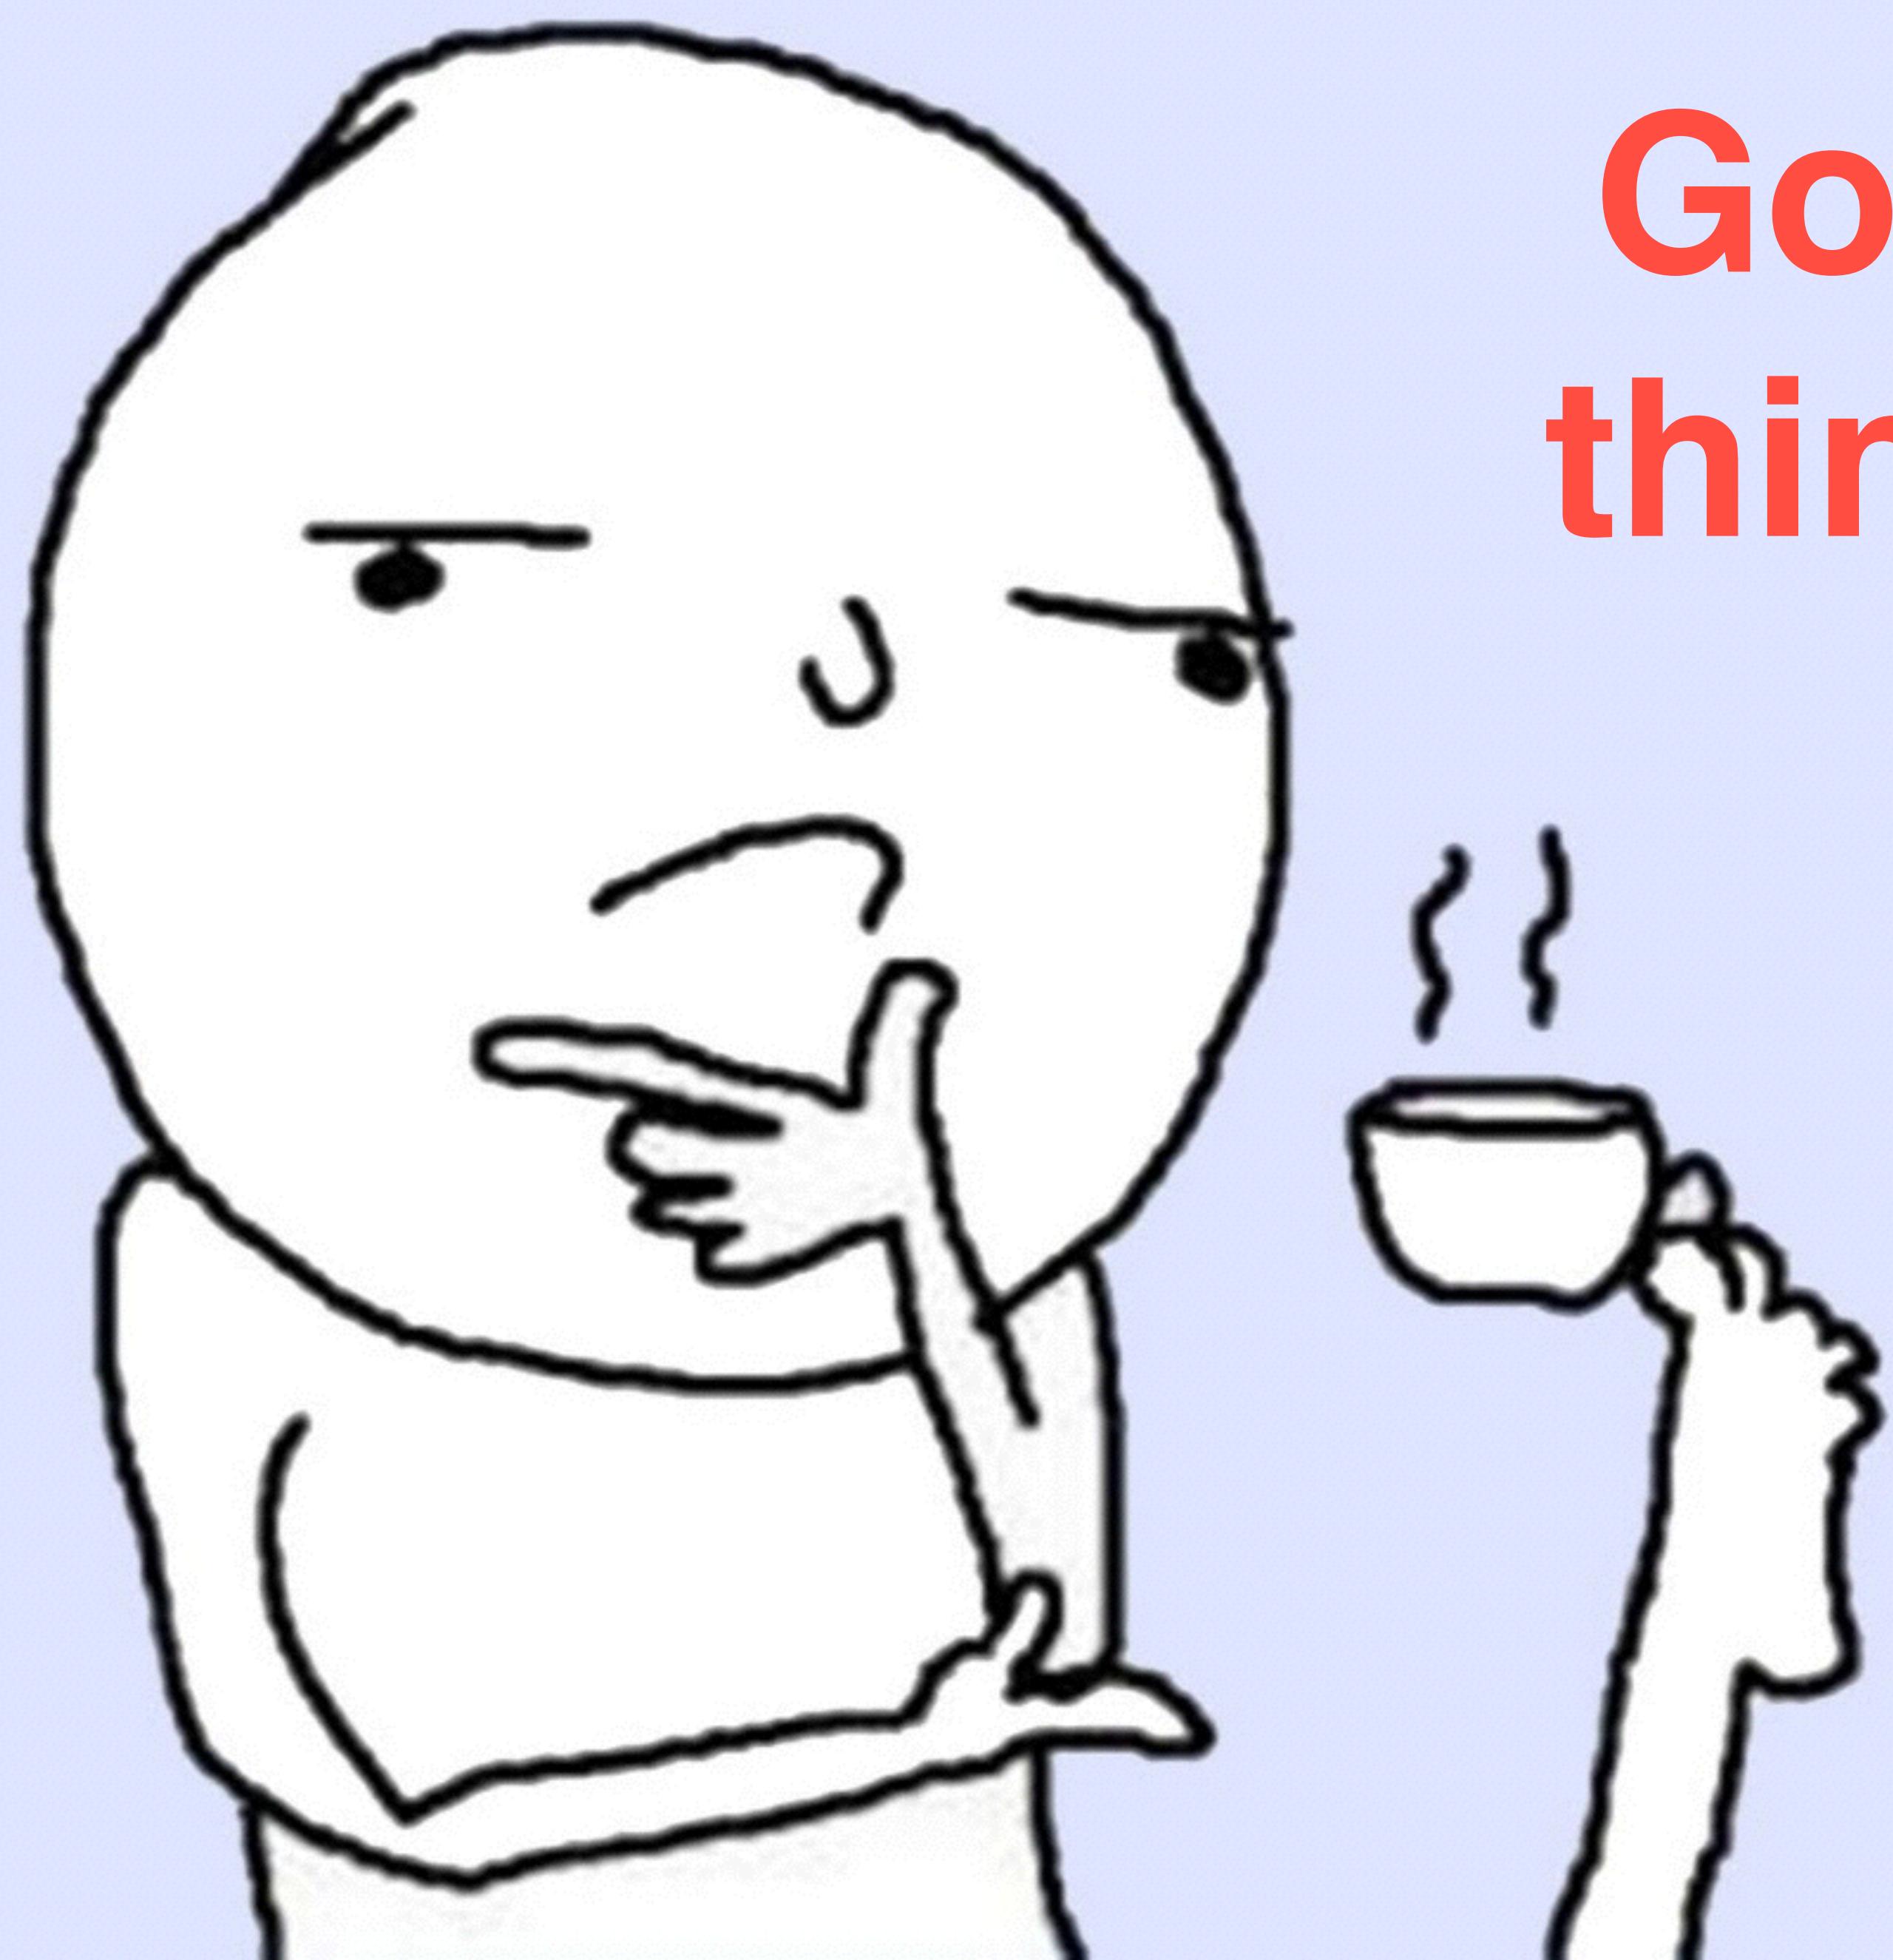


# 2014



2016





Got me  
thinking

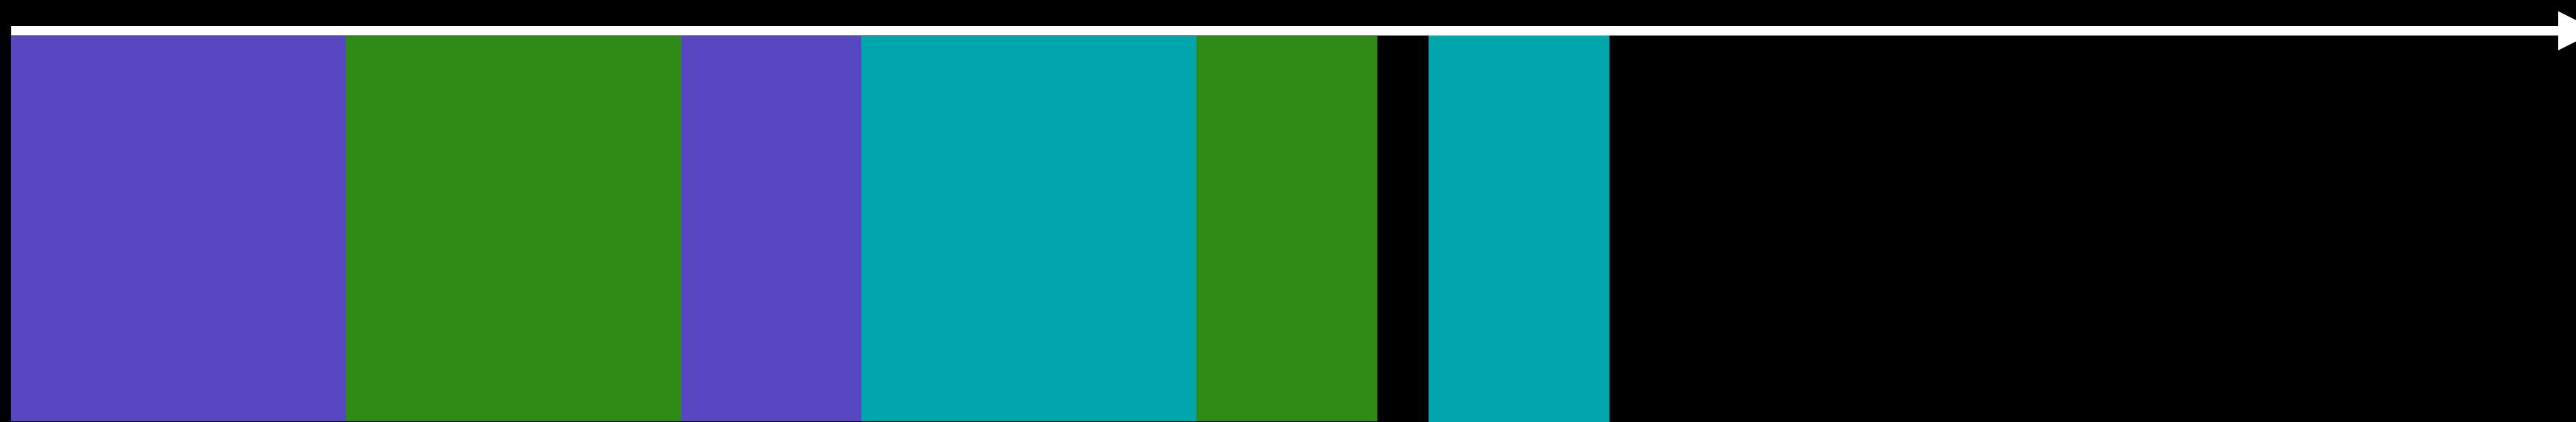
IT'S DANGEROUS TO GO  
ALONE! TAKE THIS.



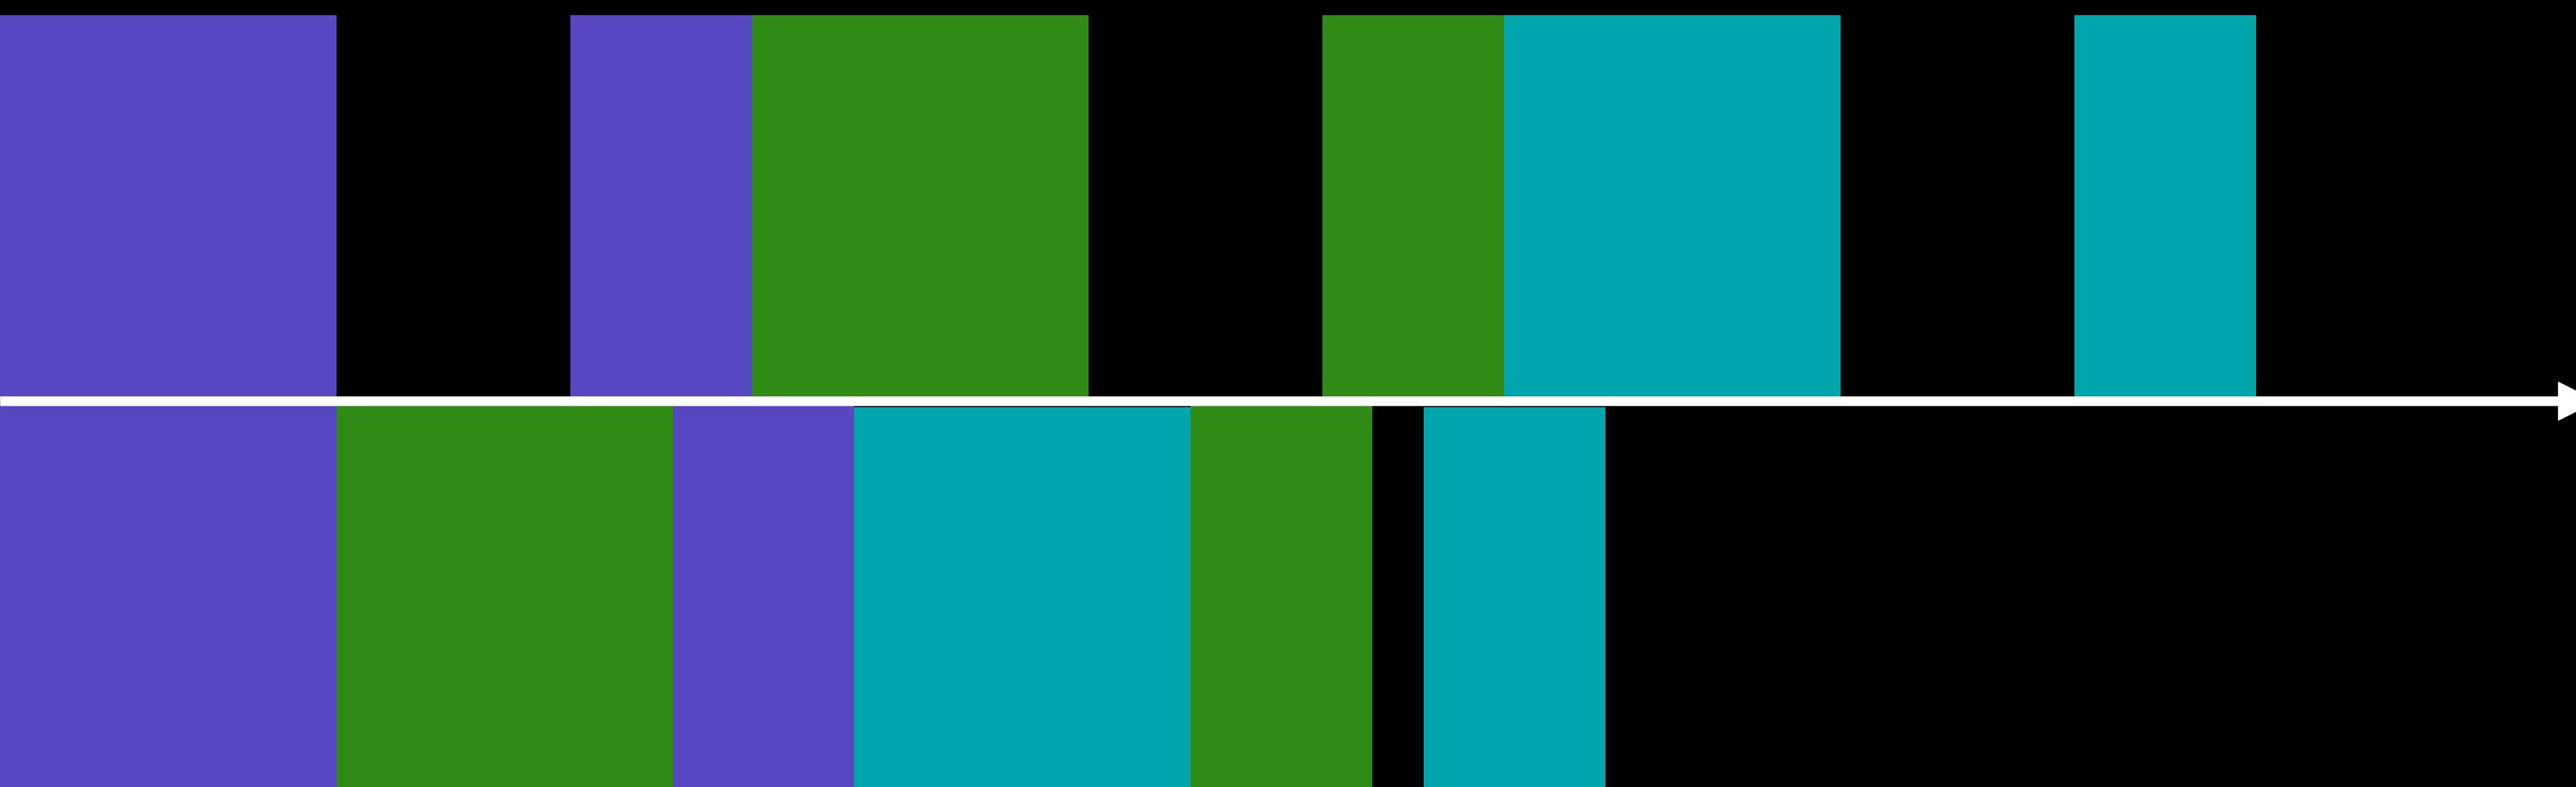
# Synchronous

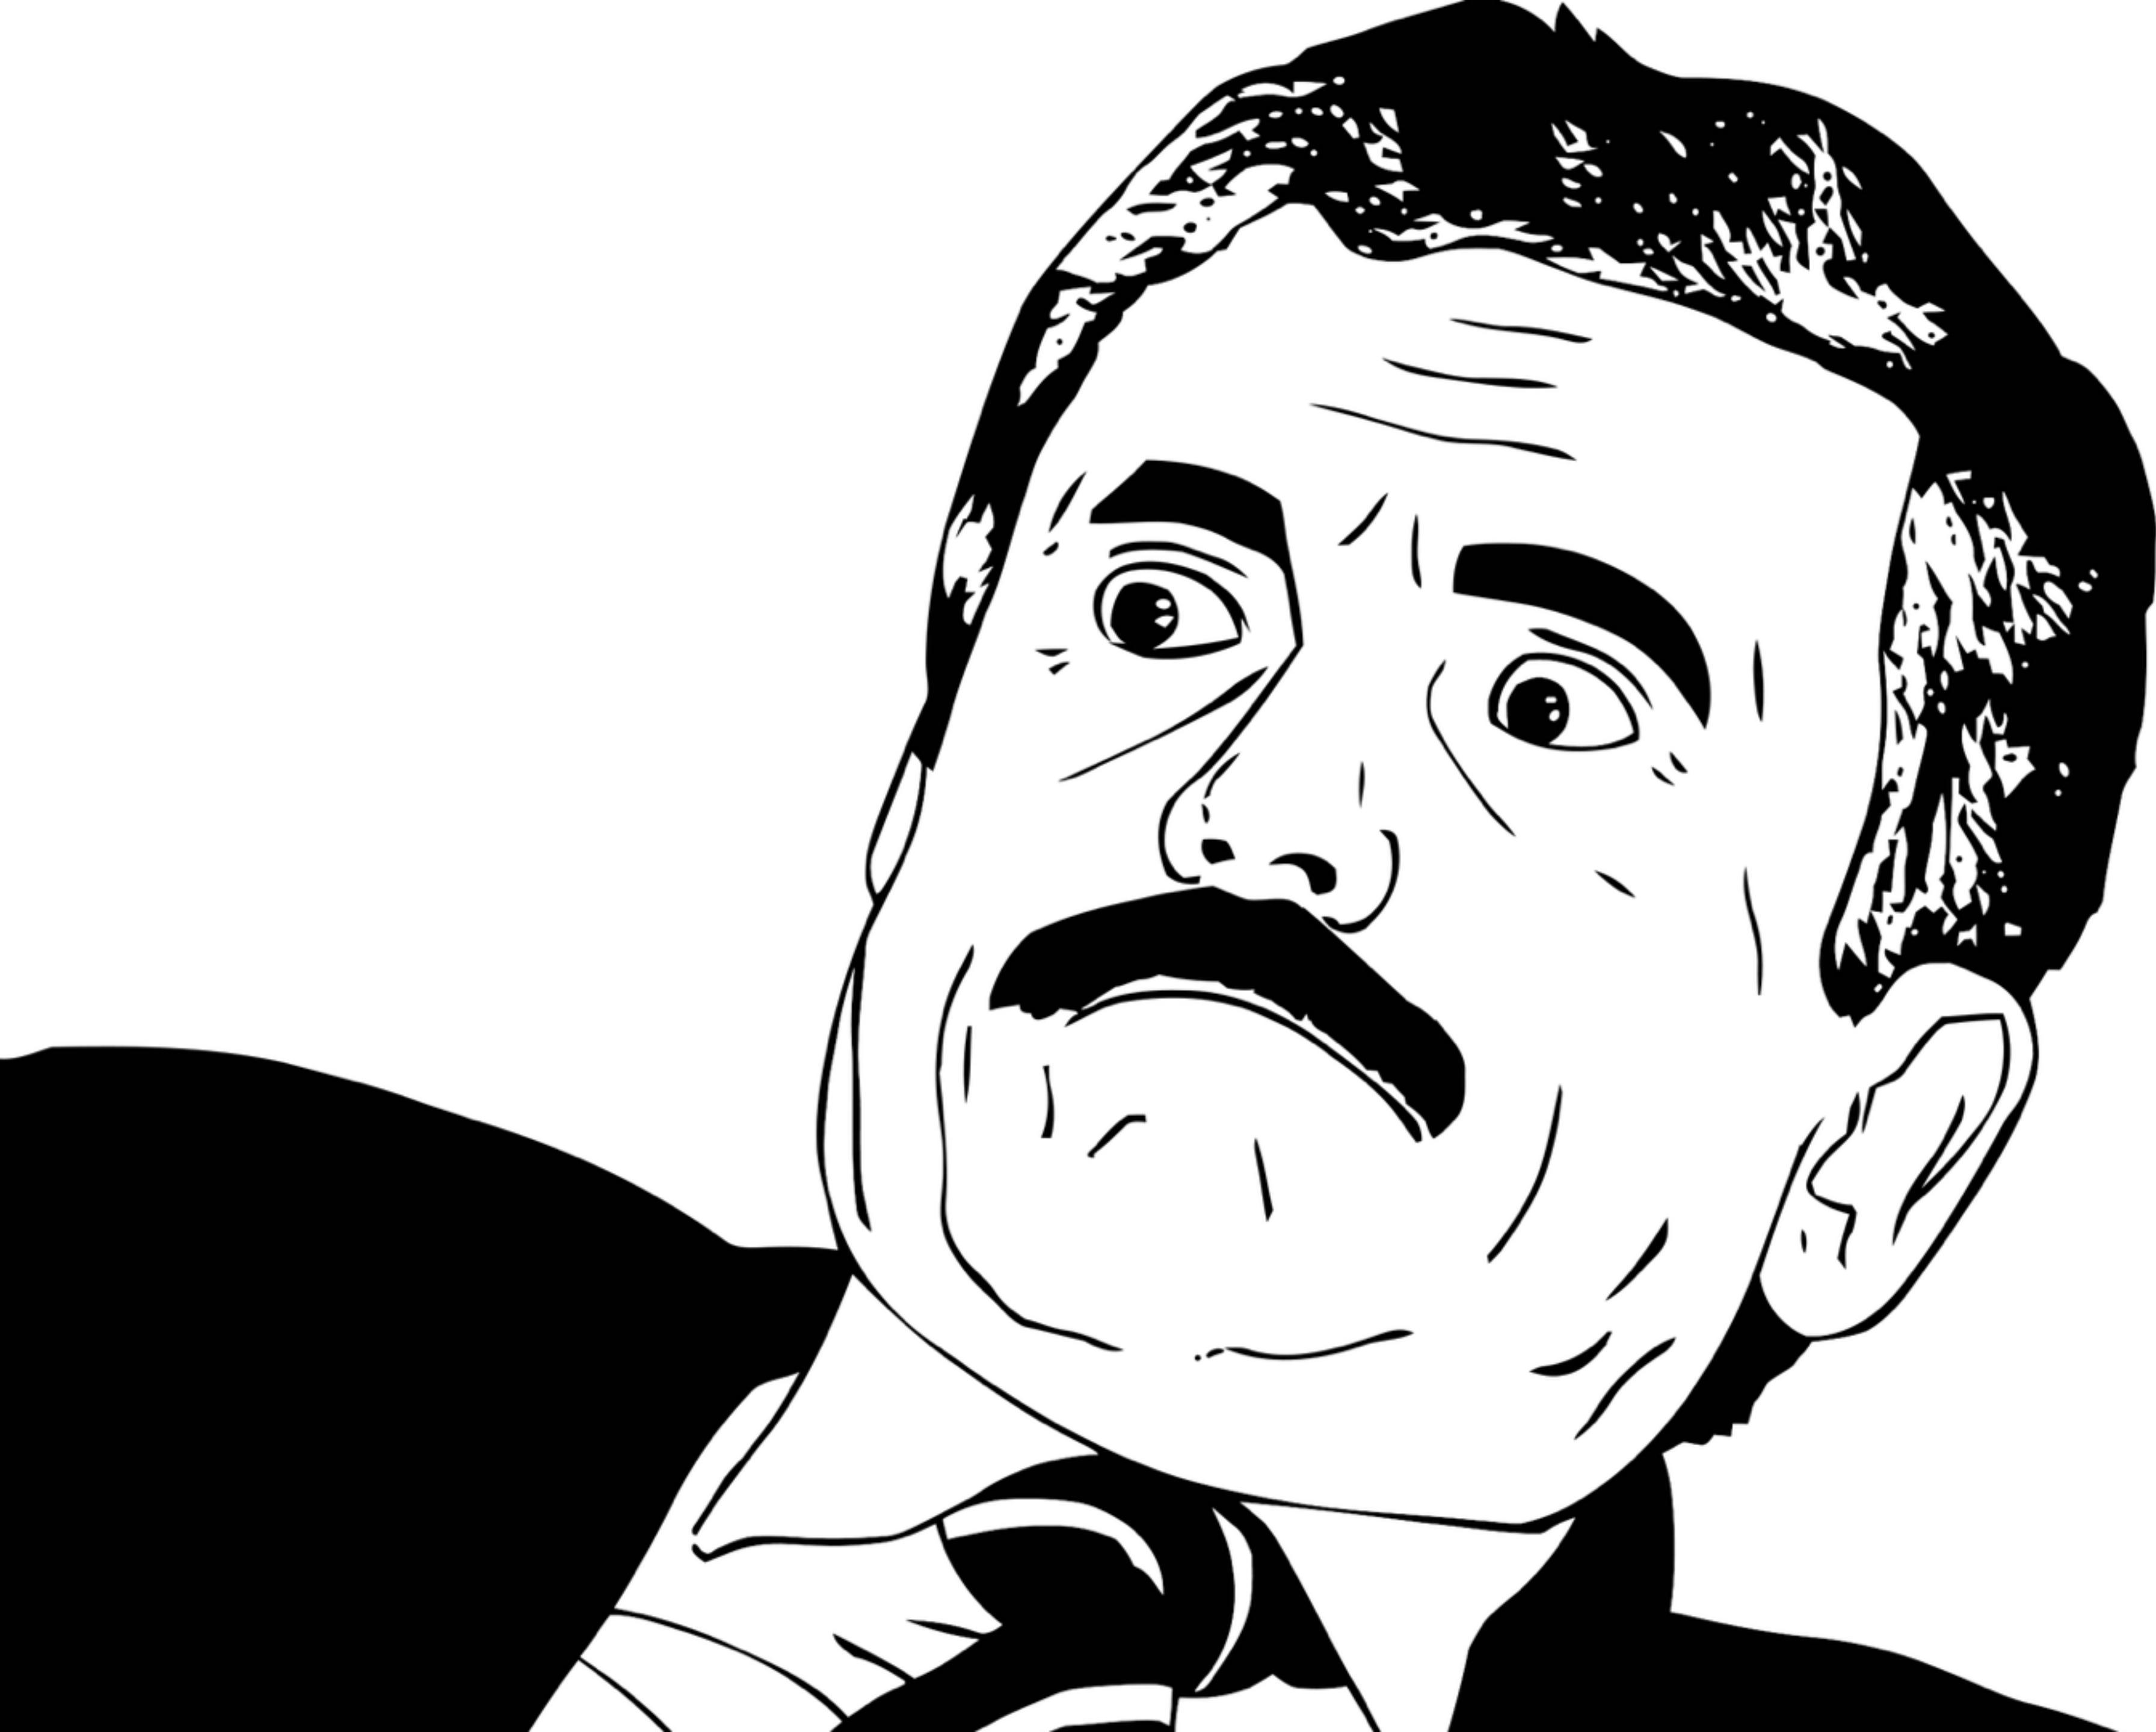


# (Single-Threaded) Async



# Sync vs. Async





Absolute  
magic



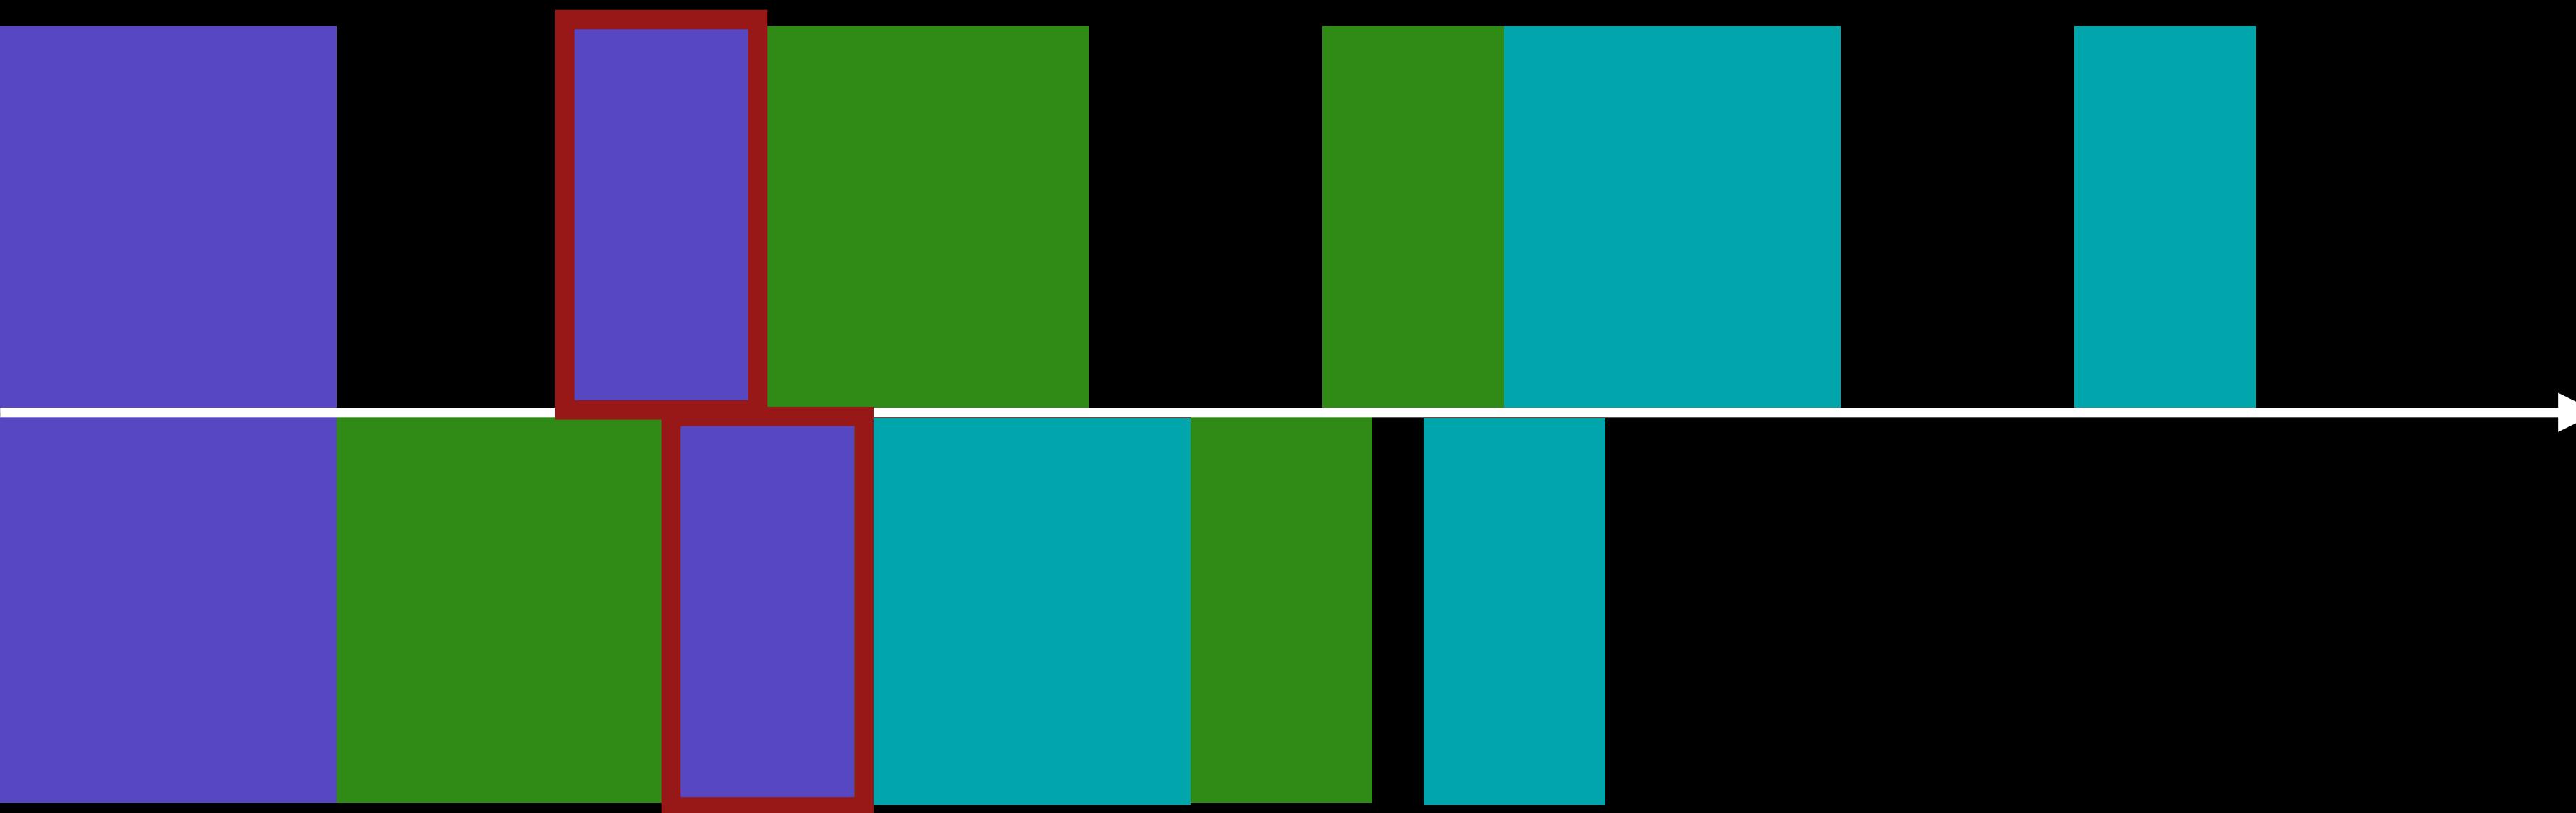
Before



After



# Sync vs. Async



# Not For You

- Infects the whole program
- Async is not parallelism
- Third party support

```
import sqlite3

def read_data(dbname):
    con = sqlite3.connect(dbname)
    cur = con.cursor()
    cur.execute('SELECT * FROM data OFFSET 0 LIMIT 1')
    data = cur.fetchone()
    cur.close()
    con.close()
    return data
```

```
import aioodbc

async def read_data(dbname):
    con = await aioodbc.connect(
        dsn='Driver=SQLite;Database={}'.format(dbname),
    )
    cur = await con.cursor()
    await cur.execute('SELECT * FROM data OFFSET 0 LIMIT 1')
    data = await cur.fetchone()
    await cur.close()
    await con.close()
    return data
```

```
from .db import read_data

def main():
    data = read_data('data.sqlite3')
    print(data)

main()
```

```
import asyncio
from .db import read_data

async def main():
    data = await read_data('db.sqlite3')
    print(data)

loop = asyncio.get_event_loop()
loop.run_until_complete(main())
loop.close()
```



U+1F937 SHRUG (Unicode 9.0)

A close-up photograph of a woman's face in the foreground, looking directly at the camera with a neutral expression. In the background, a building is engulfed in flames, with bright orange and yellow fire visible through the windows and roof. A yellow fire hose lies on the ground in front of the building. The overall atmosphere is one of concern or observation.

**RELAX PEOPLE**

**I'M JUST GETTING STARTED**

```
import asyncio
from .db import read_data

async def main():
    data = read_data('db.sqlite3')
    print(data)

loop = asyncio.get_event_loop()
loop.run_until_complete(main())
loop.close()
```

```
(demo) $ python demo.py  
<coroutine object read_data at 0x10e2d1f10>
```

```
import asyncio
from .db import read_data

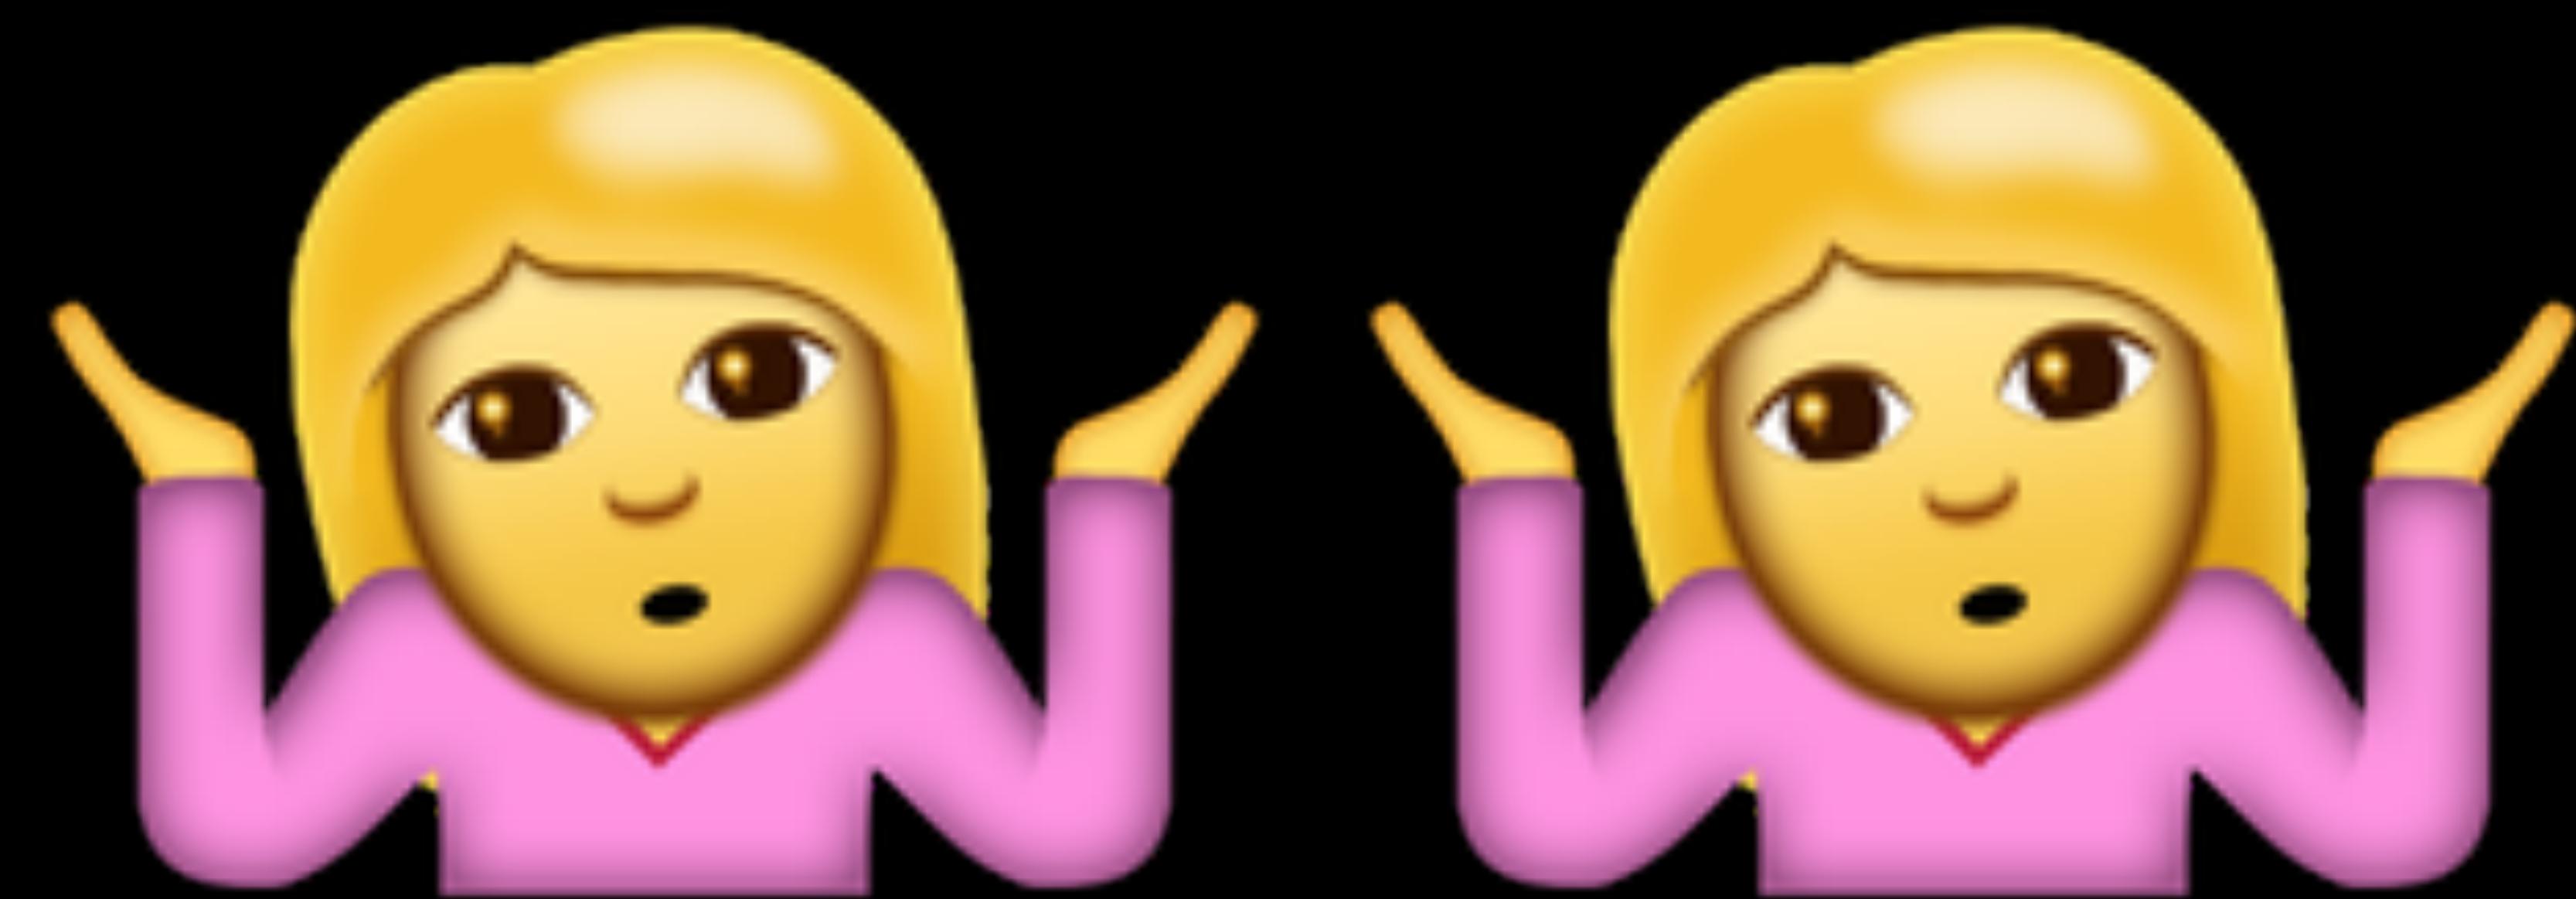
async def main():
    data = read_data('db.sqlite3')
    print(data)

loop = asyncio.get_event_loop()
loop.run_until_complete(main())
loop.close()
```

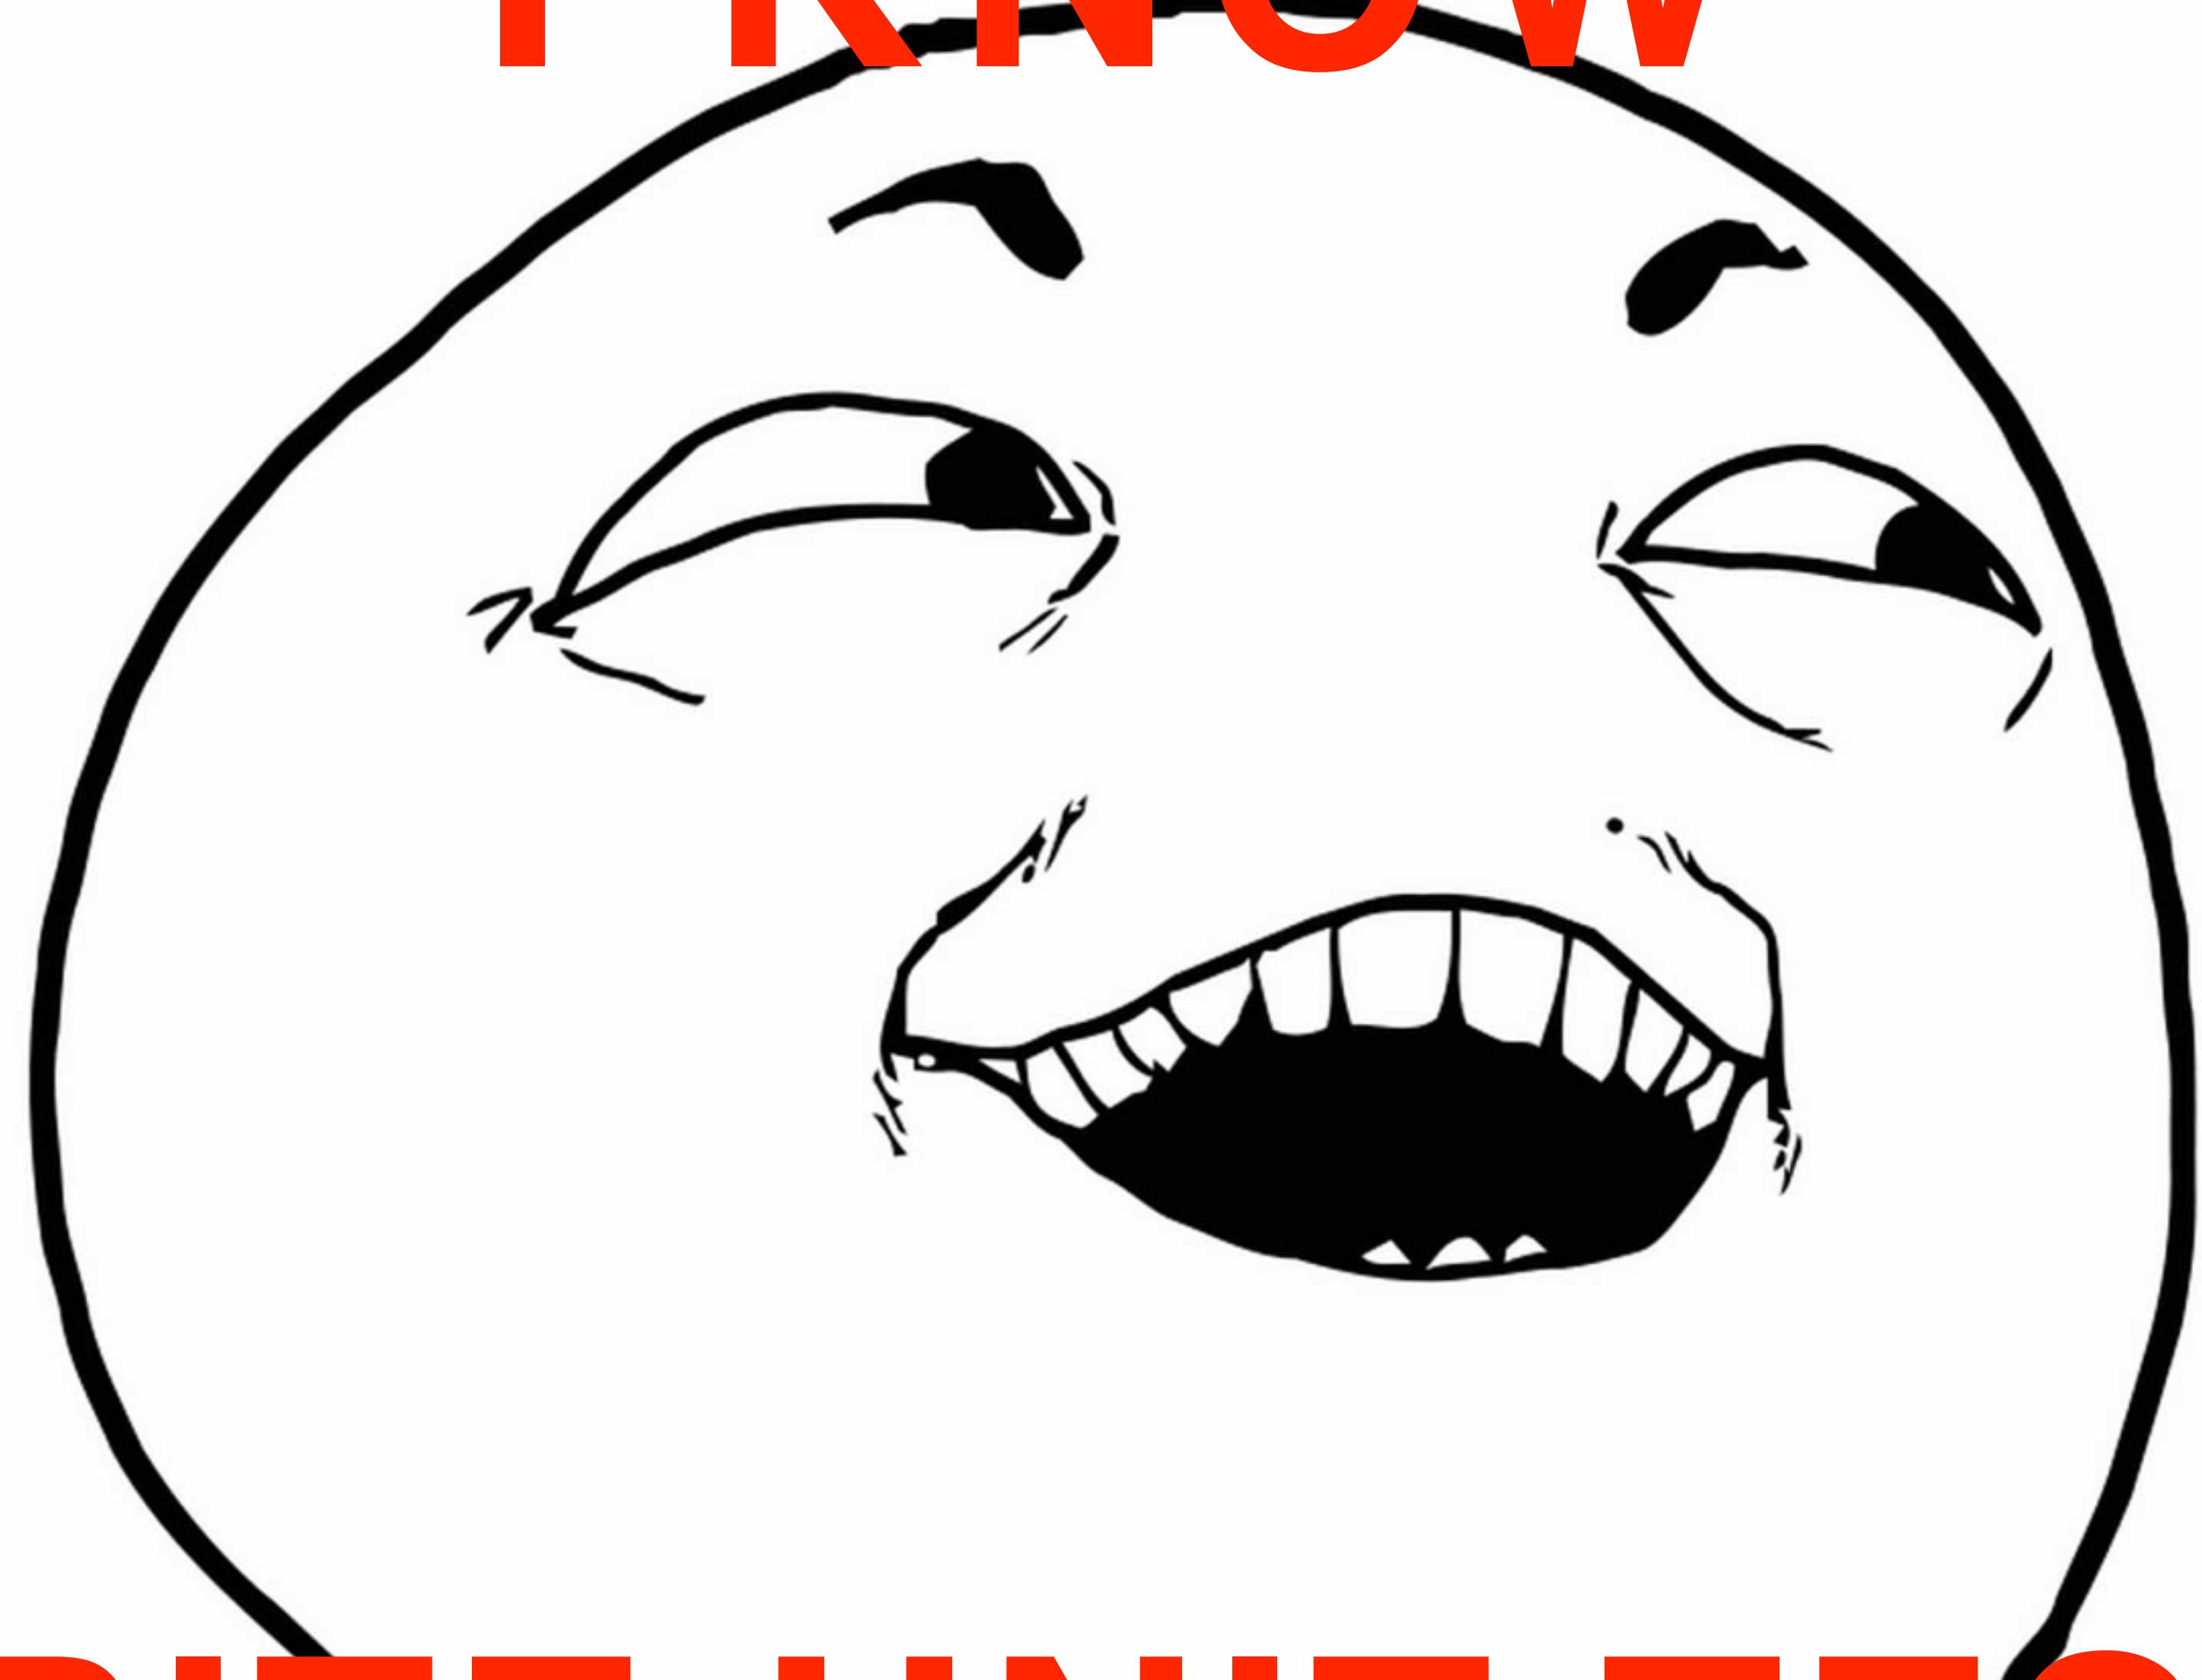
```
import asyncio
from .db import read_data

async def main():
    data = await read_data('db.sqlite3')
    print(data)

loop = asyncio.get_event_loop()
loop.run_until_complete(main())
loop.close()
```



I KNOW



WRITE UNIT TESTS

GOOD LUCK



WITH THAT

# How to test Python 3.4 asyncio code?



What's the best way to write unit tests for code using the Python 3.4 `asyncio` library? Assume I want to test a TCP client (`SocketConnection`):

31



10

```
import asyncio
import unittest

class TestSocketConnection(unittest.TestCase):
    def setUp(self):
        self.mock_server = MockServer("localhost", 1337)
        self.socket_connection = SocketConnection("localhost", 1337)

    @asyncio.coroutine
    def test_sends_handshake_after_connect(self):
        yield from self.socket_connection.connect()
        self.assertTrue(self.mock_server.received_handshake())
```

When running this test case with the default test runner, the test will always succeed as the method executes only up until the first `yield from` instruction, after which it returns before executing any assertions. This causes tests to always succeed.

Is there a prebuilt test runner that is able to handle asynchronous code like this?

# Eventually Correct

A. Jesse Jiryu Davis

@jessejiryudavis

MongoDB



```
import asyncio  
import time
```

```
async def do_something():  
    print('Before', time.monotonic())  
    await asyncio.sleep(2)  
    print('After ', time.monotonic())
```

```
>>> await do_something()  
Before 199208.190632853  
After 199210.192092425
```

```
import unittest

class MyTestCase(unittest.TestCase):
    @async def test_do_something(self):
        await do_something()

if __name__ == '__main__':
    unittest.main()
```

(demo) \$ python tests.py

.

---

Ran 1 test in 0.002s

OK

(demo) \$ python tests.py

.

---

Ran 1 test in 0.002s

OK

# What Happened?

- unittest does not know about asyncio
- Coroutine methods are executed “normally”
- Called, but not executed (`awaited`)

**WHAT IF I TOLD YOU**



**THIS IS GONNA BE TOUGH**

```
import asyncio
import functools

def asynchronous(func):

    @functools.wraps(func)
    def asynchronous_inner(*args, **kwargs):
        loop = asyncio.new_event_loop()
        asyncio.set_event_loop(loop)
        try:
            loop.run_until_complete(func(*args, **kwargs))
        finally:
            loop.close()

    return asynchronous_inner
```

```
import asyncio
import functools

def asynchronous(func):

    @functools.wraps(func)
    def asynchronous_inner(*args, **kwargs):
        loop = asyncio.new_event_loop()
        asyncio.set_event_loop(loop)
        try:
            loop.run_until_complete(func(*args, **kwargs))
        finally:
            loop.close()

    return asynchronous_inner
```

```
import unittest

class MyTestCase(unittest.TestCase):
    @asynchronous
    async def test_do_something(self):
        await do_something()

if __name__ == '__main__':
    unittest.main()
```

(demo) \$ python tests.py

Before 51728.420457105

After 51730.424515145

.

---

Ran 1 test in 2.006s

OK

# Tips

- Beware of warning output
  - Especially if you redirect
- Add coverage report to *testing code*
- Consider pip install asynctest

Or use pytest instead

```
import time

async def test_do_something():
    before = time.monotonic()
    await do_something()
    delta_t = time.monotonic() - before
    assert -0.01 < delta_t < 0.01
```

```
(demo) $ py.test tests.py
```

```
===== test session starts =====
```

```
platform darwin -- Python 3.5.1, pytest-2.9.1, py-1.4.31,
```

```
pluggy-0.3.1
```

```
collected 0 items / 1 errors
```

```
===== ERRORS =====
```

```
_____ ERROR collecting tests.py_____
```

```
>     for i, x in enumerate(self.obj()):
```

```
E       TypeError: 'coroutine' object is not iterable
```

```
python3.5/site-packages/_pytest/python.py:765: TypeError
```

```
===== 1 error in 0.15 seconds =====
```

# pytest-asyncio 0.3.0

*Pytest support for asyncio.*

Downloads ↓

pypi v0.3.0

build passing

coverage 100%

pytest-asyncio is an Apache2 licensed library, written in Python, for testing asyncio code with pytest.

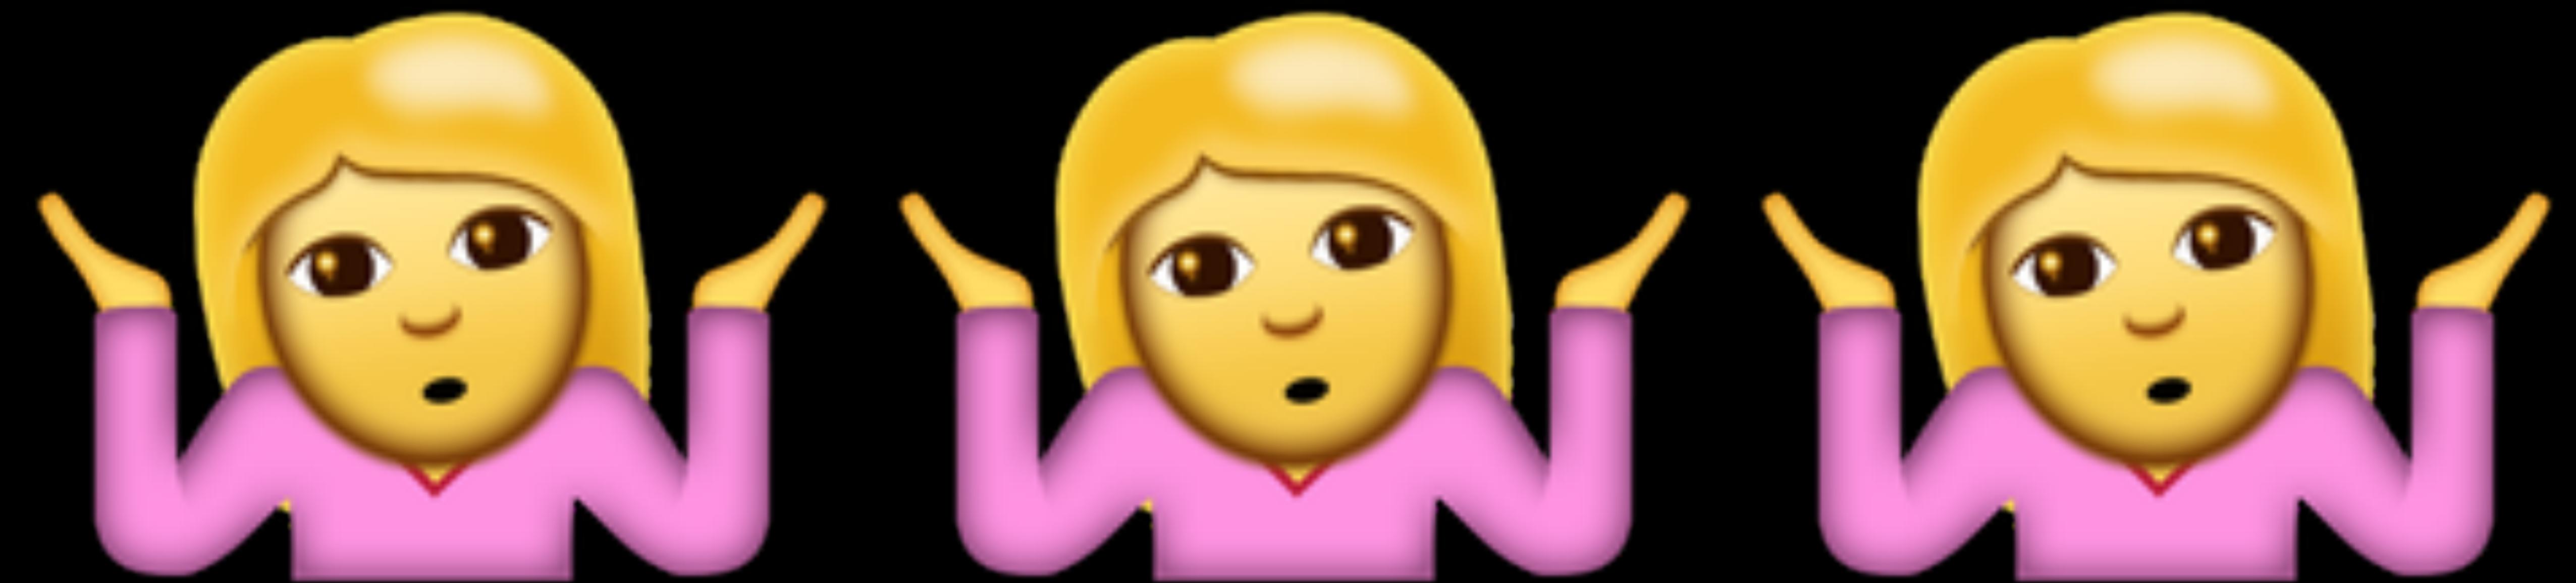
asyncio code is usually written in the form of coroutines, which makes it slightly more difficult to test using normal testing tools. pytest-asyncio provides useful fixtures and markers to make testing easier.

```
@pytest.mark.asyncio
async def test_some_asyncio_code():
    res = await library.do_something()
    assert b'expected result' == res
```

or, if you're using the pre-Python 3.5 syntax:

```
import time
import pytest

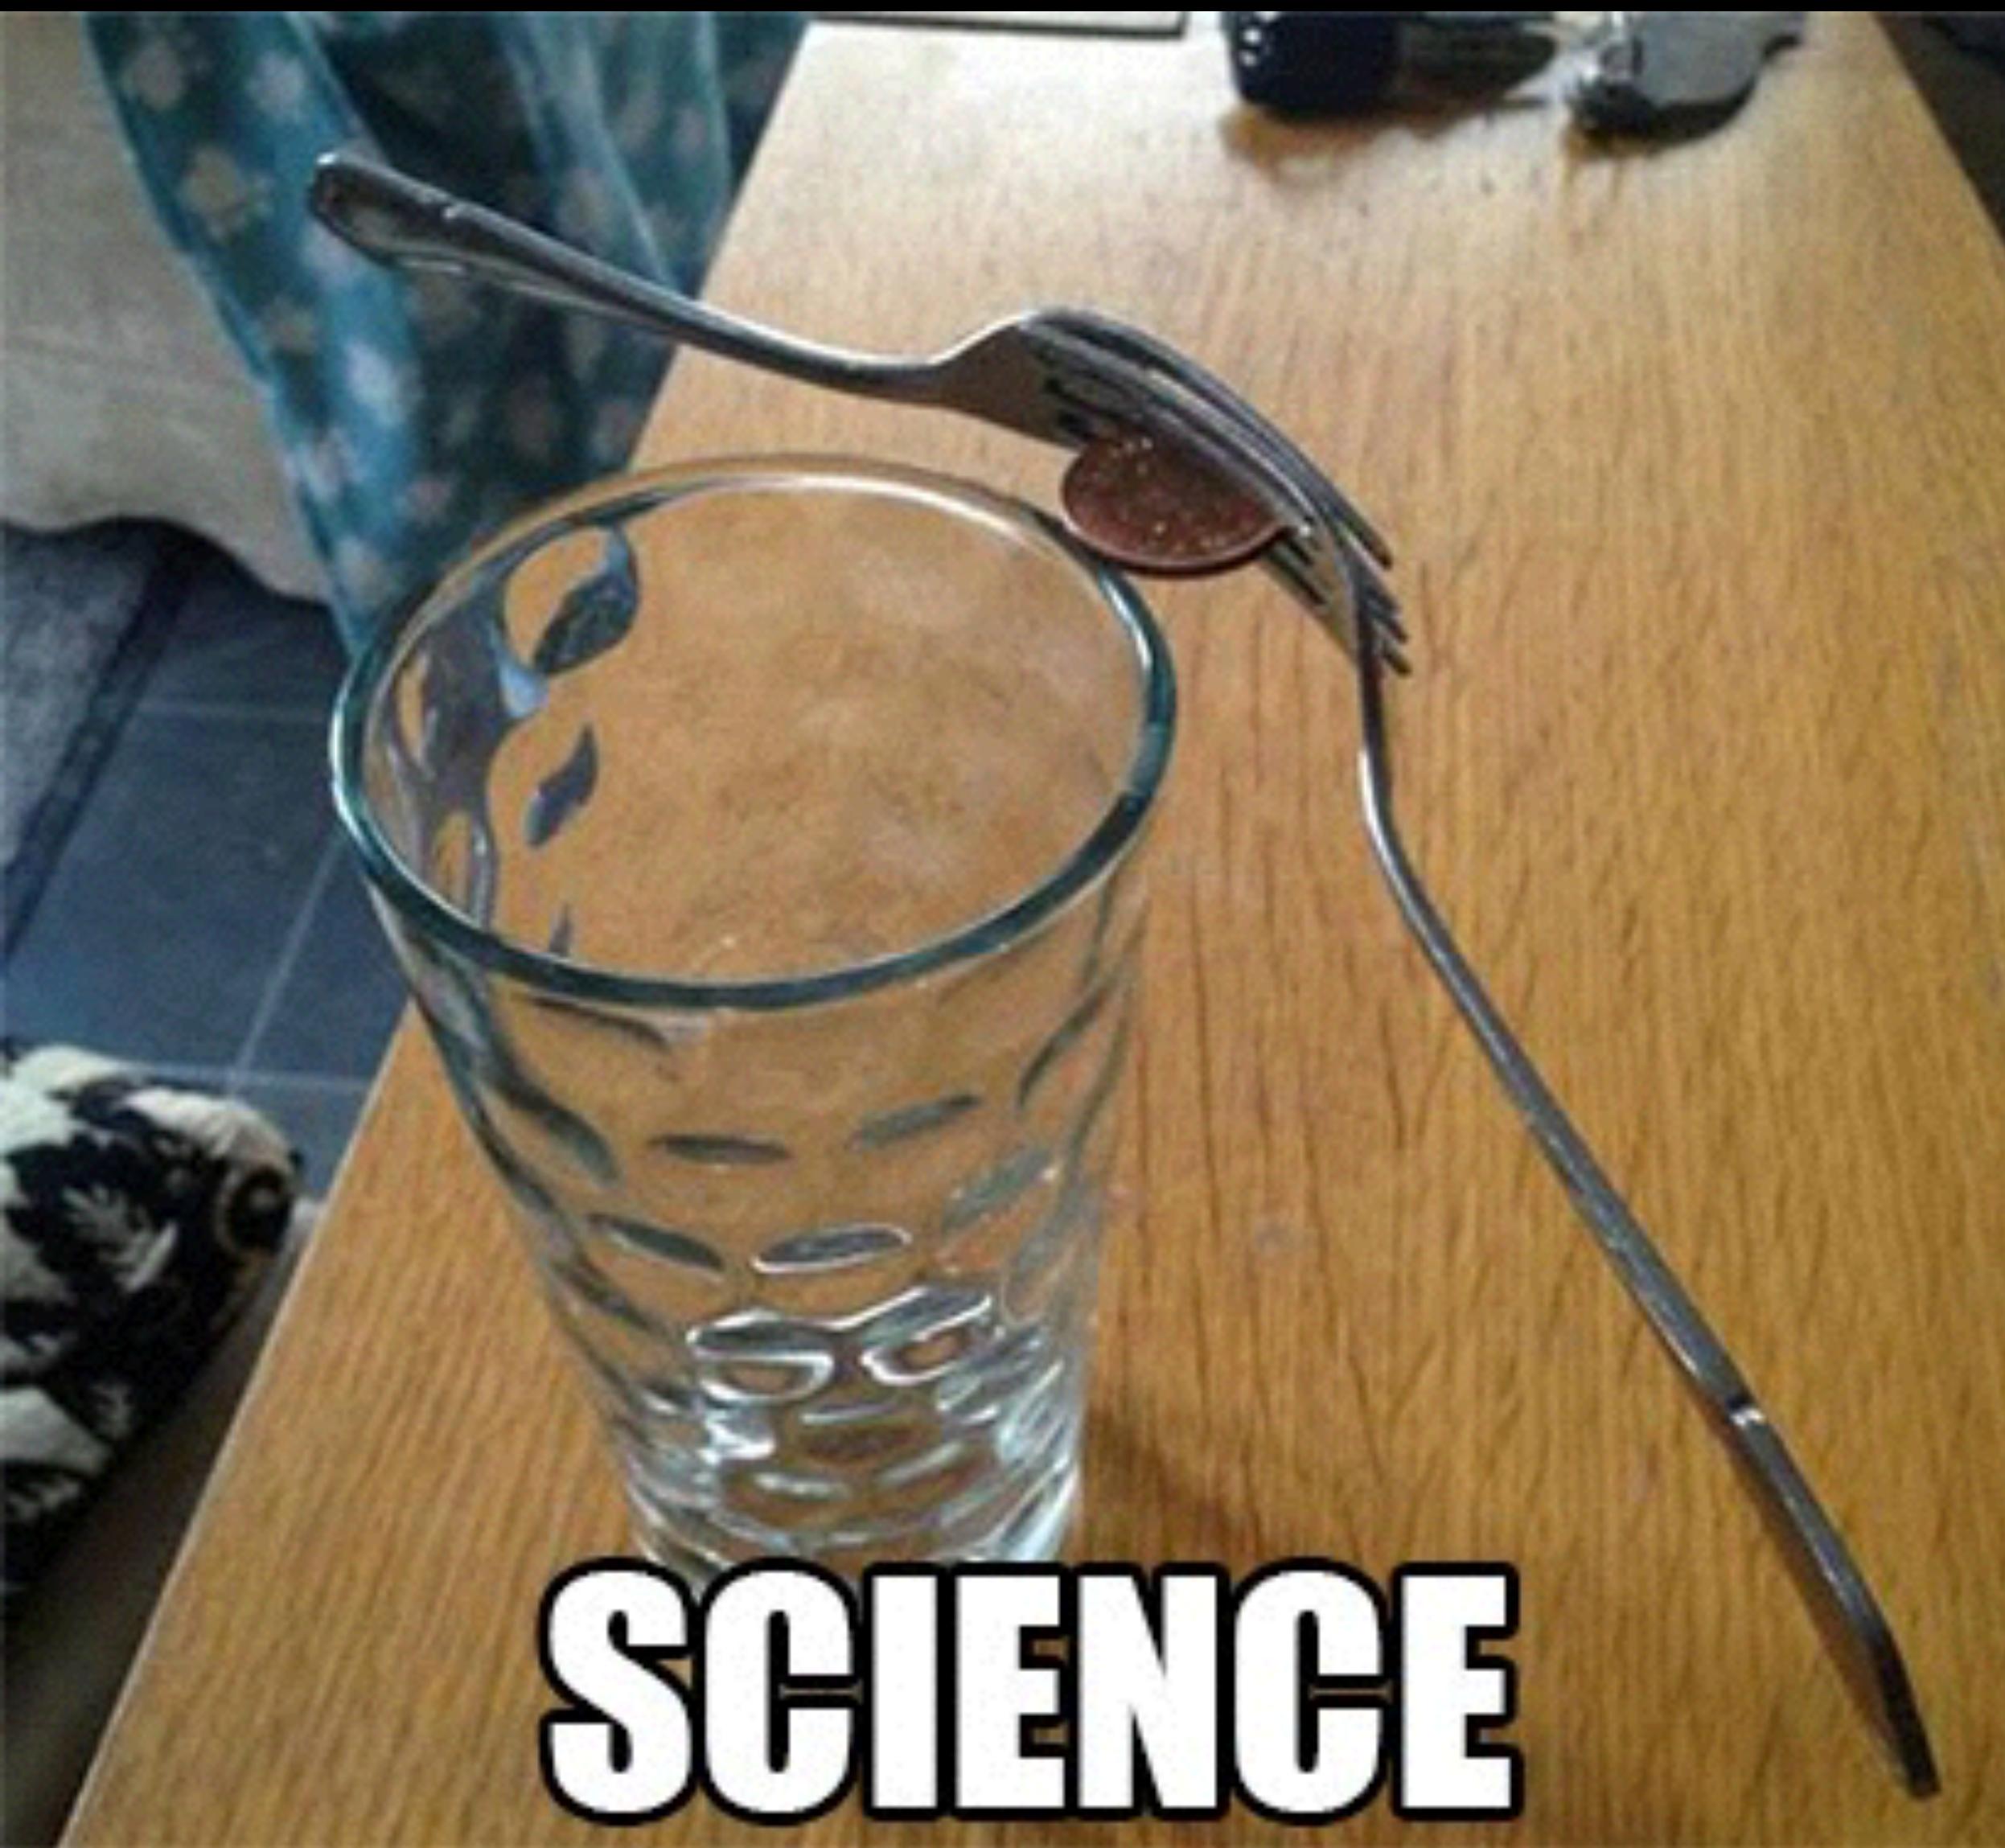
@pytest.mark.asyncio
async def test_do_something(capsys):
    before = time.monotonic()
    await do_something()
    delta_t = time.monotonic() - before
    assert -0.01 < delta_t < 0.01
```



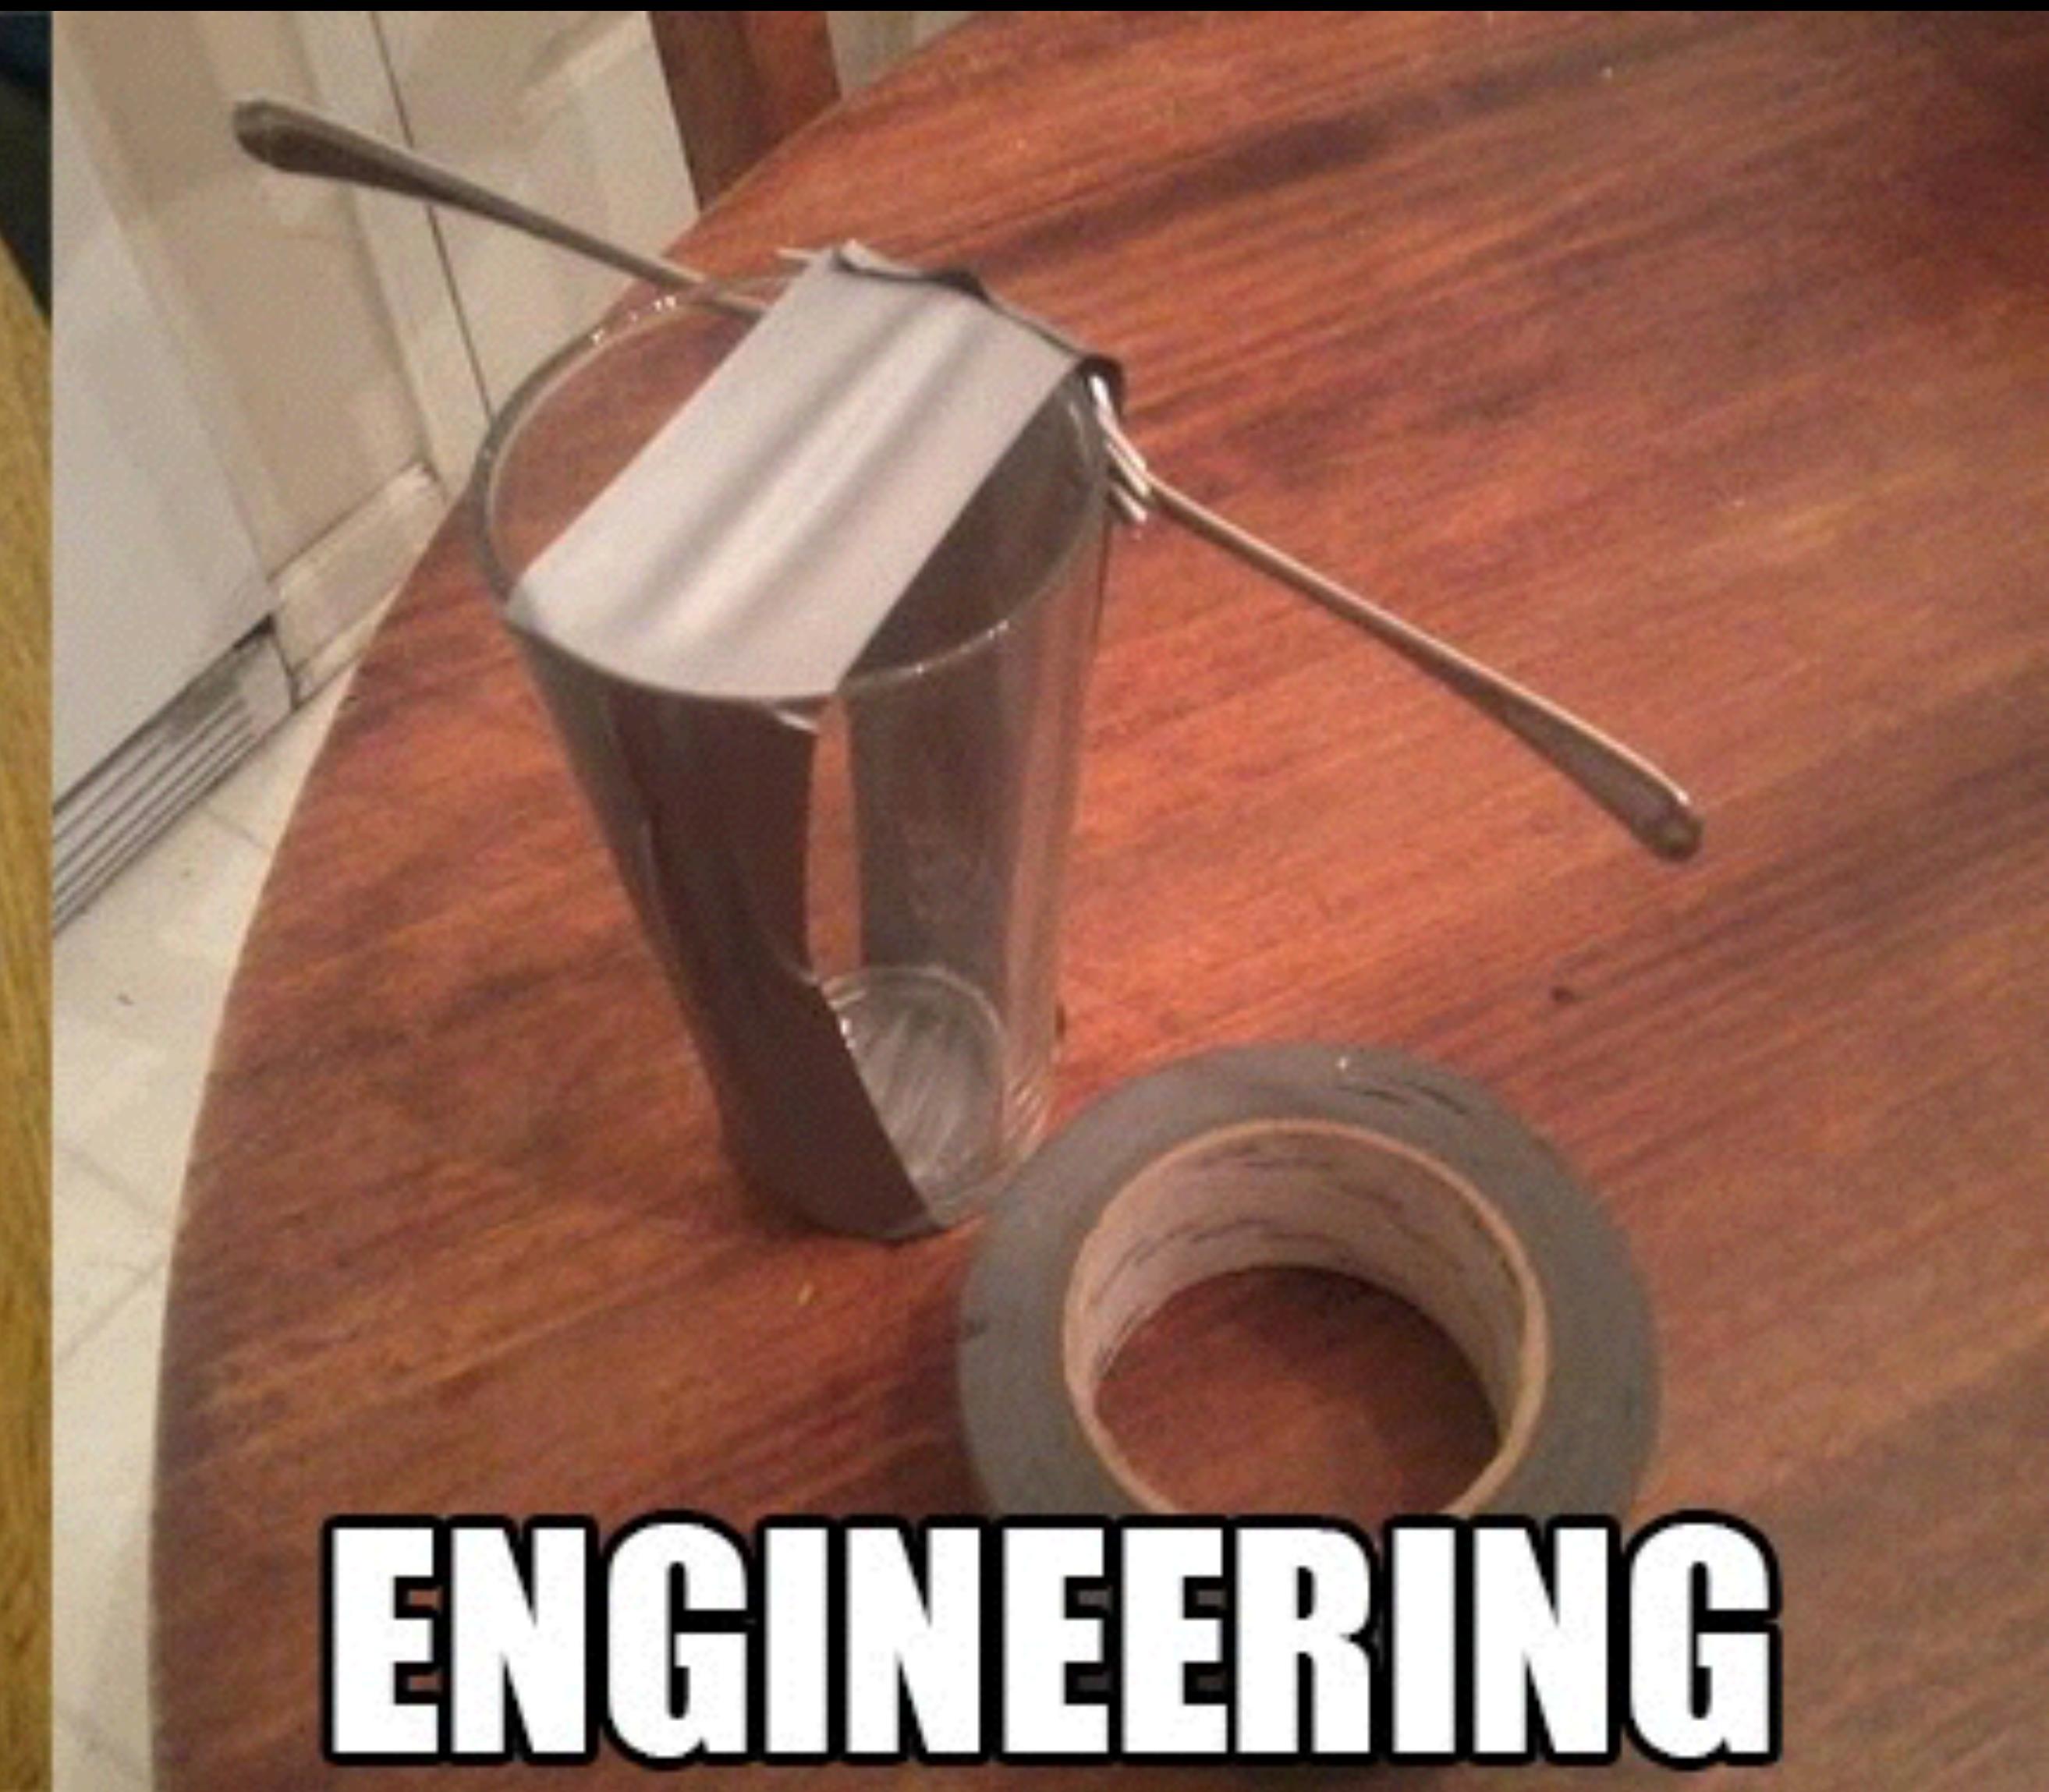
*But wait...*

**THERE'S  
MORE!!!**





**SCIENCE**



**ENGINEERING**

# asyncio

- Boilerplate
- Error-prone
- Immature API (I think)

```
import requests

def collect_contents(urls):
    contents = []
    for url in urls:
        resp = requests.get(url)
        if resp.status_code != 200:
            continue
        content = resp.text
        contents.append(content)
    return contents
```

```
import aiohttp

async def collect_contents(urls):
    contents = []
    with aiohttp.ClientSession() as session:
        for url in urls:
            async with session.get(url) as resp:
                if resp.status != 200:
                    continue
                content = await resp.text()
                contents.append(content)
    return contents
```

ONE DOES NOT SIMPLY

ADD “ASYNC” AND “AWAIT”

```
import aiohttp

async def collect_contents(urls):
    contents = []
    with aiohttp.ClientSession() as session:
        for url in urls:
            async with session.get(url) as resp:
                if resp.status != 200:
                    continue
                content = await resp.text()
                contents.append(content)
    return contents
```

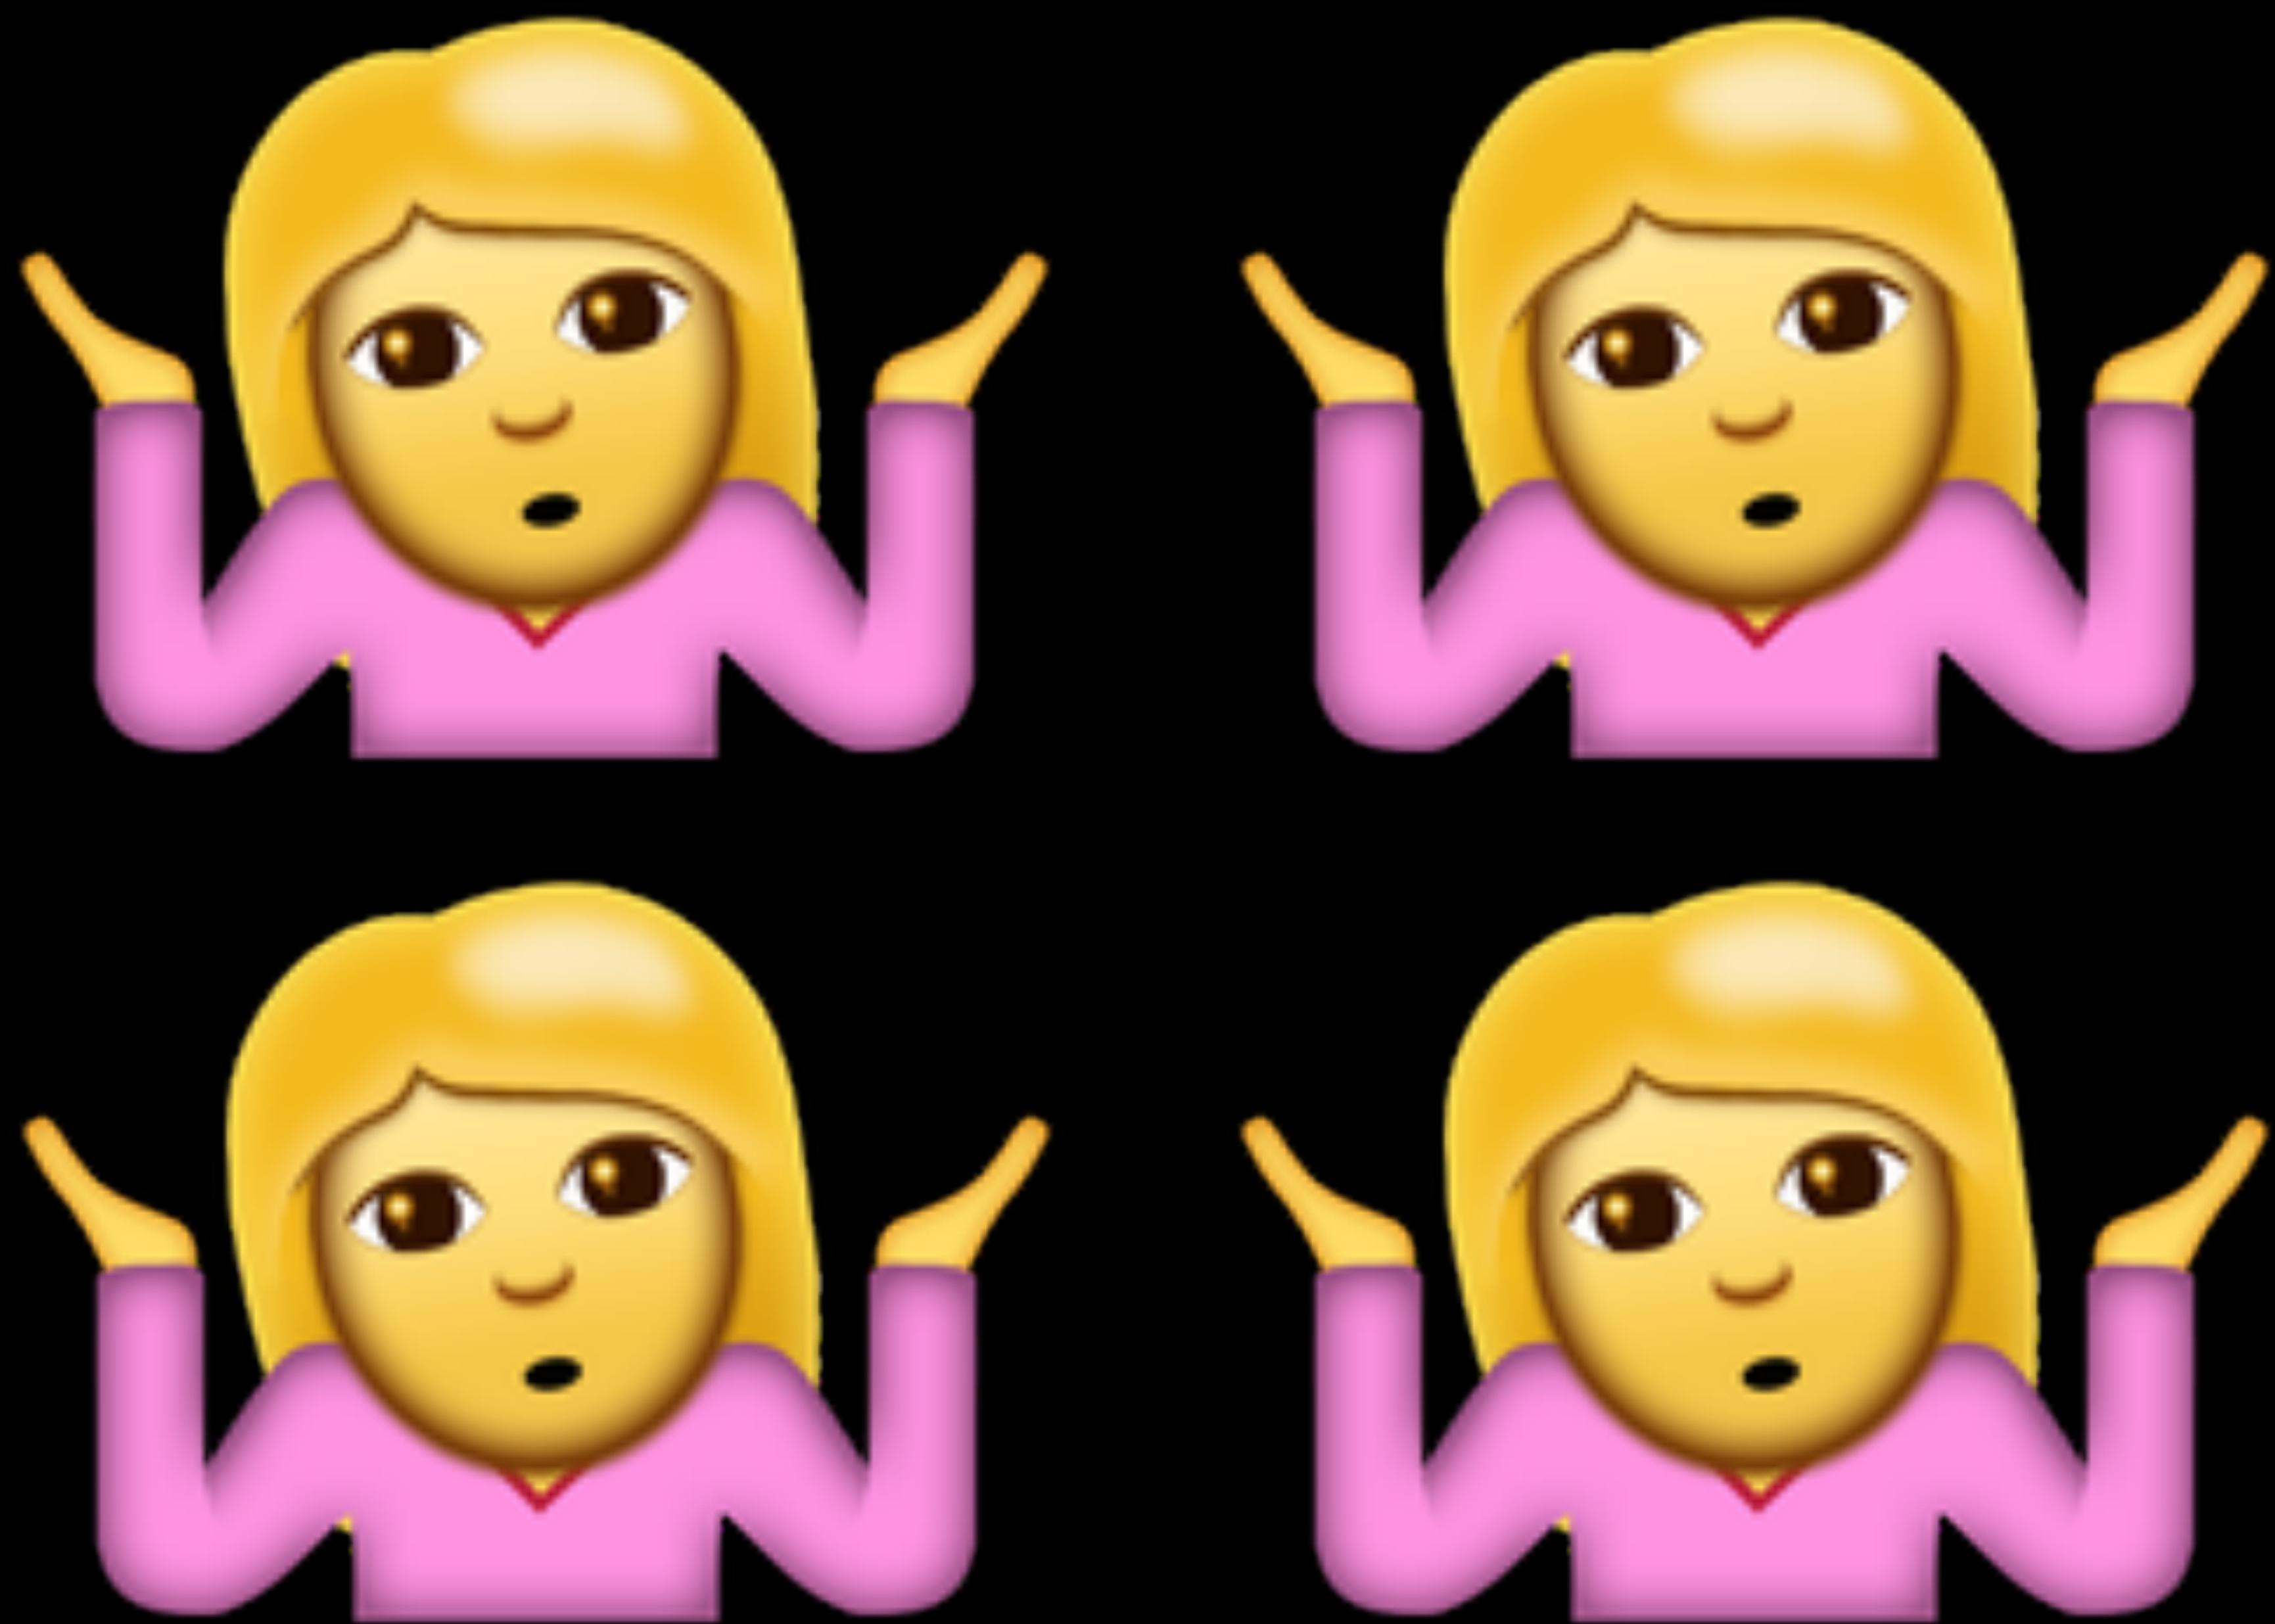


```
import aiohttp

async def collect_contents(urls):
    contents = []
    with aiohttp.ClientSession() as session:
        for url in urls:
            async with session.get(url) as resp:
                if resp.status != 200:
                    continue
                content = await resp.text()
                contents.append(content)
    return contents
```

```
import asyncio
import aiohttp

async def collect_contents(urls):
    coroutines = []
    with aiohttp.ClientSession() as session:
        for url in urls:
            async with session.get(url) as resp:
                if resp.status != 200:
                    continue
                coroutines.append(resp.text())
    contents = await asyncio.gather(*coroutines)
    return contents
```



## `coroutine asyncio.wait(futures, *, loop=None, timeout=None, return_when=ALL_COMPLETED)`

Wait for the Futures and coroutine objects given by the sequence *futures* to complete. Coroutines will be wrapped in Tasks. Returns two sets of Future: (done, pending).

The sequence *futures* must not be empty.

*timeout* can be used to control the maximum number of seconds to wait before returning. *timeout* can be an int or float. If *timeout* is not specified or `None`, there is no limit to the wait time.

*return\_when* indicates when this function should return. It must be one of the following constants of the `concurrent.futures` module:

## `asyncio.gather(*coros_or_futures, loop=None, return_exceptions=False)`

Return a future aggregating results from the given coroutine objects or futures.

All futures must share the same event loop. If all the tasks are done successfully, the returned future's result is the list of results (in the order of the original sequence, not necessarily the order of results arrival). If *return\_exceptions* is True, exceptions in the tasks are treated the same as successful results, and gathered in the result list; otherwise, the first raised exception will be immediately propagated to the returned future.

Cancellation: if the outer Future is cancelled, all children (that have not completed yet) are also cancelled. If any child is cancelled, this is treated as if it raised `CancelledError` - the outer Future is *not* cancelled in this case. (This is to prevent the cancellation of one child to cause other children to be cancelled.)

## `coroutine asyncio.wait(futures, *, loop=None, timeout=None, return_when=ALL_COMPLETED)`

Wait for the Futures and coroutine objects given by the sequence *futures* to complete. Coroutines will be wrapped in Tasks. Returns two sets of Future: (done, pending).

The sequence *futures* must not be empty.

*timeout* can be used to control the maximum number of seconds to wait before returning. *timeout* can be an int or float. If *timeout* is not specified or `None`, there is no limit to the wait time.

*return\_when* indicates when this function should return. It must be one of the following constants of the `concurrent.futures` module:

## `asyncio.gather(*coros_or_futures, loop=None, return_exceptions=False)`

Return a future aggregating results from the given coroutine objects or futures.

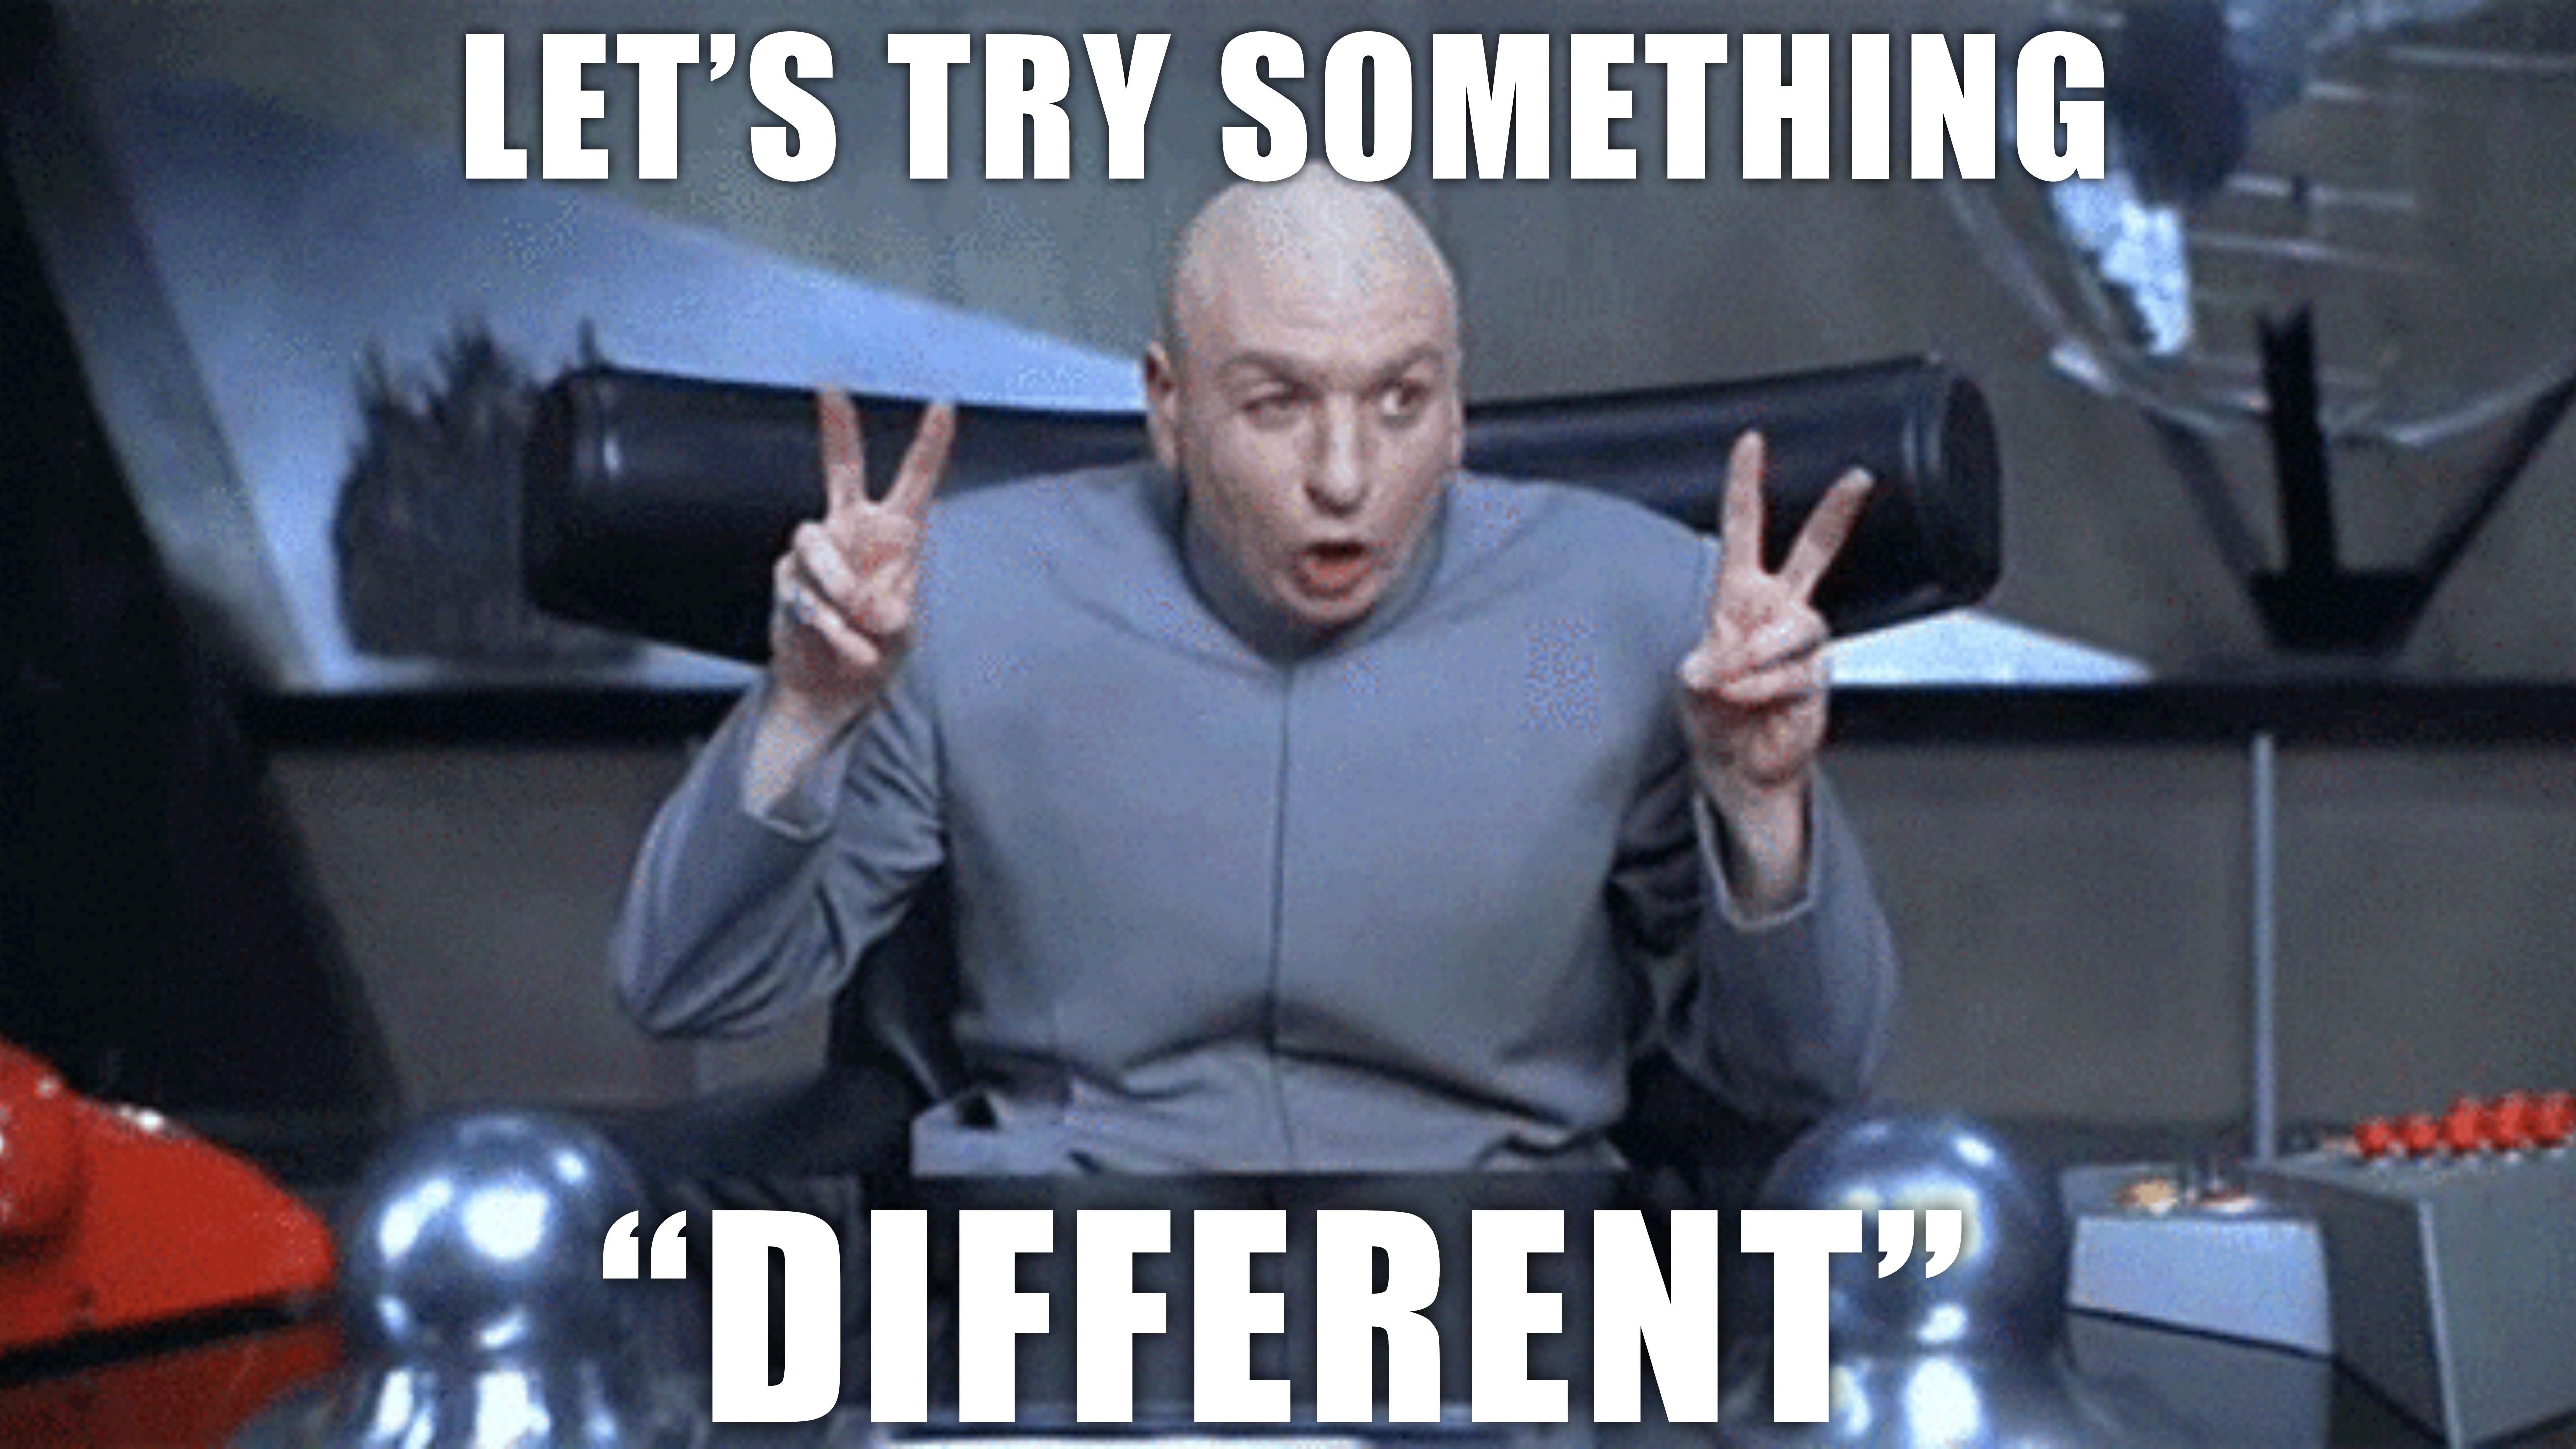
All futures must share the same event loop. If all the tasks are done successfully, the returned future's result is the list of results (in the order of the original sequence, not necessarily the order of results arrival). If *return\_exceptions* is True, exceptions in the tasks are treated the same as successful results, and gathered in the result list; otherwise, the first raised exception will be immediately propagated to the returned future.

Cancellation: if the outer Future is cancelled, all children (that have not completed yet) are also cancelled. If any child is cancelled, this is treated as if it raised `CancelledError` - the outer Future is *not* cancelled in this case. (This is to prevent the cancellation of one child to cause other children to be cancelled.)



**WHY!!???**

# LET'S TRY SOMETHING



## “DIFFERENT”

# Alternatives

- concurrent & multiprocessing
- *Greenlets*
  - Similar idea, but less infectious
  - C extension
- threading
  - Standard I/O

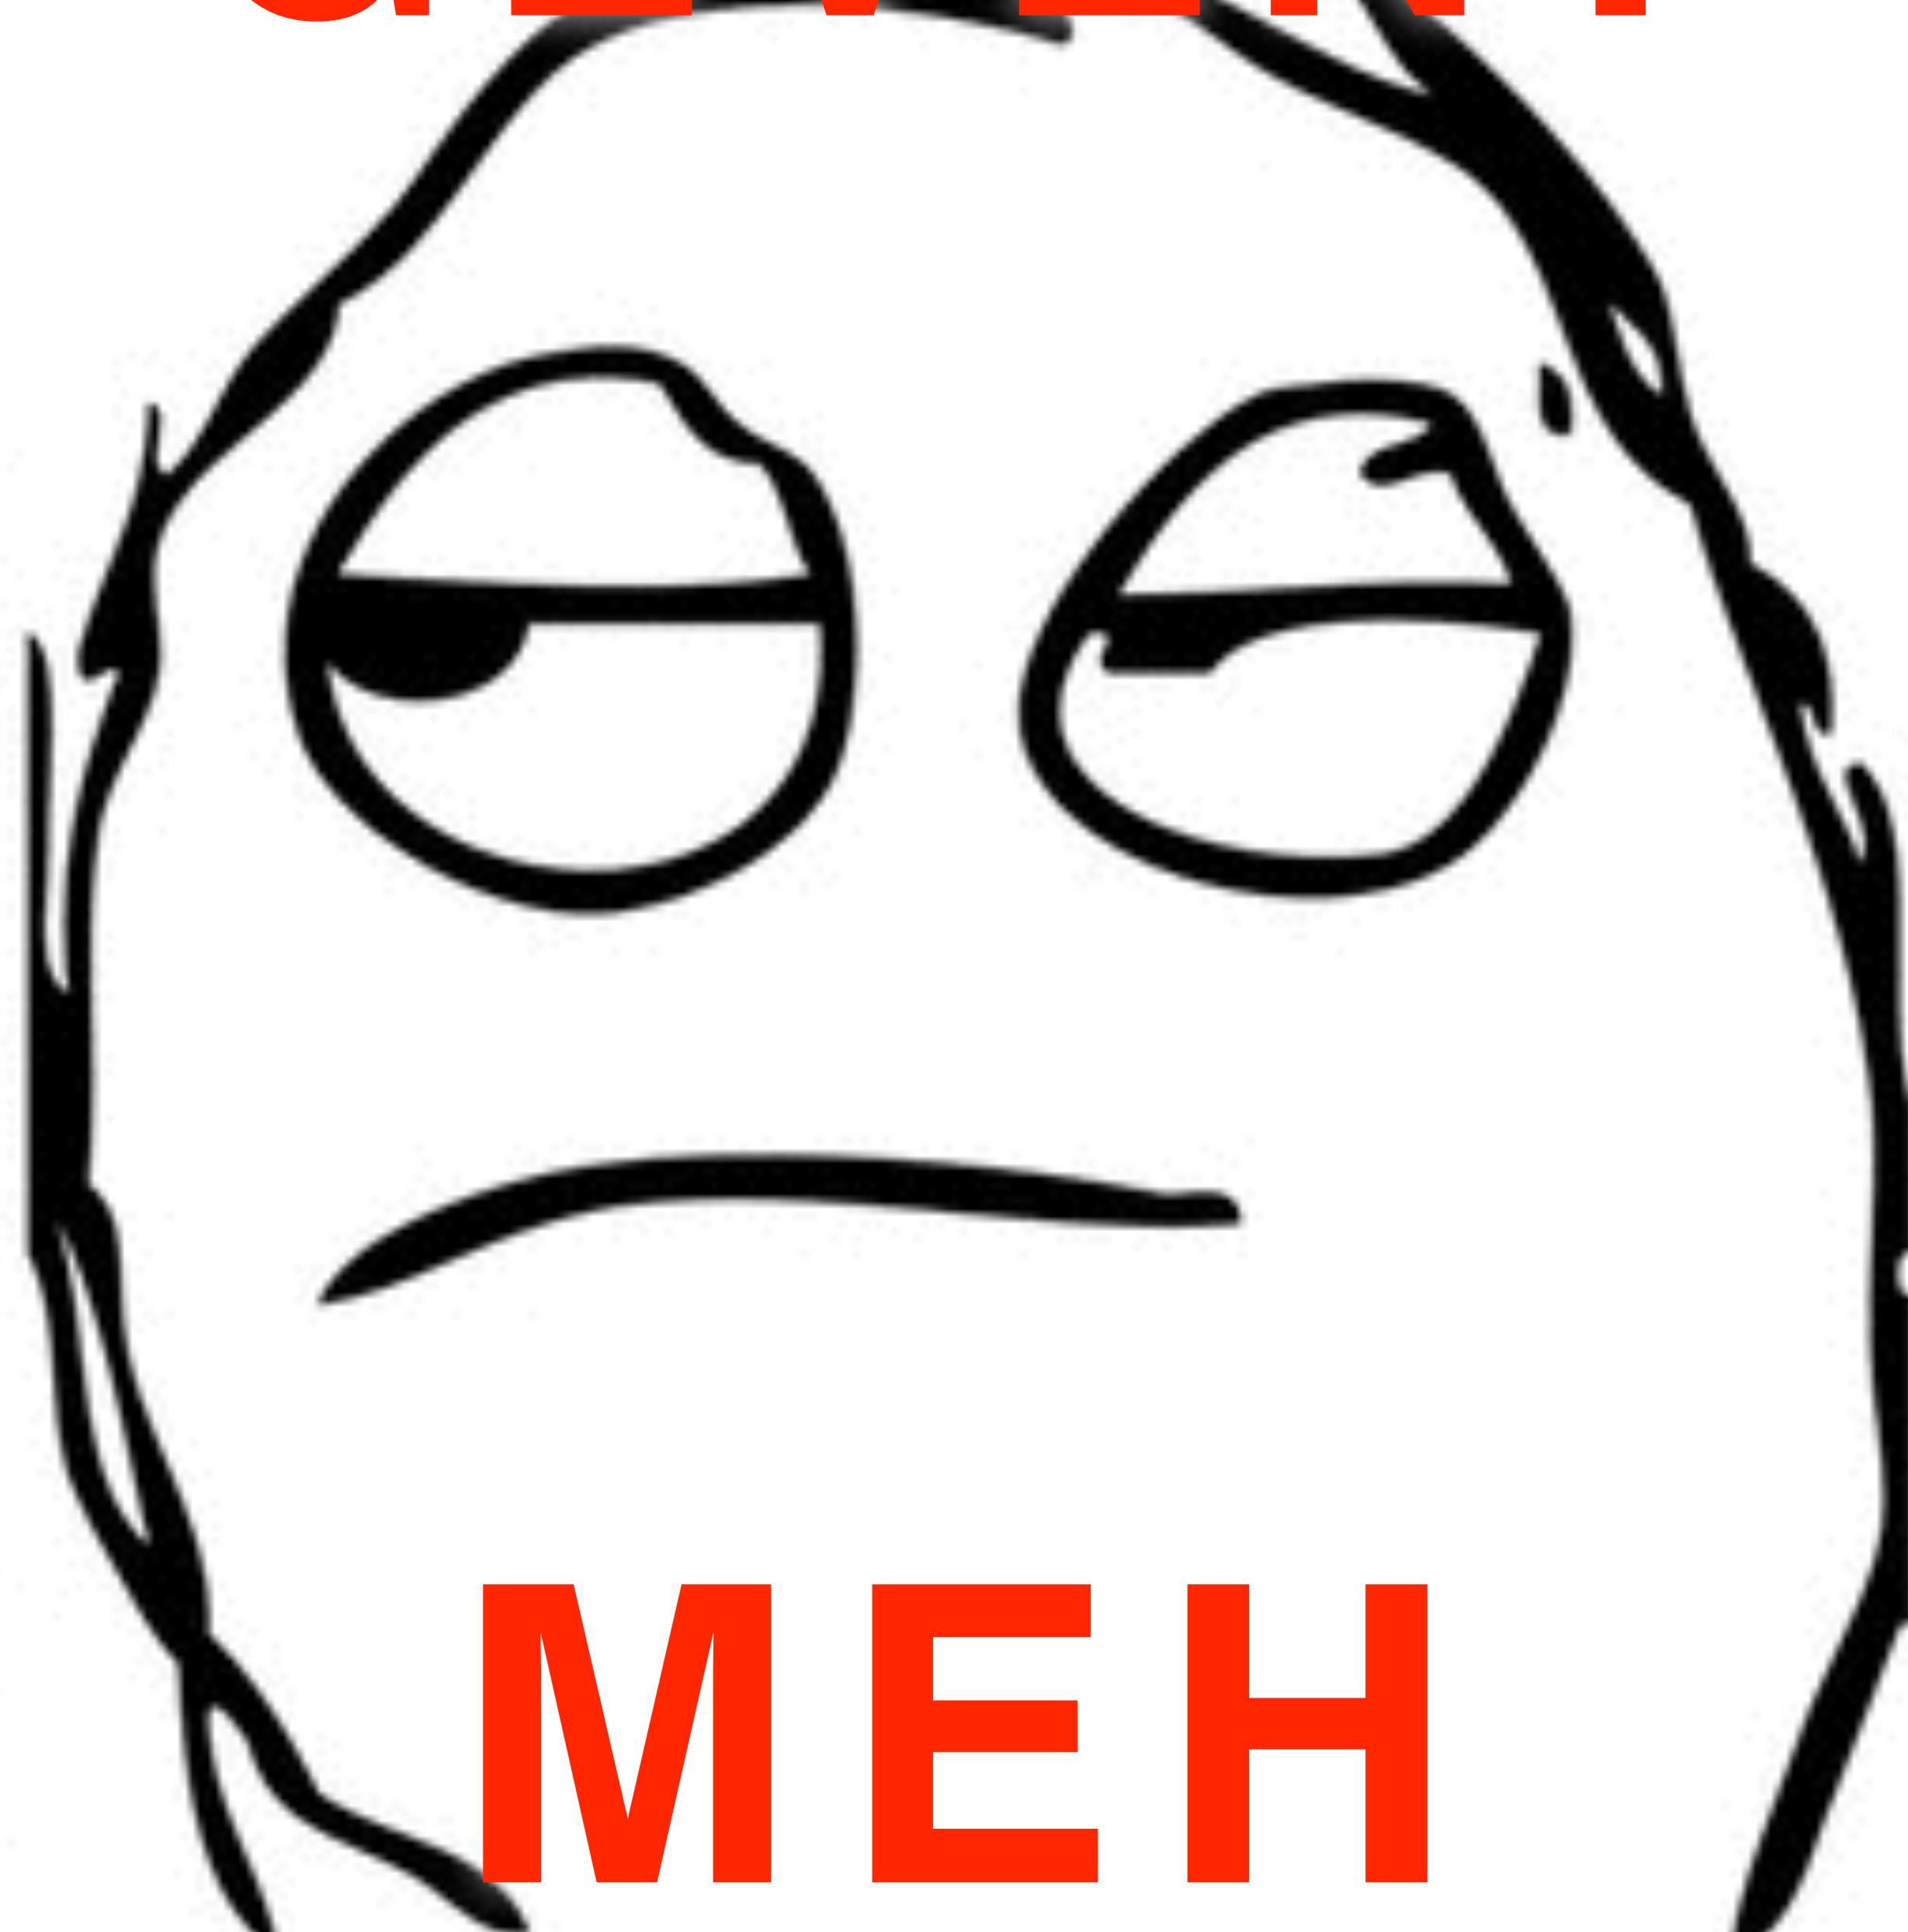
# greenlet: Lightweight concurrent programming

## Motivation

The “greenlet” package is a spin-off of [Stackless](#), a version of CPython that supports micro-threads called “tasklets”. Tasklets run pseudo-concurrently (typically in a single or a few OS-level threads) and are synchronized with data exchanges on “channels”.

A “greenlet”, on the other hand, is a still more primitive notion of micro-thread with no implicit scheduling; coroutines, in other words. This is useful when you want to control exactly when your code runs. You can build custom scheduled micro-threads on top of greenlet; however, it seems that greenlets are useful on their own as a way to make advanced control flow structures. For example, we can recreate generators; the difference with Python’s own generators is that our generators can call nested functions and the nested functions can yield values too. (Additionally, you don’t need a “yield” keyword. See the example at <http://greenlet.readthedocs.io>.)

GEVENT



MEH

# ROUTINES



# ROUTINES EVERYWHERE

```
package main
import ("fmt"; "time")

func doSomething() {
    time.Sleep(2 * time.Second)
    fmt.Println(time.Now(), "Slept")
}

func main() {
    doSomething()
    fmt.Println(time.Now(), "OK")
}
```

```
00:00:00 Slept
00:00:00 OK
```

```
package main
import ("fmt"; "time")

func doSomething() {
    time.Sleep(2 * time.Second)
    fmt.Println(time.Now(), "Slept")
}

func main() {
    go doSomething()
    fmt.Println(time.Now(), "OK")
}
```

```
00:00:00 OK
```

```
package main
import ("fmt"; "time")

var sem = make(chan bool)
func doSomething() {
    time.Sleep(2 * time.Second)
    fmt.Println(time.Now(), "Slept")
    sem <- true
}

func main() {
    go doSomething()
    fmt.Println(time.Now(), "OK")
    <-sem
}
```

```
00:00:00 OK
00:00:02 Slept
```

# The Go Model

- Boilerplate
- Error-prone
- Immature API (I think)

```
"""Do something. Synchronous version."""
```

```
import datetime  
import time
```

```
def do_something():  
    time.sleep(2)  
    print(datetime.datetime.now(), 'Slept')
```

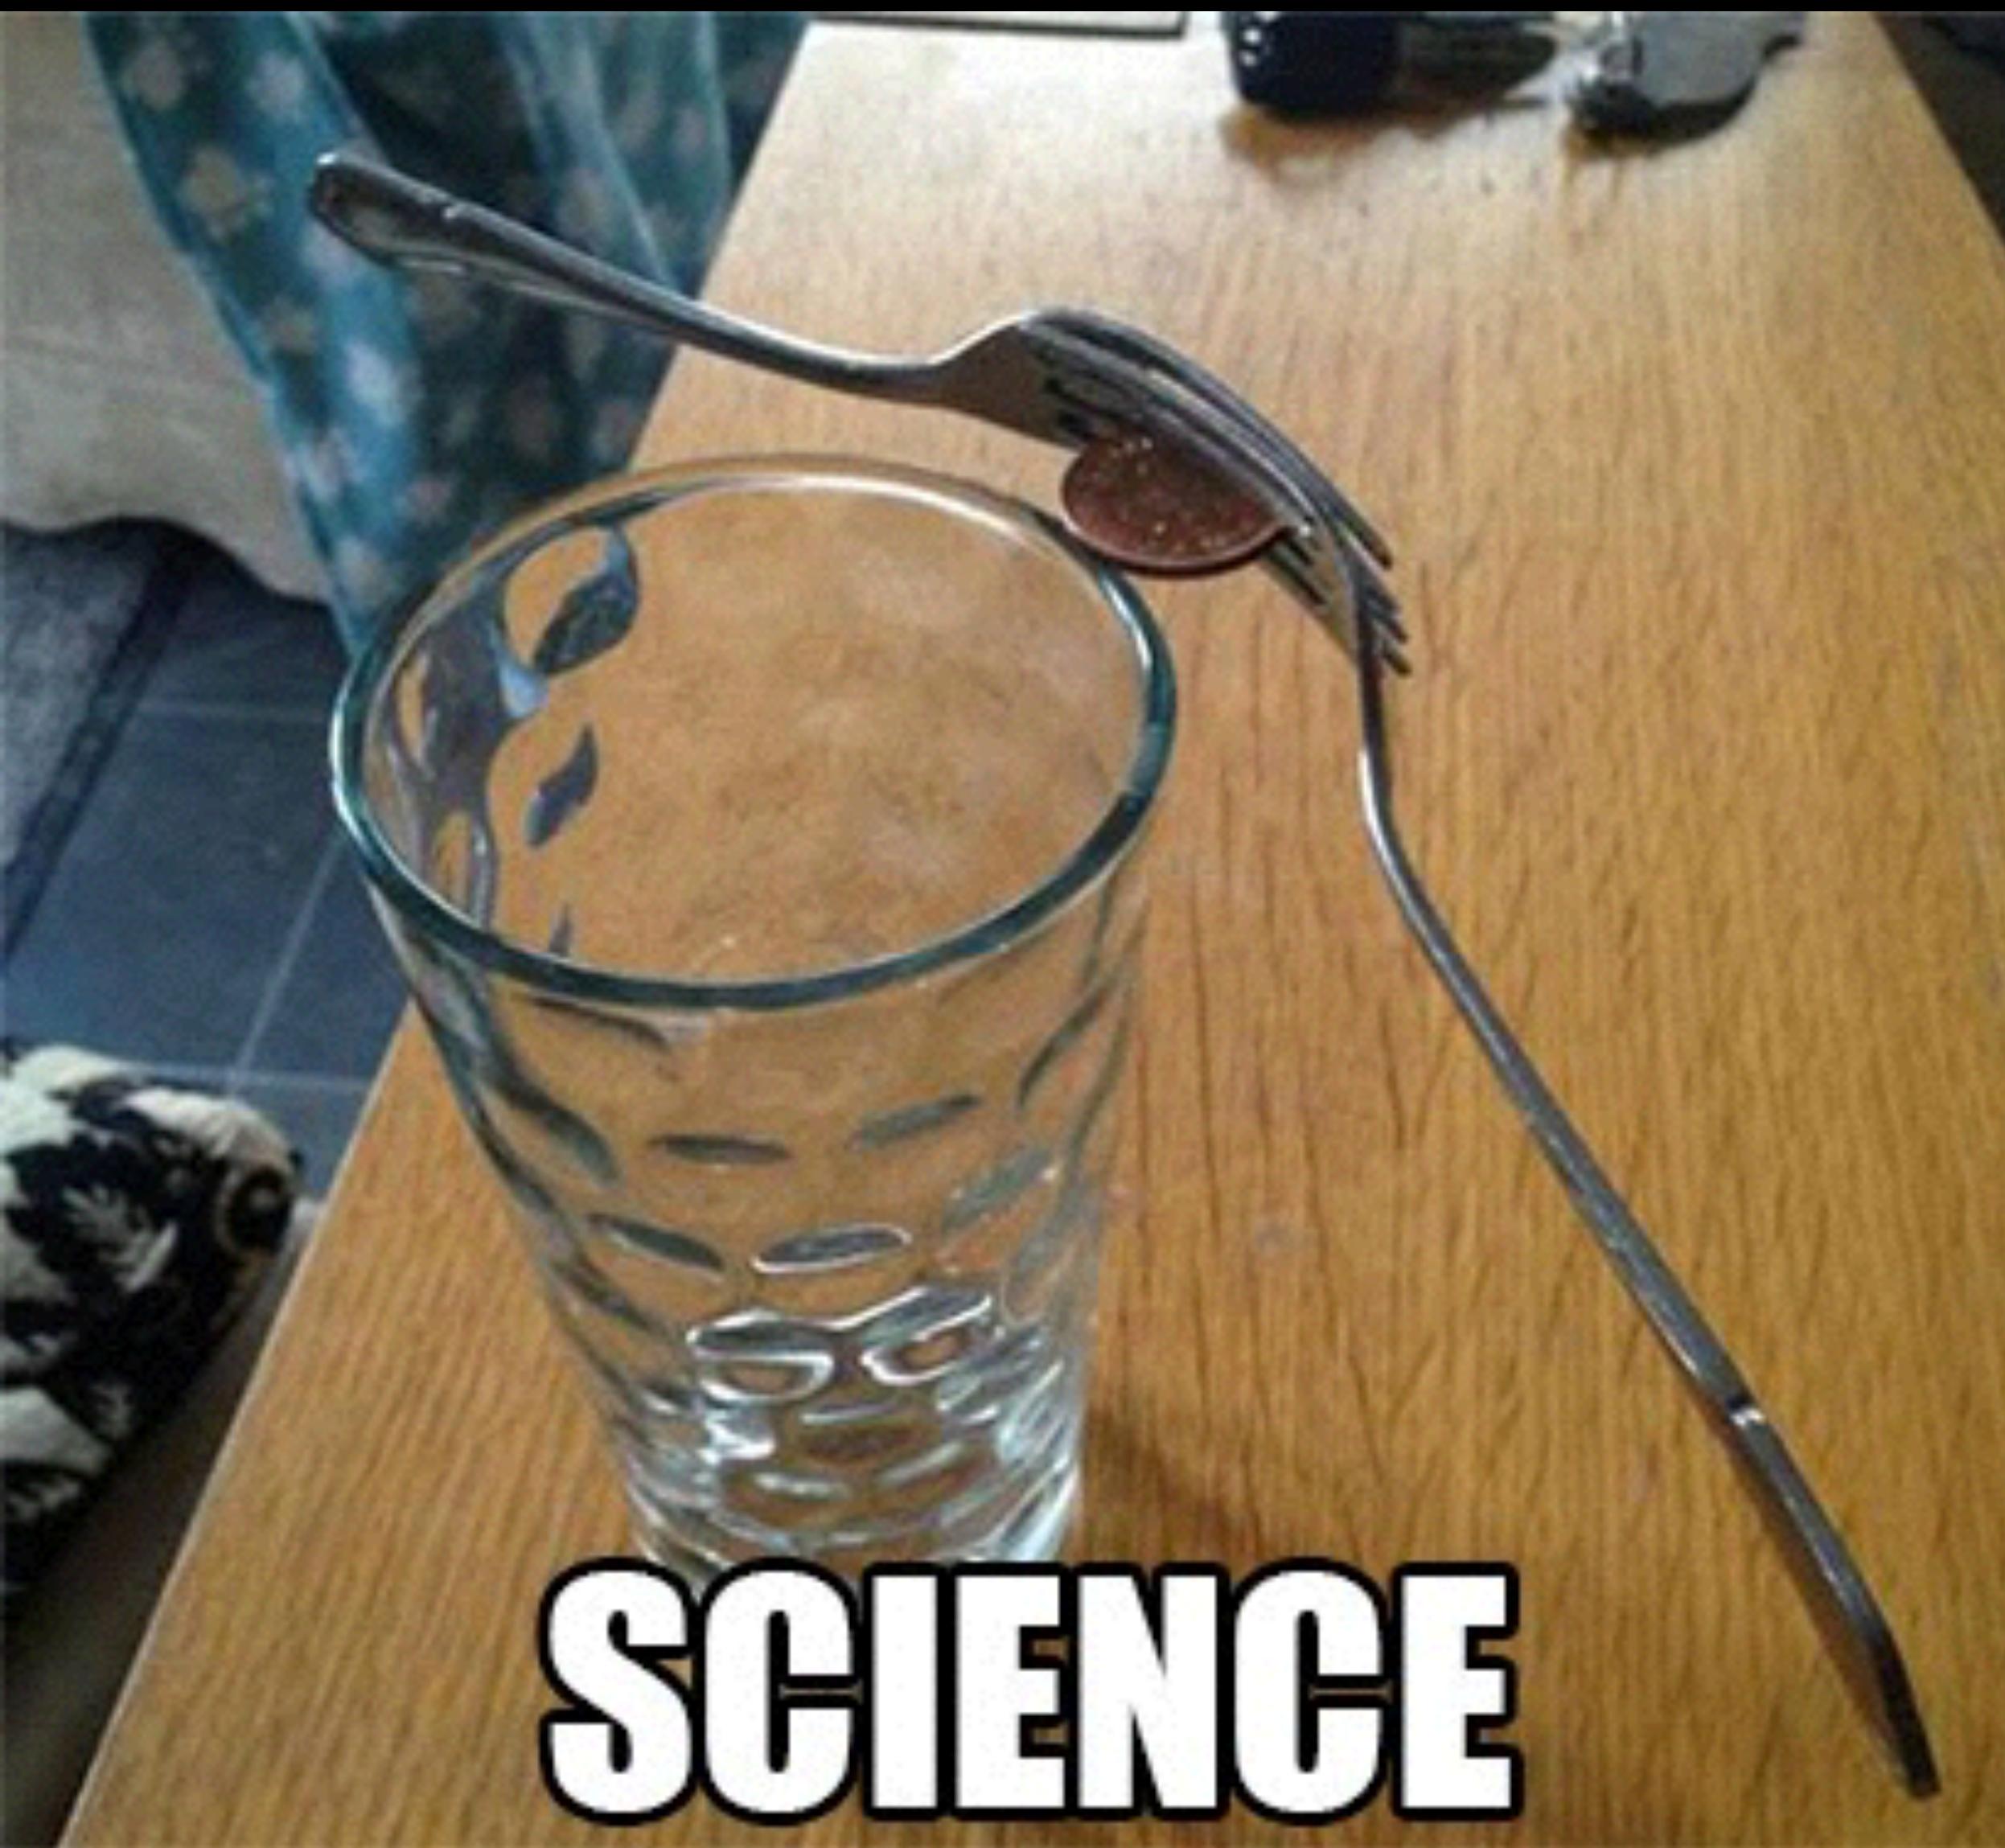
```
do_something()  
print(datetime.datetime.now(), 'Done')
```

```
"""Do something. Asynchronous version."""
```

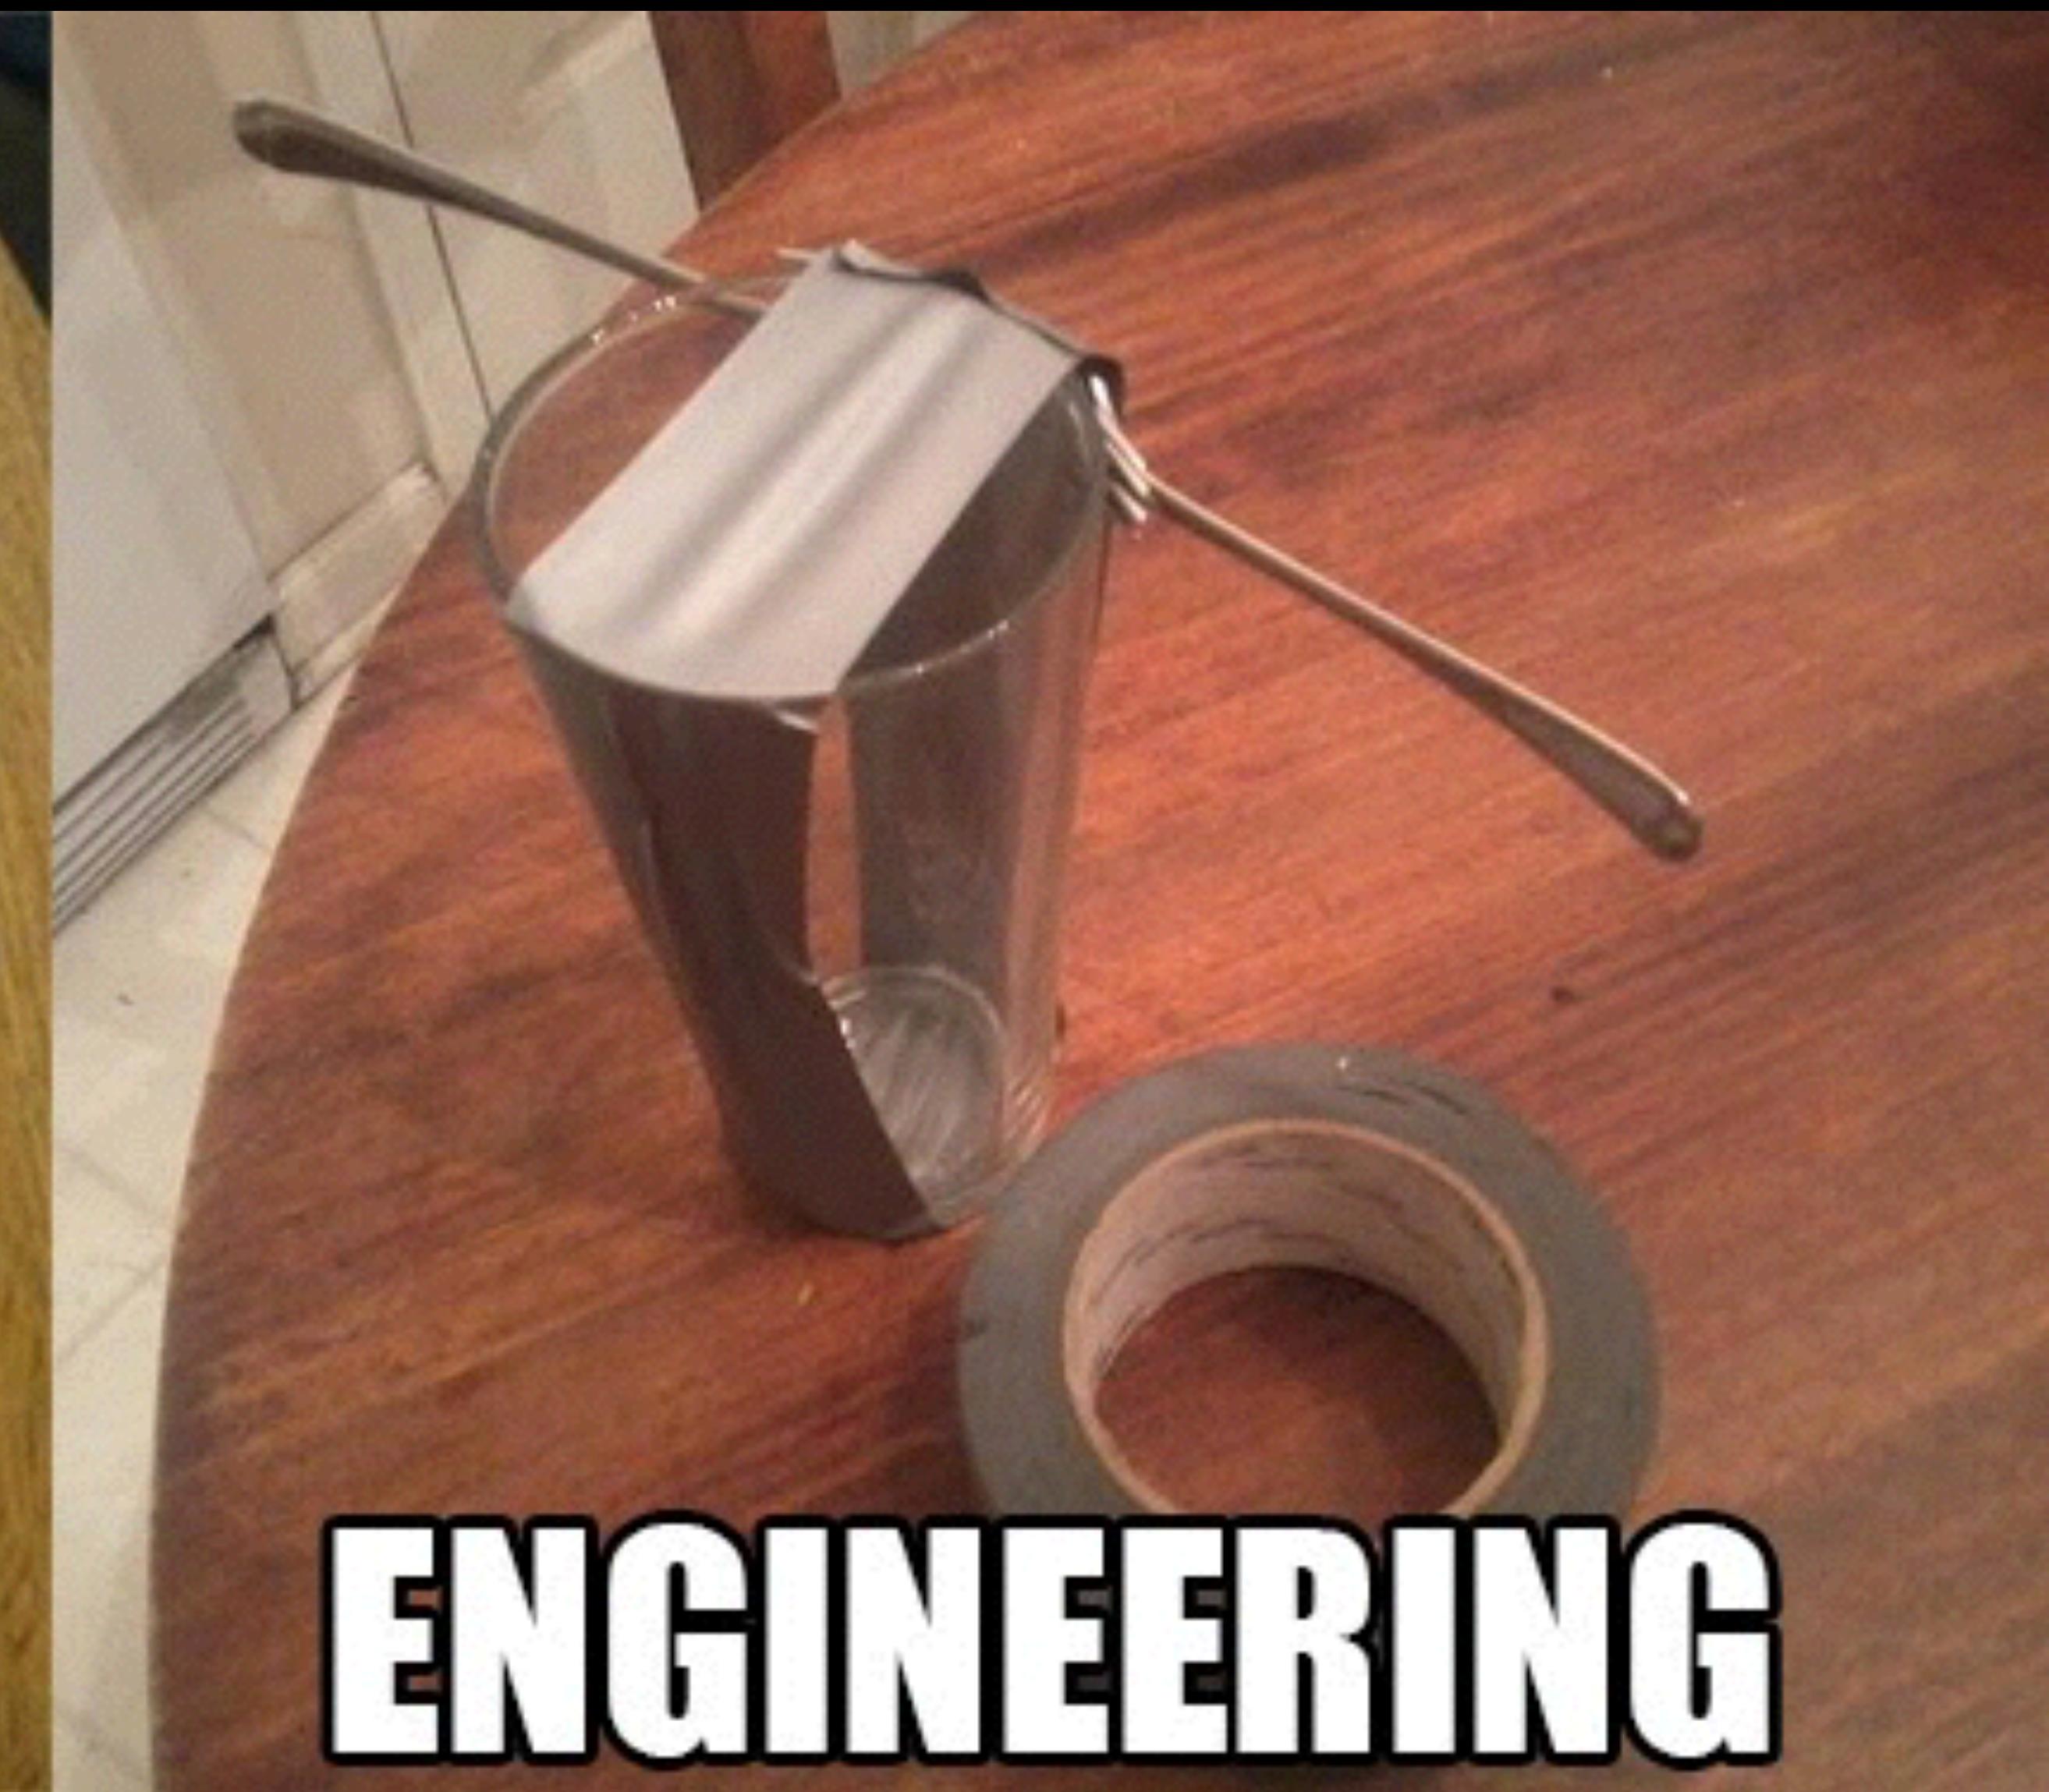
```
import asyncio
import datetime

async def do_something():
    await asyncio.sleep(2)
    print(datetime.datetime.now(), 'Slept')

loop = asyncio.get_event_loop()
task = loop.create_task(do_something())
print(datetime.datetime.now(), 'Done')
loop.run_until_complete(asyncio.wait([task]))
loop.close()
```



**SCIENCE**



**ENGINEERING**

```
"""Do something. Synchronous version."""
```

```
import datetime  
import time
```

```
def do_something():  
    time.sleep(2)  
    print(datetime.datetime.now(), 'Slept')
```

```
do_something()  
print(datetime.datetime.now(), 'Done')
```

"""What if I can just write this?"""

```
import asyncio  
import datetime  
import time
```

```
async def do_something():  
    await time.sleep(2)  
    print(datetime.datetime.now(), 'Slept')
```

```
await do_something()  
print(datetime.datetime.now(), 'Done')  
asyncio.run_event_loop()
```

I know, it's not really  
possible.

# At least we can dream.

Or wait until Python 6.0?

# But Seriously

- Asynchrony is not the silver bullet
- It makes you jump through loops
- There are alternatives
- Fingers crossed

# Recap

- Rant
- Moar rant
- Susceptible advice

# But Seriously

- Asynchrony is not the silver bullet
- It makes you jump through loops
- There are alternatives

IT'S DANGEROUS TO GO  
ALONE! TAKE THIS.

