

APAC Korea 2016

Linux Kernel Instrumentation in Python

Feng Li (李枫)

knewlife2015@yahoo.com

Aug 14, 2016

Agenda

I. eBPF & BCC

- eBPF
- BCC

II. In-Kernel Virtual Machine

- KPlugs
- Application

III. Python-assisted Data Center Dev

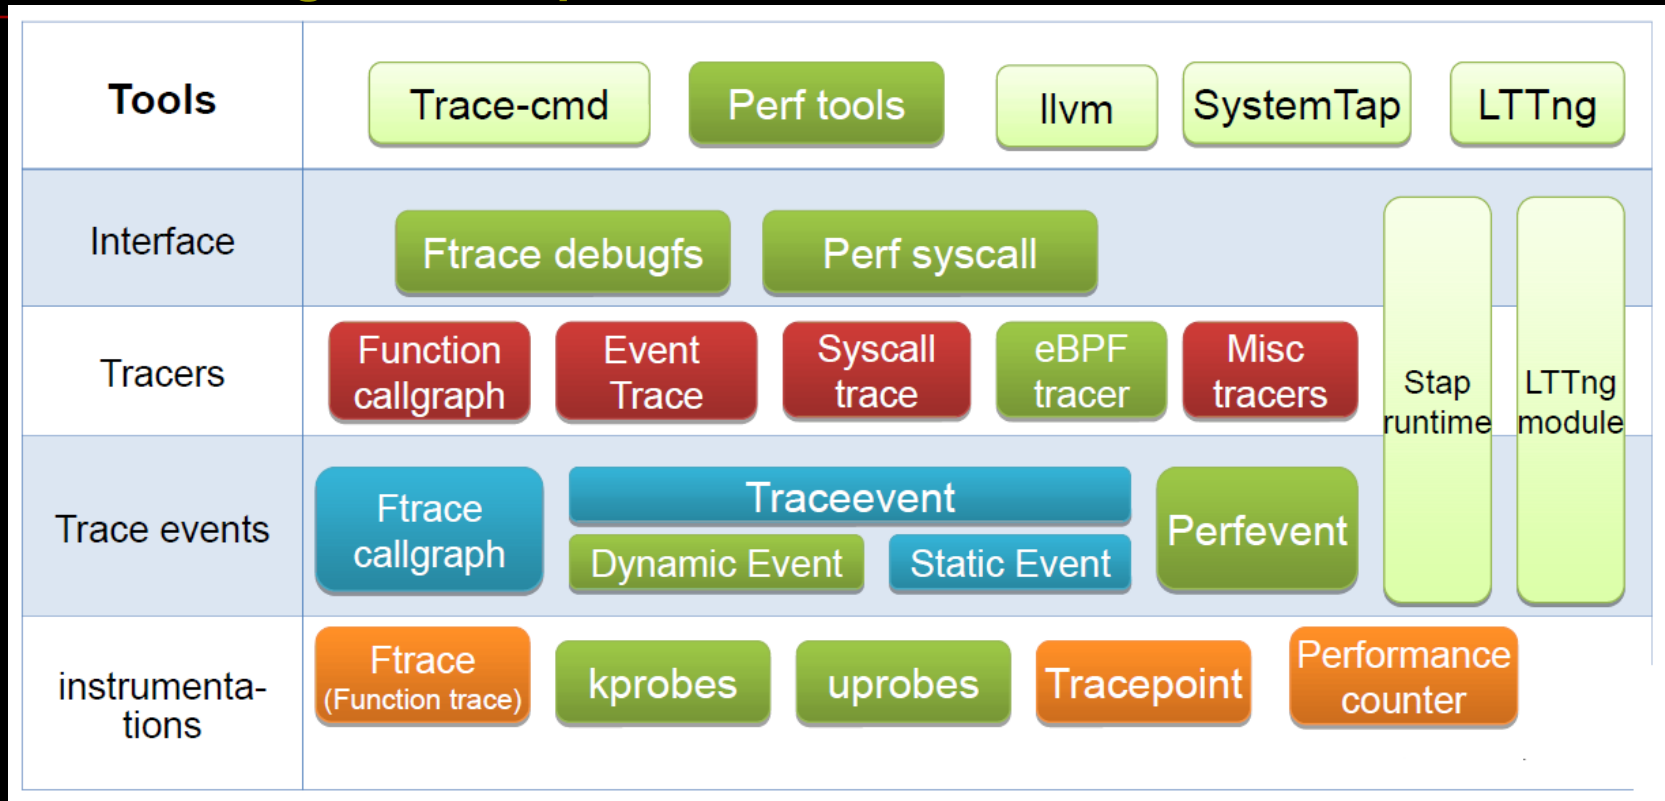
- HPC
- Software-Defined Everything

IV. Summary

I. eBPF & BCC

Background

■ The Tracing Landscape



Source: <http://tracingsummit.org/w/images/8/8c/TracingSummit2015-DynamicProbes.pdf>

1) eBPF

- https://en.wikipedia.org/wiki/Berkeley_Packet_Filter
- <http://lwn.net/>

BPF (Berkeley Packet Filter, aka cBPF)

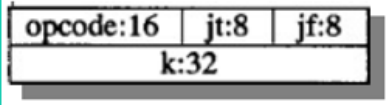
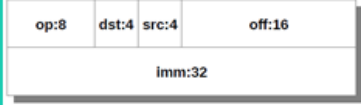
- Introduced in kernel 2.1.75 (1997)
- Originally designed for packet filtering (tcpdump...)
- Apply for seccomp filters, traffic control...

eBPF (extended BPF)

- Since Linux Kernel v3.15 and ongoing
- Aims at being a universal in-kernel virtual machine, it changes the old ways for Kernel instrumentation
- <https://lwn.net/Articles/655544/>

BPF for tracing is currently a hot area, Starovoitov said. It is a better alternative to [SystemTap](#) and runs two to three times faster than Oracle's [DTrace](#). Part of that speed comes from LLVM's optimizations plus the kernel's internal just-in-time compiler for BPF bytecode.

Comparison

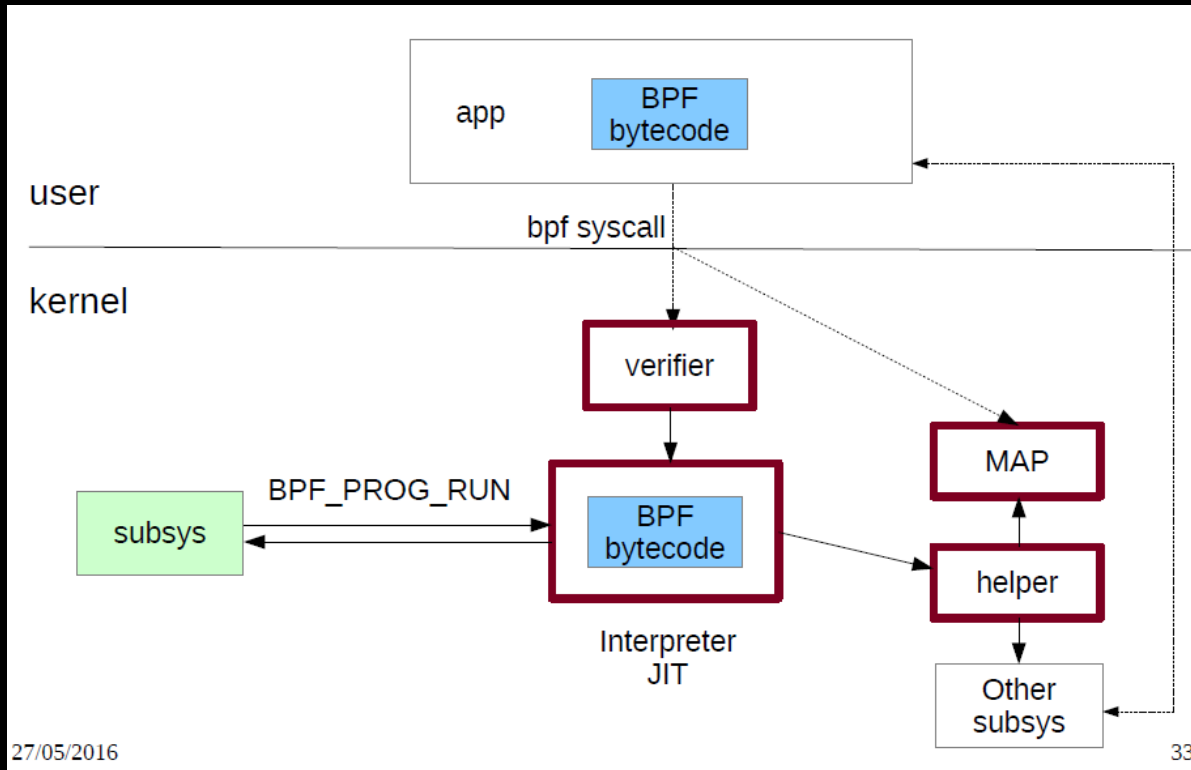
	cBPF	eBPF
Register	Two 32 bit registers: A: accumulator X: indexing	Eleven 64 bit registers: R0: return value/exit value R1-R5: arguments R6-R9: callee saved registers R10: read-only frame pointer
Instruction	~30 	~90 
JIT	Support	Support (better mapping with newer architectures for JITing)
Toolchain	GCC, tools/net	LLVM eBPF backend
Platform	x86_64, ARM, ARM64, SPARC, PowerPC, MIPS and s390	x86-64, aarch64, s390x
System Call		<pre>#include <linux/bpf.h> int bpf(int cmd, union bpf_attr *attr, unsigned int size);</pre> <p>(CALL, MAP, LOAD...)</p>

Dev

■ Dev Methods

- 1) write directly using eBPF assembly
- 2) Write it using C, and compile with LLVM
- 3) BCC

...



Source: <http://www.slideshare.net/vh21/meet-cutebetweenebpfandtracing>

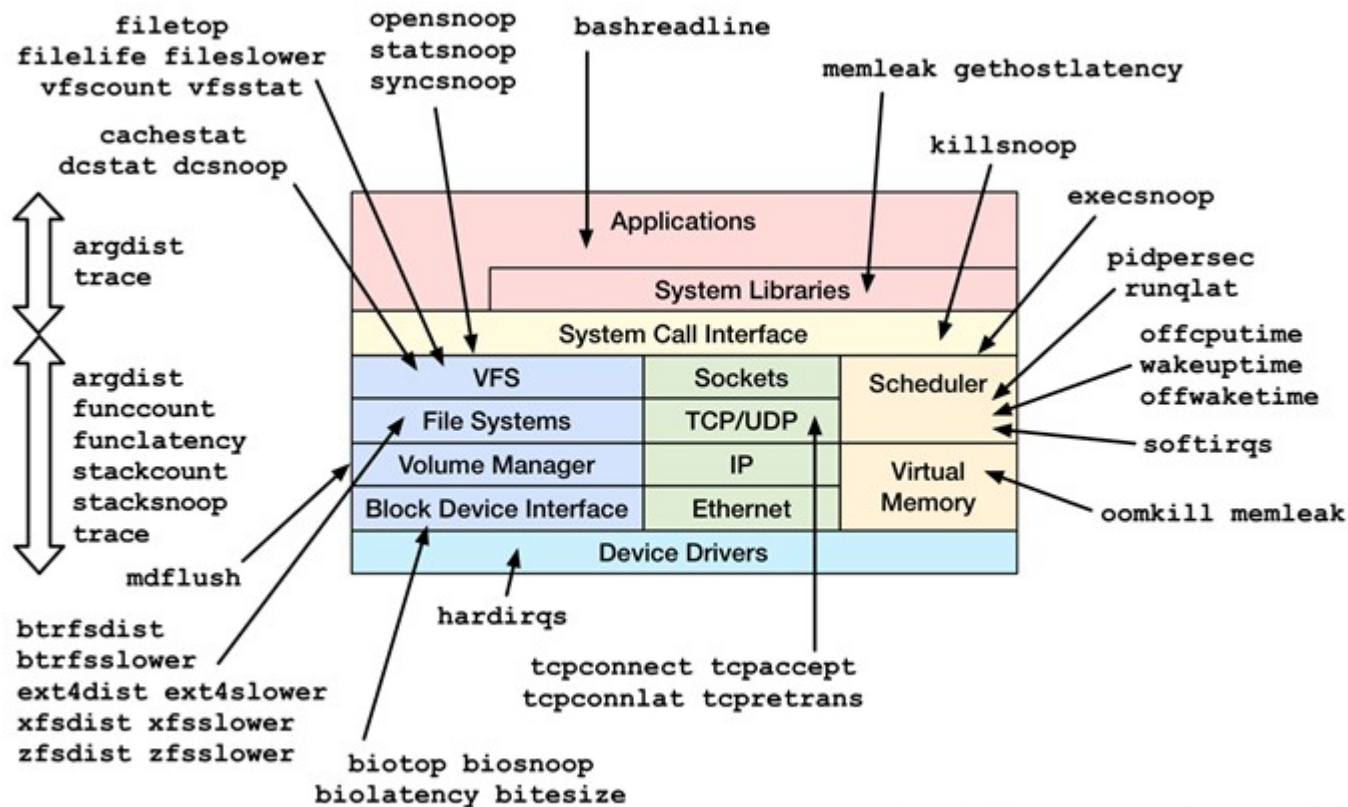
2) BCC (BPF Compiler Collection)

- <https://iovisor.github.io/bcc/>
- <https://github.com/iovisor/bcc.git>

A toolkit with Python/Lua frontend for compiling, loading, and executing BPF programs, which allows user-defined instrumentation on a live kernel image:

- **Compile BPF program from C source**
- **Attach BPF program to kprobe/uprobe/tracepoint/USDT/socket**
- **Poll data from BPF program**
- **Framework for building new tools or one-off scripts**
- **...**

BCC tools

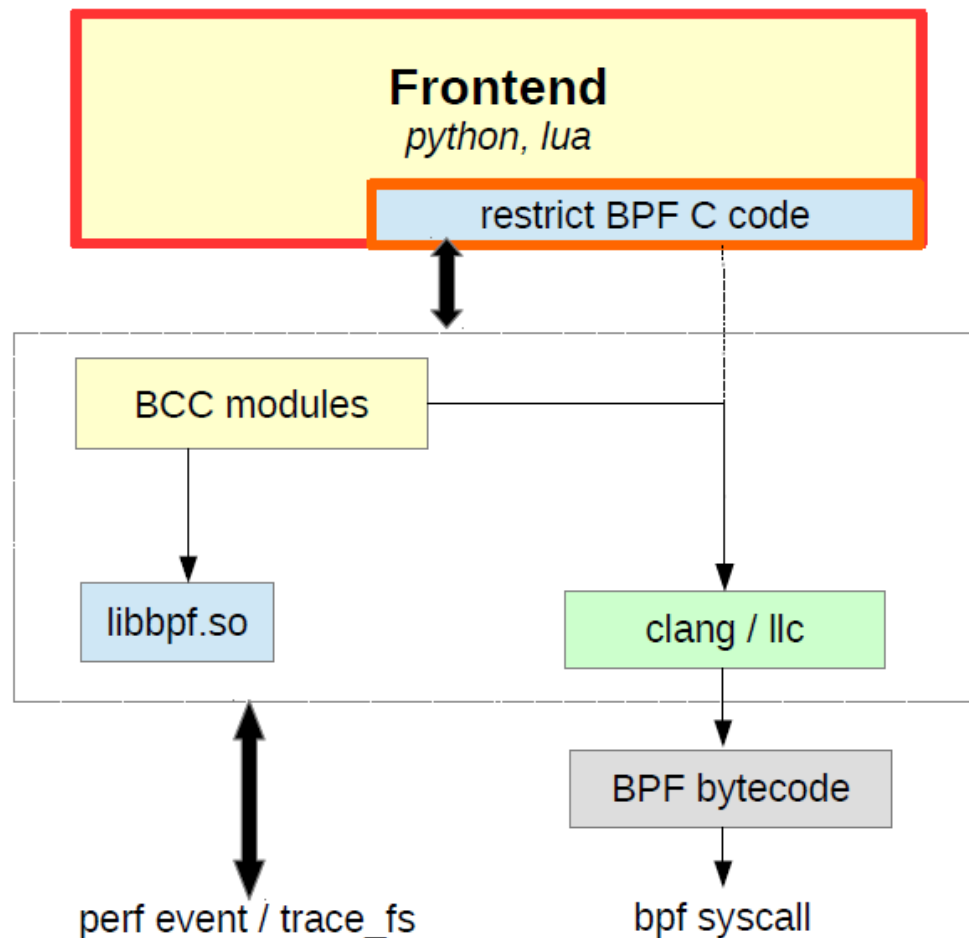


<https://github.com/iovisor/bcc#tools> 2016

Python is given the fullest support

Arch

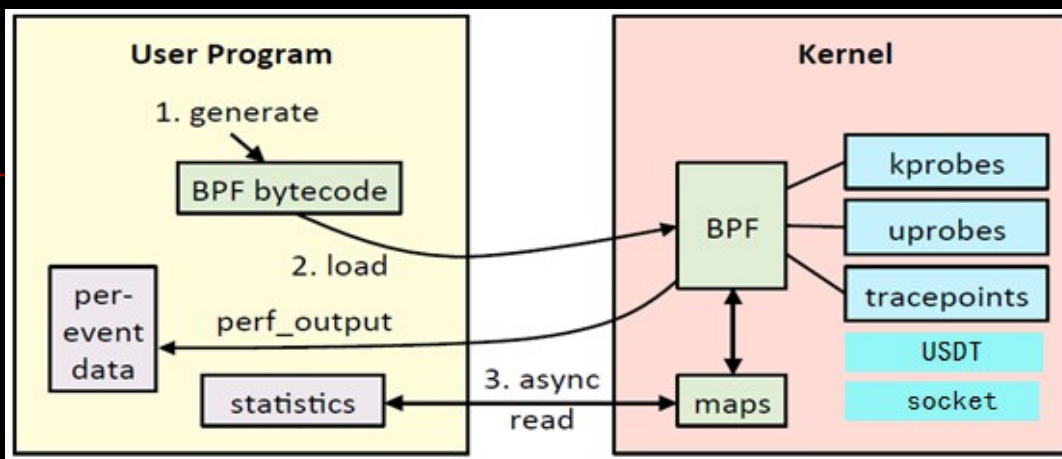
iovisor
BCC



27/05/2016

Source: <http://www.slideshare.net/vh21/meet-cutebetweenebpfandtracing>

For Tracing



Source: <http://www.slideshare.net/brendangregg/linux-bpf-superpowers>

Sample

bcc/examples/tracing/urandomread*.*

```
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/eBPF/BCC/bcc/examples/tracing# ./urandomread.py
TIME(s)      COMM      PID  GOTBITS
3031.665037000 dd        6604  8192
3031.665365000 dd        6604  8192
3031.665642000 dd        6604  8192
3031.665924000 dd        6604  8192
3031.666202000 dd        6604  8192
3095.286445000 systemd   1      128
3095.286518000 systemd   1      128
3095.286582000 systemd   1      128
3095.286671000 systemd   1      128
```

```
mydev@ubuntu:/opt/Tmp$ dd if=/dev/urandom of=/dev/null bs=1k count=5
5+0 records in
5+0 records out
5120 bytes (5.1 kB, 5.0_KiB) copied, 0.00182226 s, 2.8 MB/s
```

```

urandomread-explicit.py x
1  #!/usr/bin/python
2  #
3  # urandomread-explicit  Example of instrumenting a kernel tracepoint.
4  # For Linux, uses BCC, BPF. Embedded C.
5  #
6  # This is an older example of instrumenting a tracepoint, which defines
7  # the argument struct and makes an explicit call to attach_tracepoint().
8  # See urandomread for a newer version that uses TRACEPOINT_PROBE().
9  #
10 # REQUIRES: Linux 4.7+ (BPF_PROG_TYPE_TRACEPOINT support).
11 #
12 # Test by running this, then in another shell, run:
13 #   dd if=/dev/urandom of=/dev/null bs=1k count=5
14 #
15 # Copyright 2016 Netflix, Inc.
16 # Licensed under the Apache License, Version 2.0 (the "License")
17
18 from __future__ import print_function
19 from bcc import BPF
20
21 # define BPF program
22 bpf_text = """
23 #include <uapi/linux/ptrace.h>
24
25 struct urandom_read_args {
26     // from /sys/kernel/debug/tracing/events/random/urandom_read/format
27     u64 __unused__;
28     u32 got_bits;
29     u32 pool_left;
30     u32 input_left;
31 };
32
33 int printarg(struct urandom_read_args *args) {
34     bpf_trace_printk("%d\\n", args->got_bits);
35     return 0;
36 };
37 """
38
39 # load BPF program
40 b = BPF(text=bpf_text)
41 b.attach_tracepoint("random:urandom_read", "printarg")
42
43 # header
44 print("%-18s %-16s %-6s %s" % ("TIME(s)", "COMM", "PID", "GOTBITS"))
45
46 # format output
47 while 1:
48     try:
49         (task, pid, cpu, flags, ts, msg) = b.trace_fields()
50     except ValueError:
51         continue
52     print("%-18.9f %-16s %-6d %s" % (ts, task, pid, msg))

```

include/trace/events/random.h

```

290 TRACE_EVENT(urandom_read,
291   TP_PROTO(int got_bits, int pool_left, int input_left),
292   TP_ARGS(got_bits, pool_left, input_left),
293   TP_STRUCT__entry(
294     __field(int, got_bits)
295     __field(int, pool_left)
296     __field(int, input_left)
297   ),
298   TP_fast_assign(
299     __entry->got_bits = got_bits;
300     __entry->pool_left = pool_left;
301     __entry->input_left = input_left;
302   ),
303   TP_printk("got_bits %d nonblocking_pool_entropy_left %d "
304     "ringbit_entropy_left %d", __entry->got_bits,
305     __entry->pool_left, __entry->input_left)
306 );

```

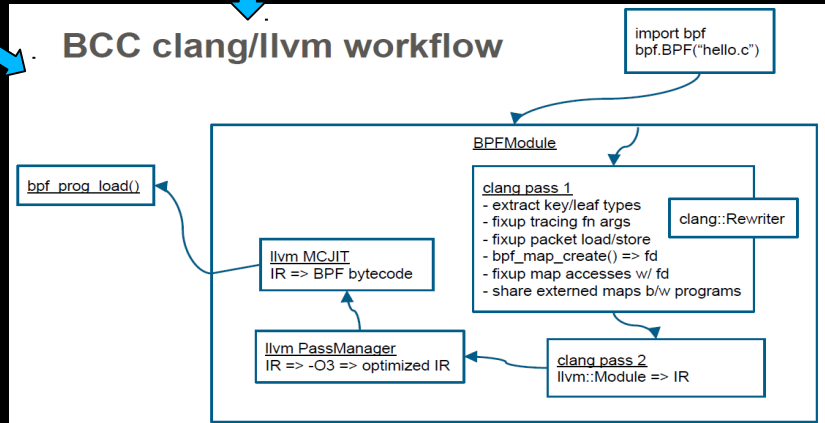
/sys/kernel/debug/tracing/trace_pipe

```

urandomread.py x
1  #!/usr/bin/python
2  #
3  # urandomread  Example of instrumenting a kernel tracepoint.
4  # For Linux, uses BCC, BPF. Embedded C.
5  #
6  # REQUIRES: Linux 4.7+ (BPF_PROG_TYPE_TRACEPOINT support).
7  #
8  # Test by running this, then in another shell, run:
9  #   dd if=/dev/urandom of=/dev/null bs=1k count=5
10 #
11 # Copyright 2016 Netflix, Inc.
12 # Licensed under the Apache License, Version 2.0 (the "License")
13
14 from __future__ import print_function
15 from bcc import BPF
16
17 # load BPF program
18 b = BPF(text="""
19 TRACEPOINT_PROBE(random, urandom_read) {
20     // args is from /sys/kernel/debug/tracing/events/random/urandom_read/format
21     bpf_trace_printk("%d\\n", args->got_bits);
22     return 0;
23 };
24 """)
25
26 # header
27 print("%-18s %-16s %-6s %s" % ("TIME(s)", "COMM", "PID", "GOTBITS"))
28
29 # format output
30 while 1:
31     try:
32         (task, pid, cpu, flags, ts, msg) = b.trace_fields()
33     except ValueError:
34         continue
35     print("%-18.9f %-16s %-6d %s" % (ts, task, pid, msg))

```

BCC clang/llvm workflow

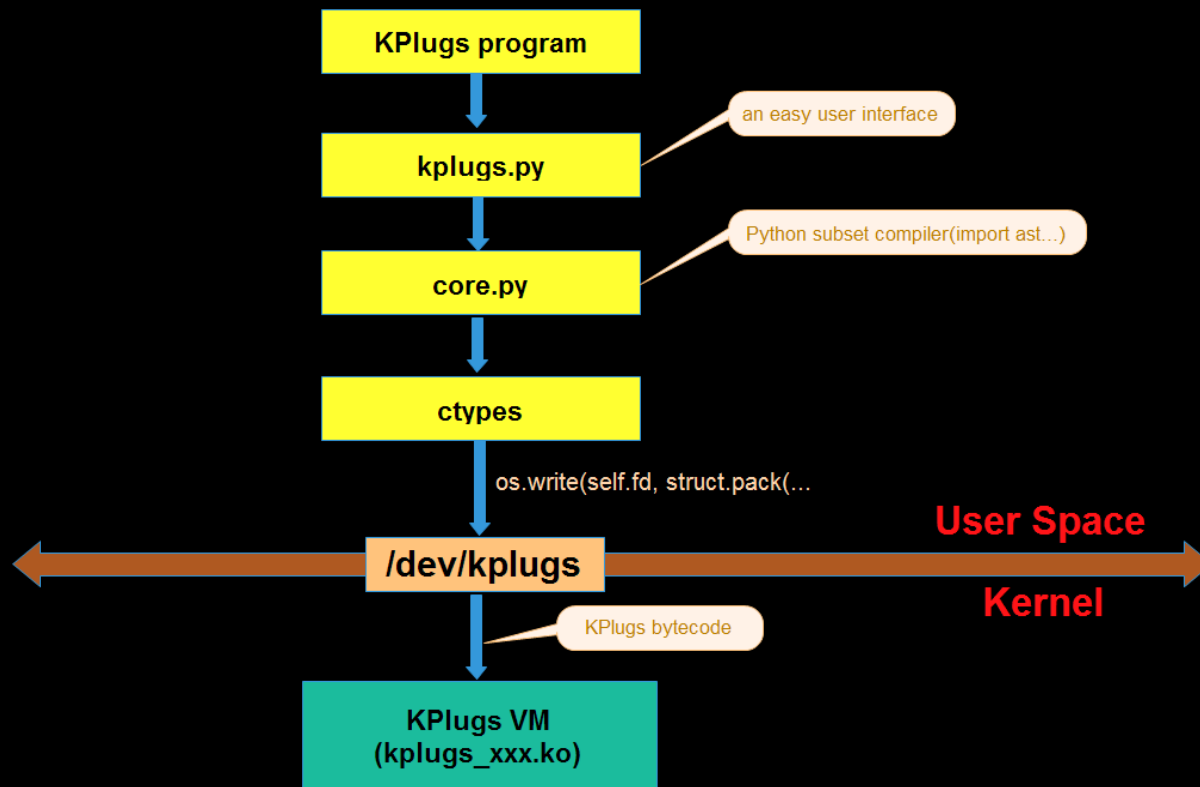


Source: http://linuxplumbersconf.org/2015/ocw//system/presentations/3249/original/bpf_llvm_2015aug19.pdf

II. In-Kernel Virtual Machine

1) KPlugs

- <https://www.kplugins.org>
- <https://github.com/avielw/kplugins>
- Arch



Hook sample

- kplugins.Mem - access all of the computer's memory of the computer (kernel and all processes' user space)
- kplugins.Symbol - resolve kernel symbols
- kplugins.Hook - hook kernel functions with a KPlugins function
- kplugins.Caller - call an exported kernel function

```
#!/usr/bin/env python
```

```
import kplugins
import time

try:
    plug = kplugins.Plug()
    hook = kplugins.Hook()

    kernel_func = r'''

def my_hook(kp, regs):
    print "The registers are stored in %p" % regs
    return 0
'''

    my_hook = plug.compile(kernel_func)[0]
    hook.hook("sys_clone", my_hook)
    time.sleep(10)
    hook.unhook(my_hook)

finally:
    kplugins.release_kplugins()
```

```
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples# lsmod |grep kplugins
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples# modprobe kplugins_release
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples# lsmod |grep kplugins
kplugins_release      53248      0
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples# dmesg -C
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples# dmesg
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples# ./kplugins_hook.py
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples# dmesg
[ 6316.625909] The registers are stored in ffff8800978ffe50
[ 6317.644612] The registers are stored in ffff8800928a3e50
[ 6318.675235] The registers are stored in ffff8800928a3e50
[ 6319.694774] The registers are stored in ffff8800928a3e50
[ 6320.744613] The registers are stored in ffff8800928a3e50
[ 6321.763556] The registers are stored in ffff8800928a3e50
root@ubuntu:/opt/MyWorkSpace/MyProjs/Open-Source/OS/In-Kernel-VM/Python/kplugins/samples#
```

```
mydev@ubuntu:/opt/MyWorkSpace/Test/Linux/Thread$ ./multithread
```

```
struct kprobe {
    struct hlist_node hlist;
    struct list_head list;
    unsigned long nmisses;
    kprobe_opcode_t *addr;
    const char *symbol_name;
    unsigned int offset;
    kprobe_pre_handler_t pre_handler;
    kprobe_post_handler_t post_handler;
    kprobe_fault_handler_t fault_handler;
    kprobe_break_handler_t break_handler;
    kprobe_opcode_t opcode;
    struct arch_specific_insn ainsn;
    u32 flags;
}
```

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>

#define THREAD_NUM 5

void *thread_func(void *argu) {
    printf("Thread %d starting...\n", (int)argu);

    return NULL;
}

int main() {
    int i;
    pthread_t thread_ids[THREAD_NUM];

    for (i = 0; i < THREAD_NUM; i++) {
        sleep(1);
        pthread_create(&thread_ids[i], NULL, &thread_func, (void *)i);
    }

    for (i = 0; i < THREAD_NUM; i++) {
        pthread_join(thread_ids[i], NULL);
    }
    printf("--end of The program...\n");

    return 0;
}
```

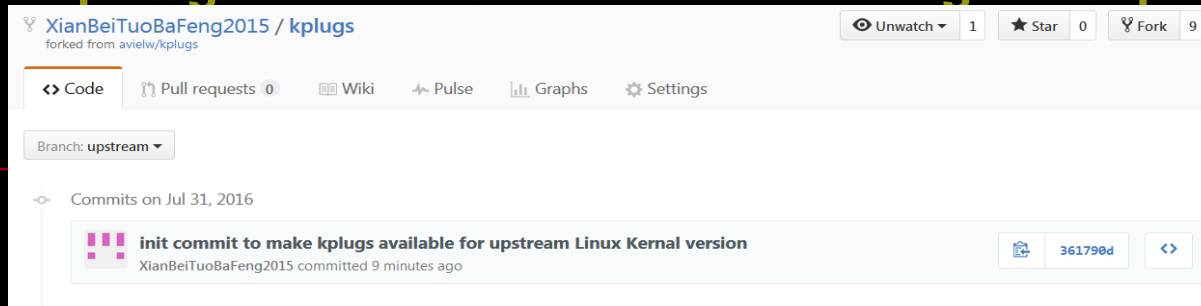
```
def hook(self, where, func):
    if self._hooks.has_key(func.addr):
        raise Exception("This function is already a callback of this class")

    # create a kprobe struct
    kp = self._mem.alloc(KPROBE_STRUCT_MAXSIZE)
    if isinstance(where, str):
        sym = self._mem.alloc(len(where) + 1)
        self._mem[sym] = where + '\0'
        self._mem[kp + KPROBE_STRUCT_SYMBOL : kp + KPROBE_STRUCT_SYMBOL + WORD_SIZE] = sym
    else:
        self._mem[kp + KPROBE_STRUCT_ADDR : kp + KPROBE_STRUCT_ADDR + WORD_SIZE] = where
        self._mem[kp + KPROBE_STRUCT_HANDLER : kp + KPROBE_STRUCT_HANDLER + WORD_SIZE] = func.addr

    # register the kprobe hook
    err = self._caller["register_kprobe"](kp)
    if err:
        raise Exception("register_kprobe failed")
    self._hooks[func.addr] = kp
```

My Fork


<https://github.com/XianBeiTuoBaFeng2015/kplugins>



Debugging

1) `python -m trace --trace ./kplugins_hello.py`

2) VisualKernel (Debug Linux Kernel with Visual Studio)



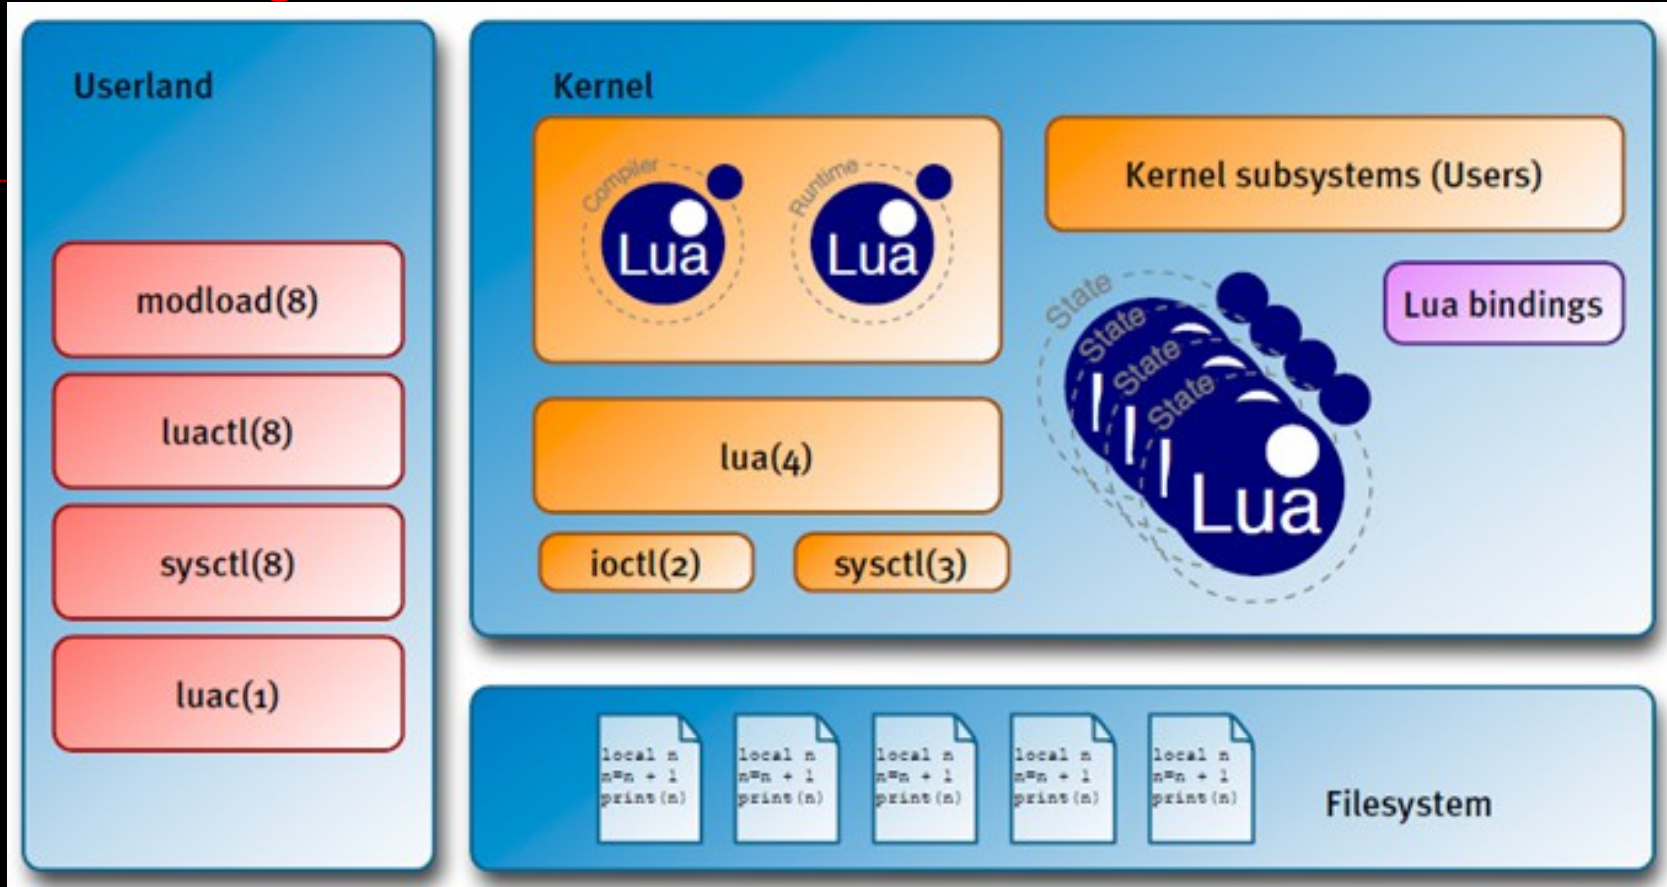
2) Application

- <http://www.netbsd.org>
 - <http://www.lua.org>
 - **NetBSD Kernel scripting with Lua**
-



- be part of NetBSD 6 (Userland)
- be part of NetBSD 7 (Kernel)

The Big Picture



Source: https://archive.fosdem.org/2013/schedule/event/lua_in_the_netbsd_kernel/

■ still lack of JIT(Just-in-Time) support, but available on Linux


■ Why not Python?

- Huge library
- Memory consumption
- Difficult object mapping

VM size at megabytes level, but the key idea behind KPlugs is...

Conclusions

Reduce the size of Kernel source code



Source: https://www.openhub.net/p/linux/analyses/latest/languages_summary

Deliver a higher-level programming environment to the Kernel

Let users explore the system in an easy way

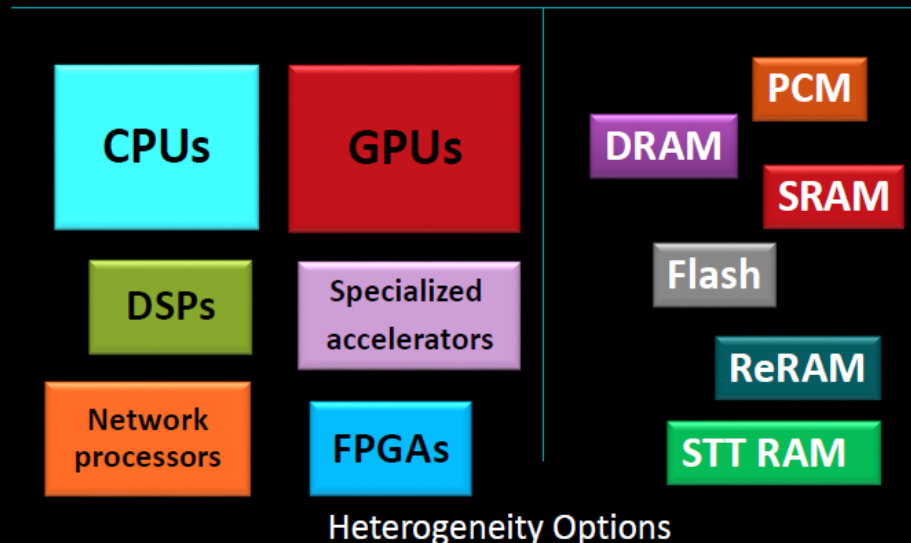
Great innovation in OS development!

III. Python-assisted Data Center Dev

1) HPC

- High Performance Computing or Heterogeneous Parallel Computing

C/C++	FORTTRAN	Java
UPC/UPC++	python	MPI
Kokkos/RAJA	OpenMP	OpenACC



Anaconda

- <https://www.continuum.io/>
- Open Source Modern Analytics Platform Powered by Python
- Swissknife for Big Data, AI, HPC, Cloud/Web, Exploration & Viz

APP	Notebooks		Embeddable Dashboards		Data Services		Visual Apps		
VIZ	Plots	Interactive Viz		Big Data	Maps & GIS	3D	Streaming	Graphs	
STORYBOARD	Notebooks		Interactive Exploration		Visual Programming		Data IDEs		
ANALYTICS	Data Prep		Stats	ML & Ensembles		Deep Learning		Simulation & Optimization	
	Geospatial		Text & NLP		Video/Image/Audio Mining		Graph & Network		
DATA	Hadoop & Hive		Spark	NoSQL	DW & SQL		Files & Web Services		
HW	Servers		Clusters	GPUs & High End Workstations					

Intel

■ Intel Distribution for Python

■ <https://software.intel.com/en-us/python-distribution>

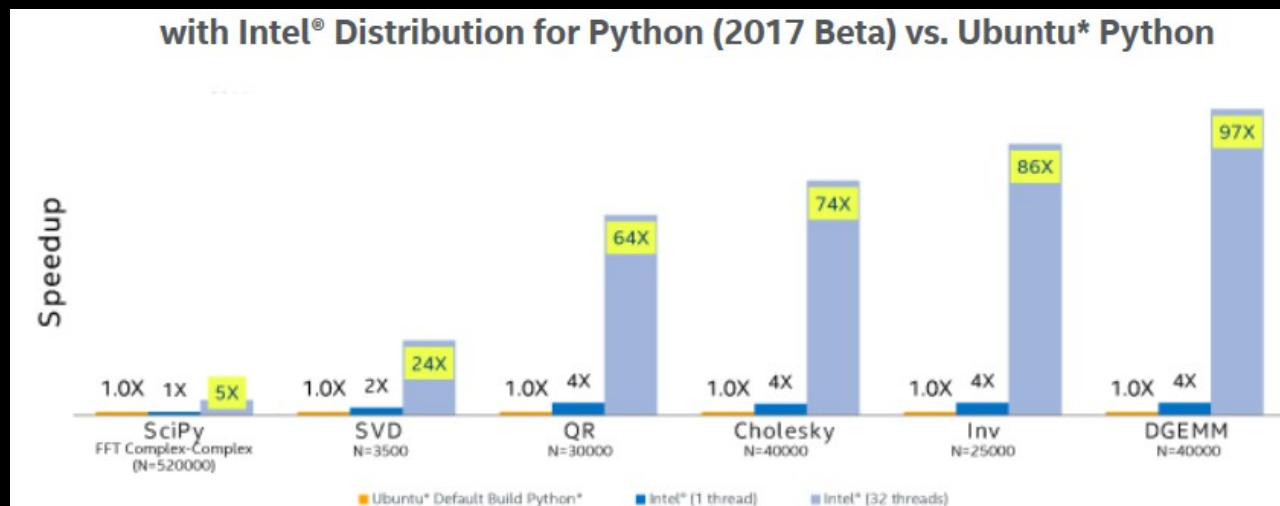
■ accelerated with MKL, MPI, TBB, DAAL

■ support of Xeon Phi

■ packages:

- Included Python* packages: NumPy*, SciPy*, Pandas*, Matplotlib*, IPython*, SymPy*, NumExpr*, Scikit-learn* (Linux*, Microsoft Windows*, and OS X* operating systems), mpi4py* (Linux*), DistArray* (Linux*), PyTables (Linux* and Windows*), Numba* (Linux*, Windows*, and OS X*), Conda* package management tool
- pyDAAL: Python package for the [Intel Data Analytics Acceleration Library](#)

■ Performance Boost on Select Numerical Functions

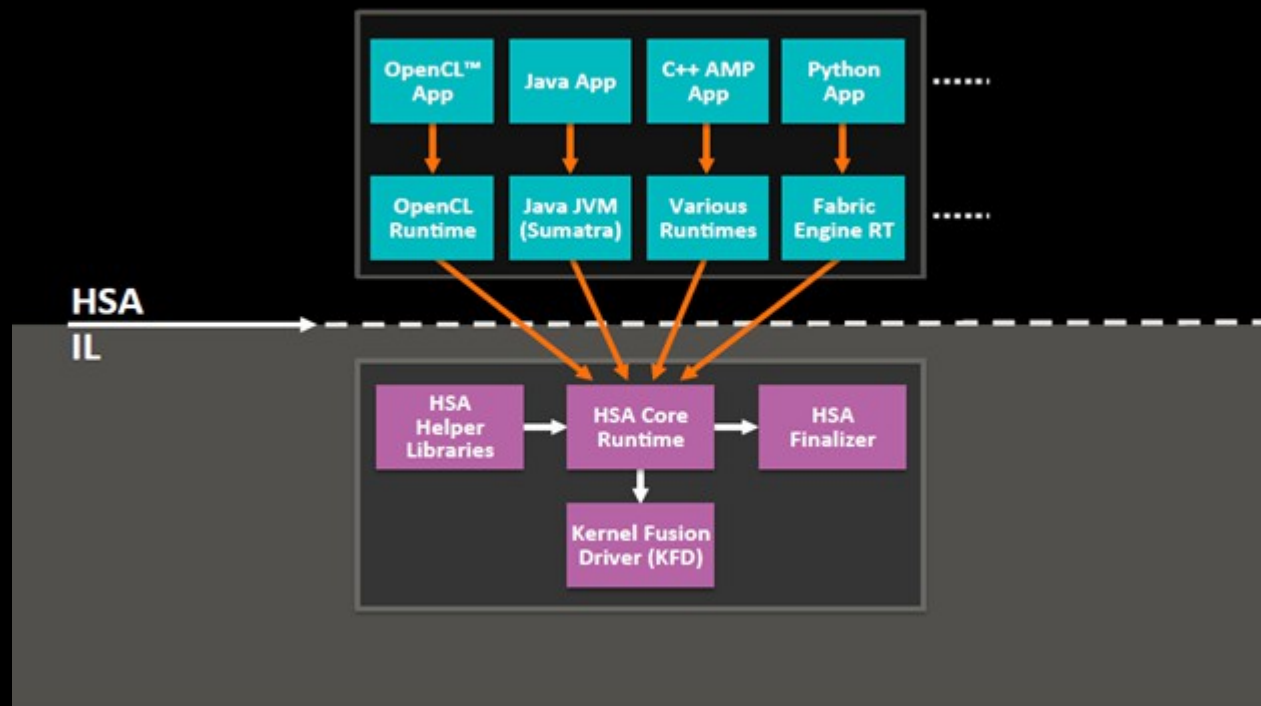


AMD

■ Heterogeneous System Architecture

■ <http://www.hsafoundation.com/>

PROGRAMMING LANGUAGES PROLIFERATING ON HSA



■ <https://github.com/ContinuumIO/Numba-HSA-Webinar/>

2) Software-Defined Everything (SDE)

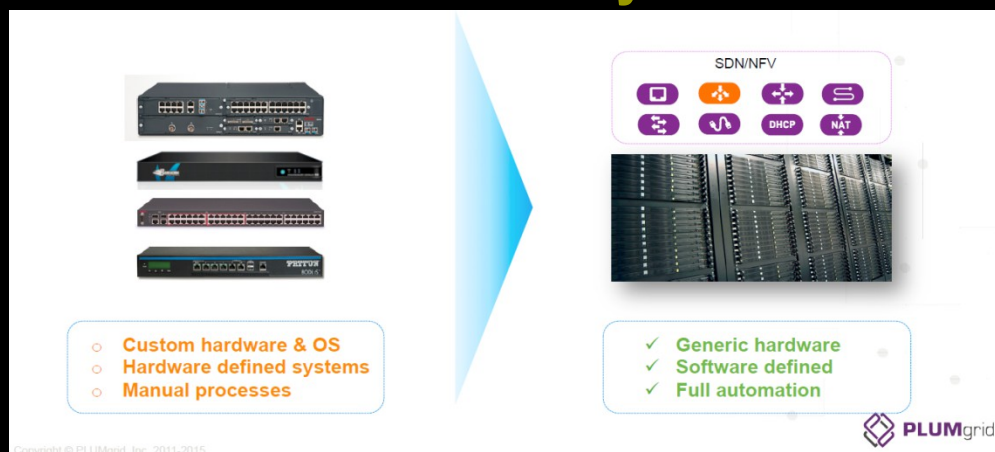
- Software Defined Data Center (SDDC)
- Software-Defined Networking (SDN)
- Software-Defined Storage (SDS)

- https://en.wikipedia.org/wiki/Software-defined_data_center

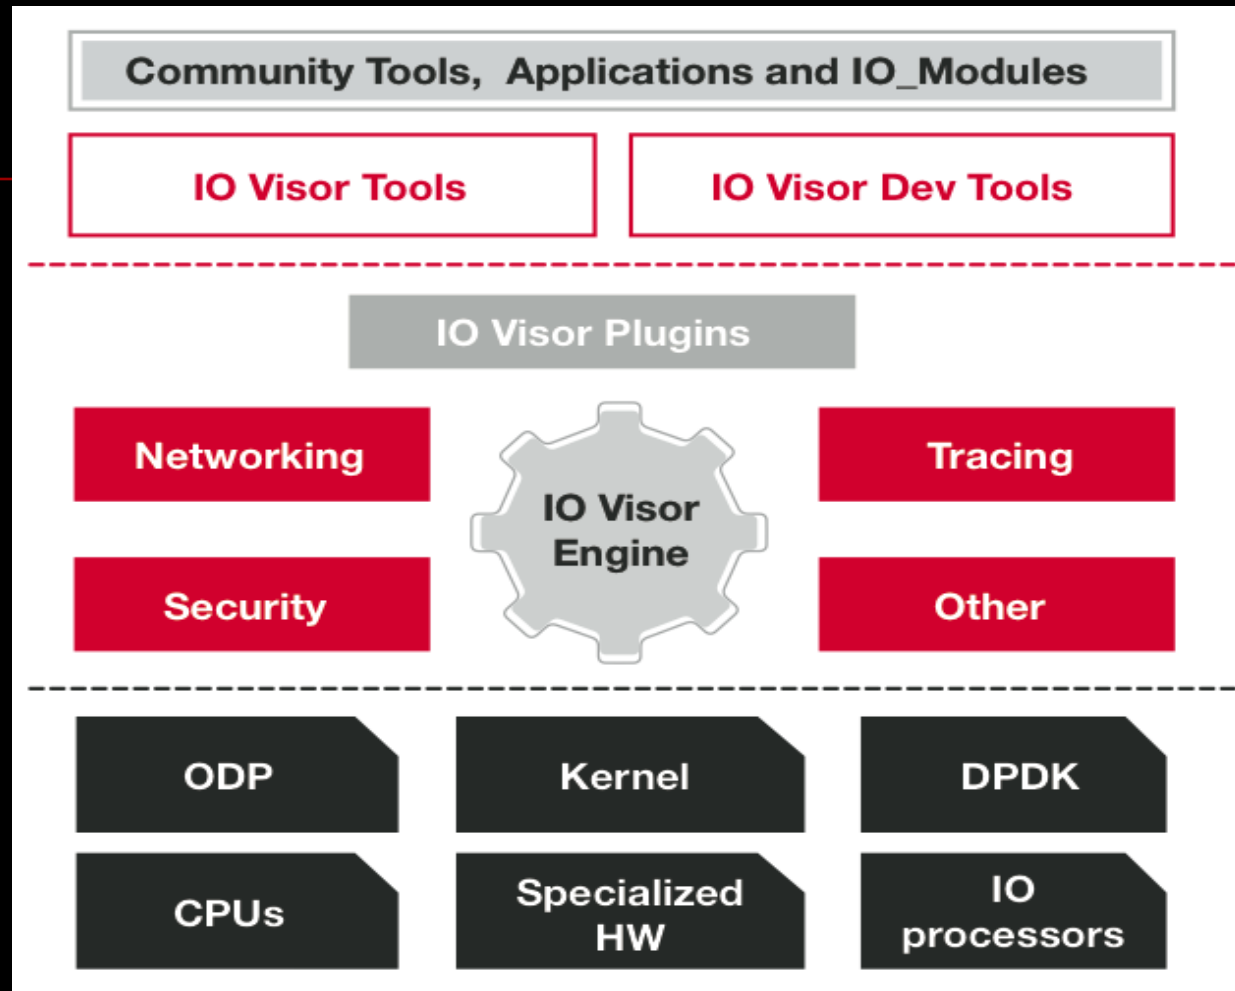
The software-defined data center encompasses a variety of concepts and data center infrastructure components, and each component can be provisioned, operated, and managed through an **application programming interface (API)**.^[4] The core architectural components that comprise the software-defined data center^[5] include the following:

- Computer **virtualization**,^[6] which is a software implementation of a computer.
- **Software-defined networking** (SDN), which includes **network virtualization**, is the process of merging hardware and software resources and networking functionality into a software-based **virtual network**.^[5]
- **Software-defined storage** (SDS), which includes **storage virtualization**, suggests a service interface to provision capacity and SLAs (Service Level Agreements) for storage, including performance and durability.
- Management and automation software, enabling an administrator to provision, control, and manage all software-defined data center components.^[7]

■ Traditional to Software Defined Systems



- <https://www.iovisor.org/>



<https://github.com/iovisor/bcc/tree/master/src/cc/frontends/p4/>

IV. Summary

- User space/Kernel space Repartition & Unifying
 - write your function in user space, while run it in kernel
 - user space drivers
 - scripting your OS
- Compilation technologies are rapidly evolving
 - compiler infrastructure like LLVM
 - source-to-source, bytecode-to-bytecode, and bytecode to native code compiler
- Python-based Domain Specific Languages & language subset
- Python is sure to play an important role in next generation Software Defined Systems!

Q & A

Thanks!



Reference

Slides/materials from many and varied sources:

- <http://en.wikipedia.org/wiki/>
- <http://www.slideshare.net/>

- <https://www.kernel.org/doc/Documentation/>
- <http://tracingsummit.org>
- <http://www.brendangregg.com/blog/>
- <https://www.python.org>
- <http://llvm.org>
- https://en.wikipedia.org/wiki/Just-in-time_compilation
- <http://sysprogs.com/VisualKernel/>
- <https://www.netbsd.org/gallery/presentations/>
- [https://en.wikipedia.org/wiki/Anaconda_\(Python_distribution\)](https://en.wikipedia.org/wiki/Anaconda_(Python_distribution))
- <https://www.opennetworking.org/>
- <https://www.opnfv.org/>
- <http://pypy.org>
- https://en.wikipedia.org/wiki/Runtime_system
- https://en.wikipedia.org/wiki/Intermediate_representation
- ...