



Python Profiling and Performance Tuning

Python at Pinterest

Joe Gordon

August 14st, 2016

About Me

- **Site Reliability Engineer (SRE) @ Pinterest**
 - Availability
 - Performance
 - Efficiency
 - Developer Velocity
- **SRE embedded in Search Team**



Problem

Gevent

Profiling

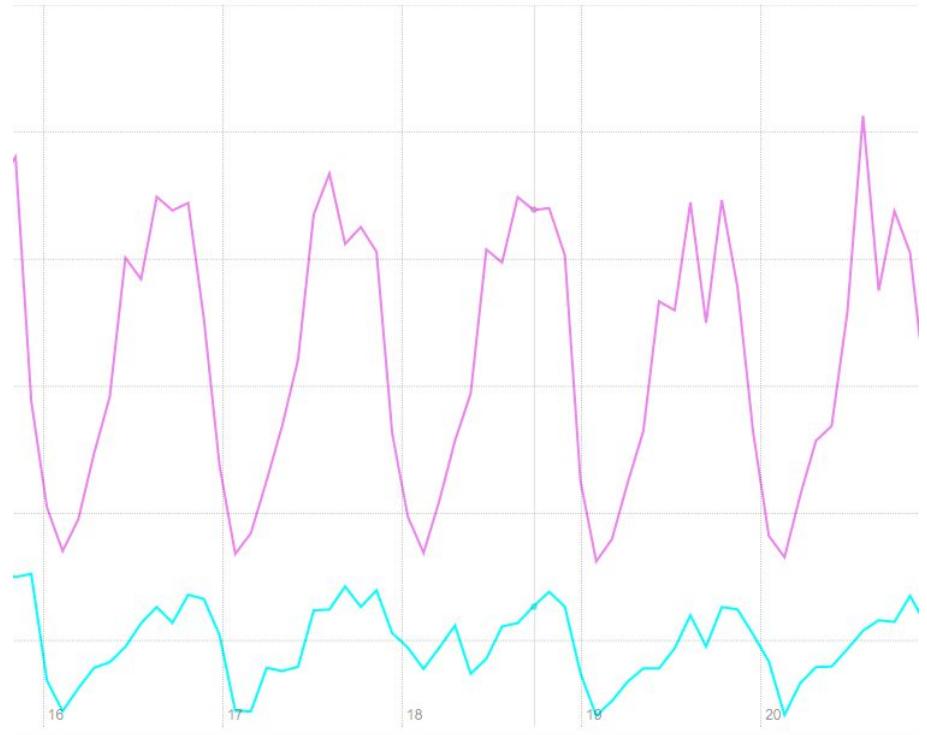
Impact



Pin Search P90 Latency

Q4 2015

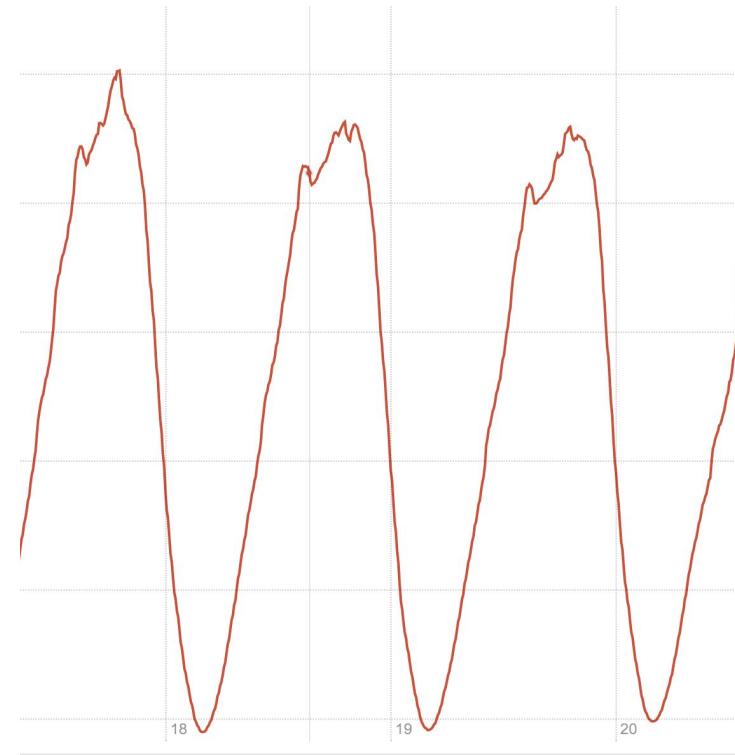
- **Latency**
 - Varnish: 2200ms
 - Search system: 640ms
- **70% of latency outside of Search system**
- **Height of Varnish latency >> height of asterix latency**



GET API Latency

Q4 2015

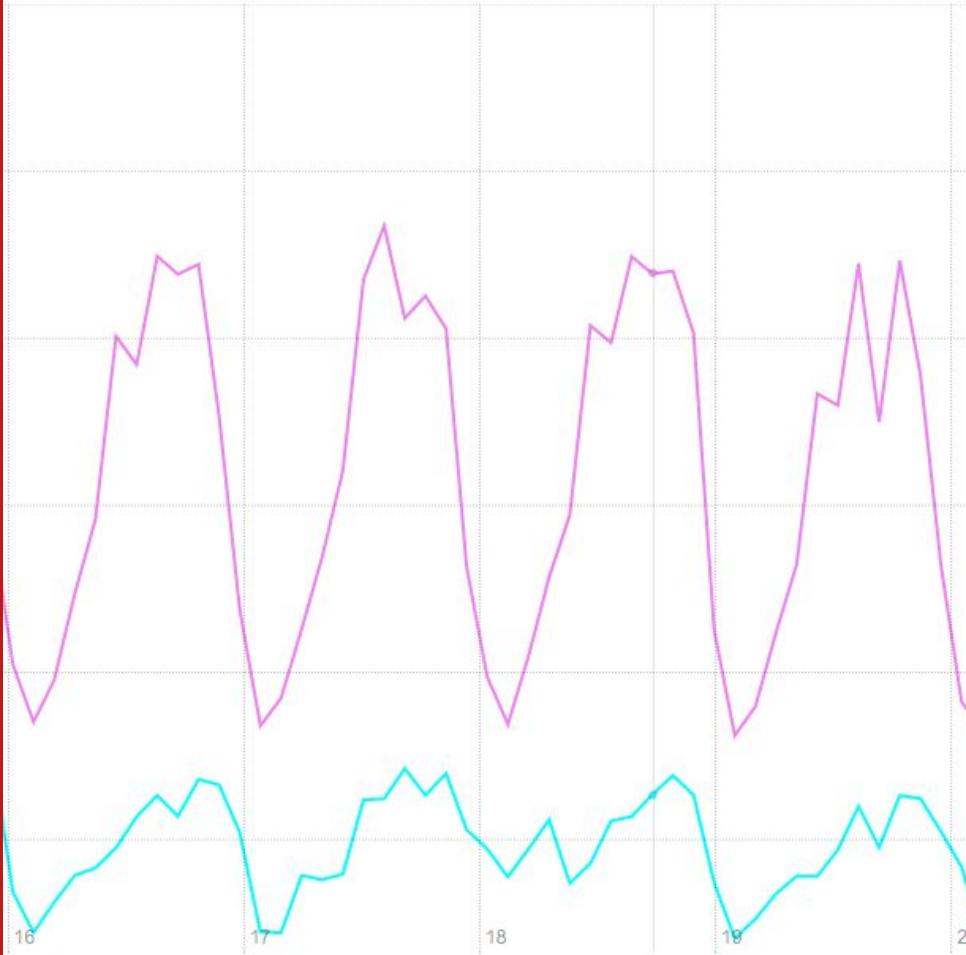
- Same pattern
- Height: 50%



Why?

- *Latency from frontend*
- *Height of latency*

Pins P90





Problem



Gevent

Profiling

Impact

What is gevent?

gevent is a [coroutine](#)-based Python networking library that uses [greenlet](#) to provide a high-level synchronous API on top of the libev event loop.

Features include:

- **Fast event loop** based on libev (epoll on Linux, kqueue on FreeBSD).
- **Lightweight execution units** based on greenlet.

What is gevent?

gevent is a `coroutine`-based Python networking library that uses `greenlet` to provide a high-level synchronous API on top of the libev event loop.

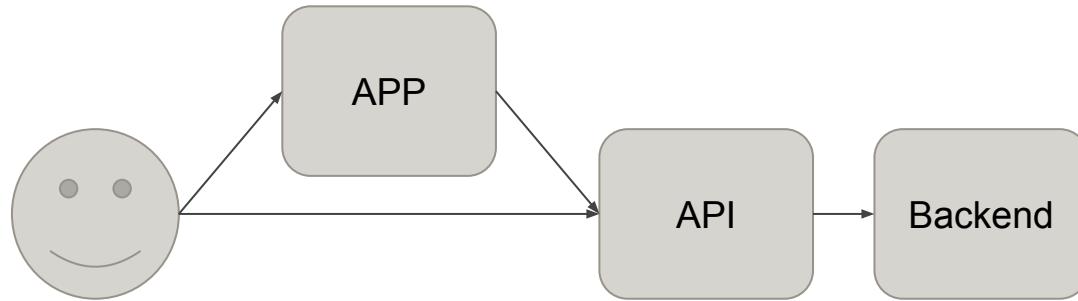
Example

The following example shows how to run tasks concurrently.

```
>>> import gevent
>>> from gevent import socket
>>> urls = ['www.google.com', 'www.example.com', 'www.python.org']
>>> jobs = [gevent.spawn(socket.gethostbyname, url) for url in urls]
>>> gevent.joinall(jobs, timeout=2)
>>> [job.value for job in jobs]
['74.125.79.106', '208.77.188.166', '82.94.164.162']
```

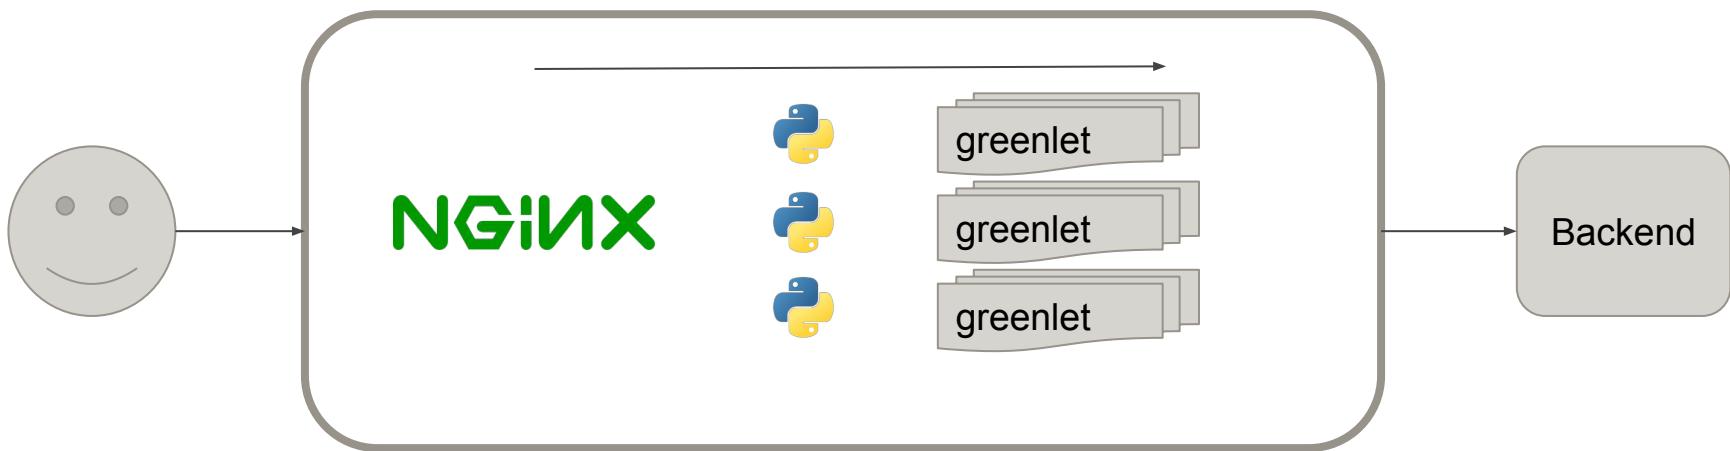
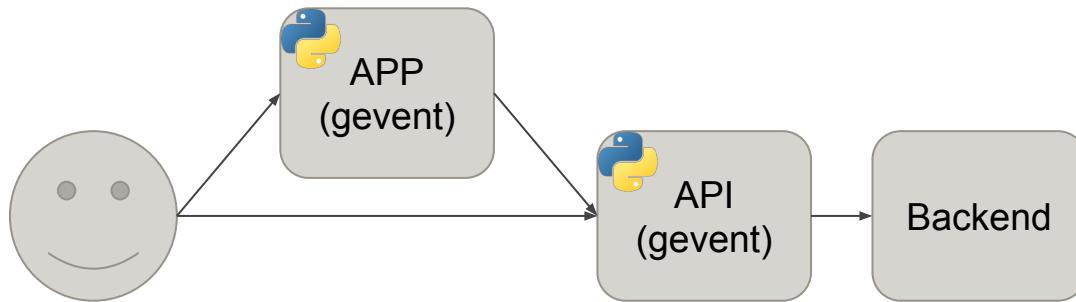
Software Stack

gevent



Software Stack

gevent



gevent

Cooperative multitasking

Cooperative multitasking

The greenlets all run in the same OS thread and are scheduled cooperatively. This means that until a particular greenlet gives up control, (by calling a blocking function that will switch to the Hub), other greenlets won't get a chance to run. This is typically not an issue for an I/O bound app, but one should be aware of this when doing something CPU intensive, or when calling blocking I/O functions that bypass the libev event loop.

gevent

Cooperative multitasking

Cooperative multitasking

The greenlets all run in the same OS thread and are scheduled cooperatively. This means that until a particular greenlet gives up control, (by calling a blocking function that will switch to the Hub), other greenlets won't get a chance to run. This is typically not an issue for an I/O bound app, but one should be aware of this when doing something CPU intensive, or when calling blocking I/O functions that bypass the libev event loop.

What happens if we mix CPU intensive and network bound workloads together?

Gevent Experiment

setup

Server:

```
^+ 62 def workload(request_id):
63     gevent_yield(request_id)
64     block(request_id) ←
65     network_call(request_id)
66     log.debug("%s: done" % request_id)
67
68
69 def main():
70     print('Serving on 8088...')
71     WSGIServer((), 8088, application).serve_forever()
72
```

Concurrent Client:

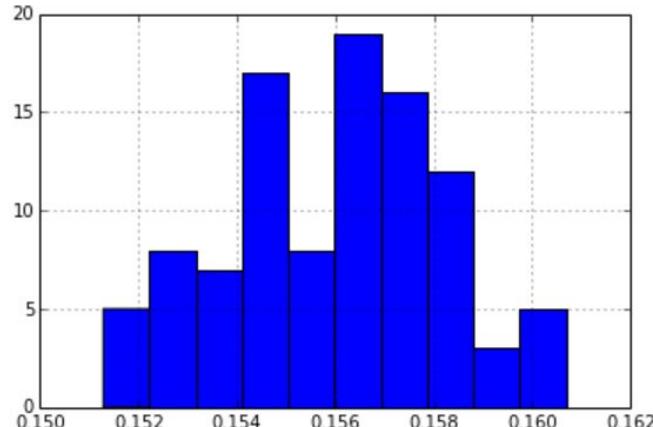
```
9  def main.concurrent, count):
10    session = grequests.Session()
11    urls = ['http://localhost:8088']*count
12    rs = (grequests.get(u, session=session) for u in urls)
13    responses = grequests.map(rs, size=concurrent)
14    for response in responses:
15        response.close() ←
16    print "%s calls, %s at a time" % (count, concurrent)
17    r = requests.get('http://localhost:8088/stats')
18    return r.json()
19
```

blocking time: 50ms
non blocking time: 100ms

Gevent Experiment

results

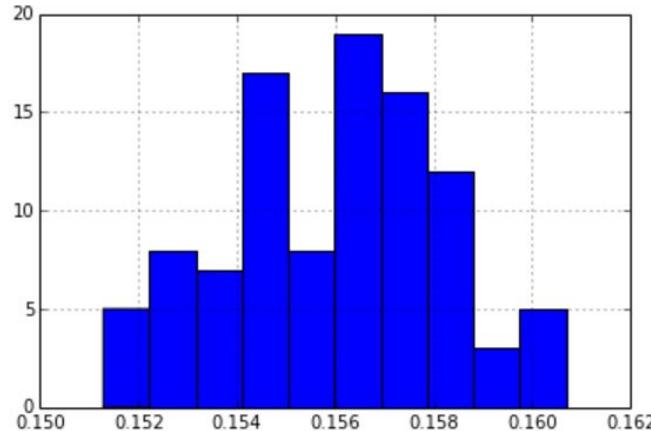
```
100 calls, 1 at a time
count    100.000000
mean     0.155956
std      0.002204
min     0.151283
25%     0.154440
50%     0.156152
75%     0.157455
max     0.160684
dtype: float64
```



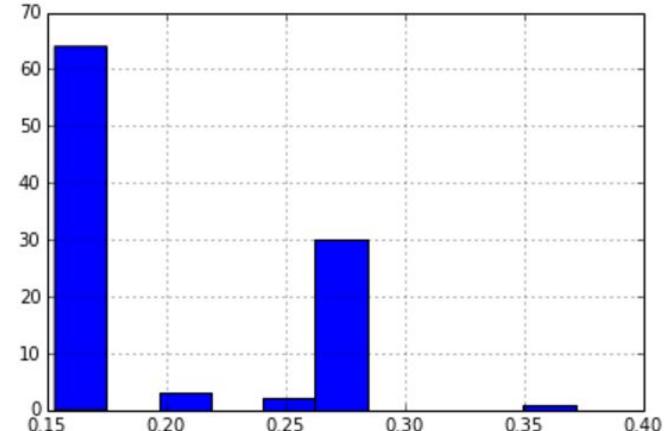
Gevent Experiment

results

```
100 calls, 1 at a time
count      100.000000
mean       0.155956
std        0.002204
min        0.151283
25%        0.154440
50%        0.156152
75%        0.157455
max        0.160684
dtype: float64
```



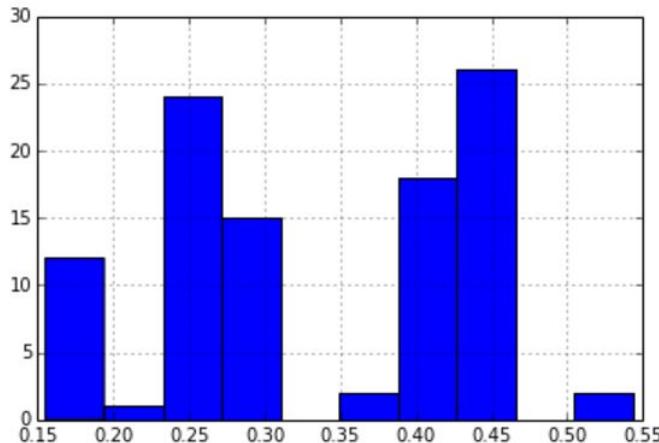
```
100 calls, 5 at a time
count      100.000000
mean       0.199159
std        0.052894
min        0.153300
25%        0.161063
50%        0.164095
75%        0.265418
max        0.371659
dtype: float64
```



Gevent Experiment

results

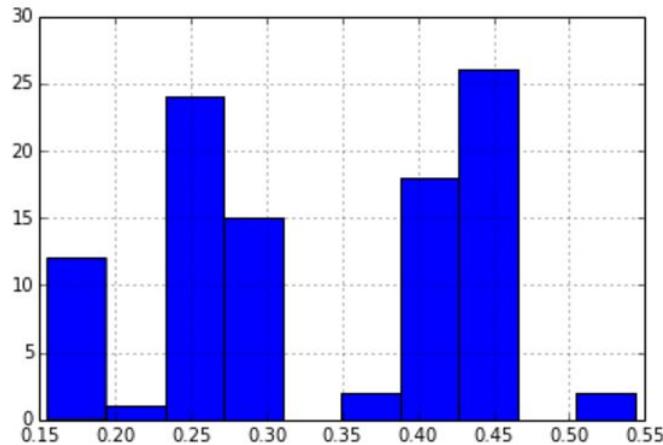
```
100 calls, 10 at a time
count    100.000000
mean     0.333240
std      0.102039
min     0.154991
25%     0.265487
50%     0.274264
75%     0.427456
max     0.543295
dtype: float64
```



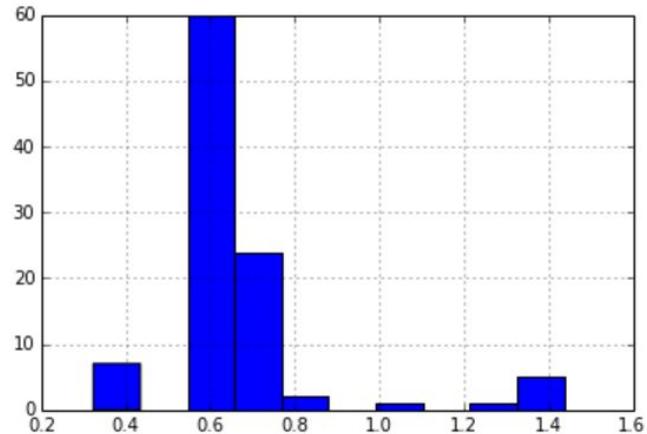
Gevent Experiment

results

```
100 calls, 10 at a time
count    100.000000
mean     0.333240
std      0.102039
min     0.154991
25%     0.265487
50%     0.274264
75%     0.427456
max     0.543295
dtype: float64
```



```
100 calls, 15 at a time
count    100.000000
mean     0.687463
std      0.194353
min     0.323455
25%     0.635788
50%     0.642643
75%     0.693844
max     1.437414
dtype: float64
```



Gevent Experiment: Server Log

5 concurrent requests

```
2016-02-19 15:50:40,334 - 92: enter gevent_yield
2016-02-19 15:50:40,334 - 91: exit gevent_yield
2016-02-19 15:50:40,334 - 91: enter block
2016-02-19 15:50:40,389 - 91: exit block
2016-02-19 15:50:40,389 - 91: enter network_call
2016-02-19 15:50:40,389 - 92: exit gevent_yield
2016-02-19 15:50:40,389 - 92: enter block
2016-02-19 15:50:40,443 - 92: exit block
2016-02-19 15:50:40,444 - 92: enter network_call
2016-02-19 15:50:40,444 - 89: exit network_call
2016-02-19 15:50:40,444 - 89: done
127.0.0.1 -- [2016-02-19 15:50:40] "GET / HTTP/1.1" 200
119 0.275503

2016-02-19 15:50:40,444 - 90: exit network_call
2016-02-19 15:50:40,444 - 90: done
127.0.0.1 -- [2016-02-19 15:50:40] "GET / HTTP/1.1" 200
119 0.165371

2016-02-19 15:50:40,445 - 93: enter gevent_yield
2016-02-19 15:50:40,445 - 93: exit gevent_yield
2016-02-19 15:50:40,445 - 93: enter block
2016-02-19 15:50:40,496 - 93: exit block
2016-02-19 15:50:40,496 - 93: enter network_call
2016-02-19 15:50:40,496 - 91: exit network_call
2016-02-19 15:50:40,496 - 91: done
127.0.0.1 -- [2016-02-19 15:50:40] "GET / HTTP/1.1" 200
119 0.163295
```

```
2016-02-19 15:50:40,497 - 94: enter gevent_yield
2016-02-19 15:50:40,497 - 95: enter gevent_yield
2016-02-19 15:50:40,497 - 94: exit gevent_yield
2016-02-19 15:50:40,497 - 94: enter block
2016-02-19 15:50:40,549 - 94: exit block
2016-02-19 15:50:40,549 - 94: enter network_call
2016-02-19 15:50:40,550 - 95: exit gevent_yield
2016-02-19 15:50:40,550 - 95: enter block
2016-02-19 15:50:40,603 - 95: exit block
2016-02-19 15:50:40,603 - 95: enter network_call
2016-02-19 15:50:40,603 - 92: exit network_call
2016-02-19 15:50:40,603 - 92: done
```

- Extra 50ms block in first gevent_yield
- network_call takes 50 extra ms (159 total)
- 100ms extra from blocks from other requests
- ~250ms (269ms) instead of expected 150ms

Gevent Experiment: Server Log

5 concurrent requests

```
2016-02-19 15:50:40,334 - 92: enter gevent_yield
2016-02-19 15:50:40,334 - 91: exit gevent_yield
2016-02-19 15:50:40,334 - 91: enter block
2016-02-19 15:50:40,389 - 91: exit block
2016-02-19 15:50:40,389 - 91: enter network_call
2016-02-19 15:50:40,389 - 92: exit gevent_yield
2016-02-19 15:50:40,389 - 92: enter block
2016-02-19 15:50:40,443 - 92: exit block
2016-02-19 15:50:40,444 - 92: enter network_call
2016-02-19 15:50:40,444 - 89: exit network_call
2016-02-19 15:50:40,444 - 89: done
127.0.0.1 -- [2016-02-19 15:50:40] "GET / HTTP/1.1" 200
119 0.275503
2016-02-19 15:50:40,444 - 90: exit network_call
2016-02-19 15:50:40,444 - 90: done
127.0.0.1 -- [2016-02-19 15:50:40] "GET / HTTP/1.1" 200
119 0.165371
2016-02-19 15:50:40,445 - 93: enter gevent_yield
2016-02-19 15:50:40,445 - 93: exit gevent_yield
2016-02-19 15:50:40,445 - 93: enter block
2016-02-19 15:50:40,496 - 93: exit block
2016-02-19 15:50:40,496 - 93: enter network_call
2016-02-19 15:50:40,496 - 91: exit network_call
2016-02-19 15:50:40,496 - 91: done
127.0.0.1 -- [2016-02-19 15:50:40] "GET / HTTP/1.1" 200
119 0.163295
```

```
2016-02-19 15:50:40,497 - 94: enter gevent_yield
2016-02-19 15:50:40,497 - 95: enter gevent_yield
2016-02-19 15:50:40,497 - 94: exit gevent_yield
2016-02-19 15:50:40,497 - 94: enter block
2016-02-19 15:50:40,549 - 94: exit block
2016-02-19 15:50:40,549 - 94: enter network_call
2016-02-19 15:50:40,550 - 95: exit gevent_yield
2016-02-19 15:50:40,550 - 95: enter block
2016-02-19 15:50:40,603 - 95: exit block
2016-02-19 15:50:40,603 - 95: enter network_call
2016-02-19 15:50:40,603 - 92: exit network_call
2016-02-19 15:50:40,603 - 92: done
```

- Extra 50ms block in first gevent_yield
- network_call takes 50 extra ms (159 total)
- 100ms extra from blocks from other requests
- ~250ms (269ms) instead of expected 150ms

Gevent

Impact

- **Performance degrades with increased**
 - CPU work
 - Requests per process
 - **Autoscale frontend at 40% CPU usage**
-
-
-
- **Reduce latency by doing fewer requests per process**
 - config changes

Gevent

Goal

- Frontend clusters are some of the most expensive
- High latency

Increase throughput (RPS / host) and decrease latency by decreasing CPU usage

- Goal: 33% efficiency gain



Problem

Gevent



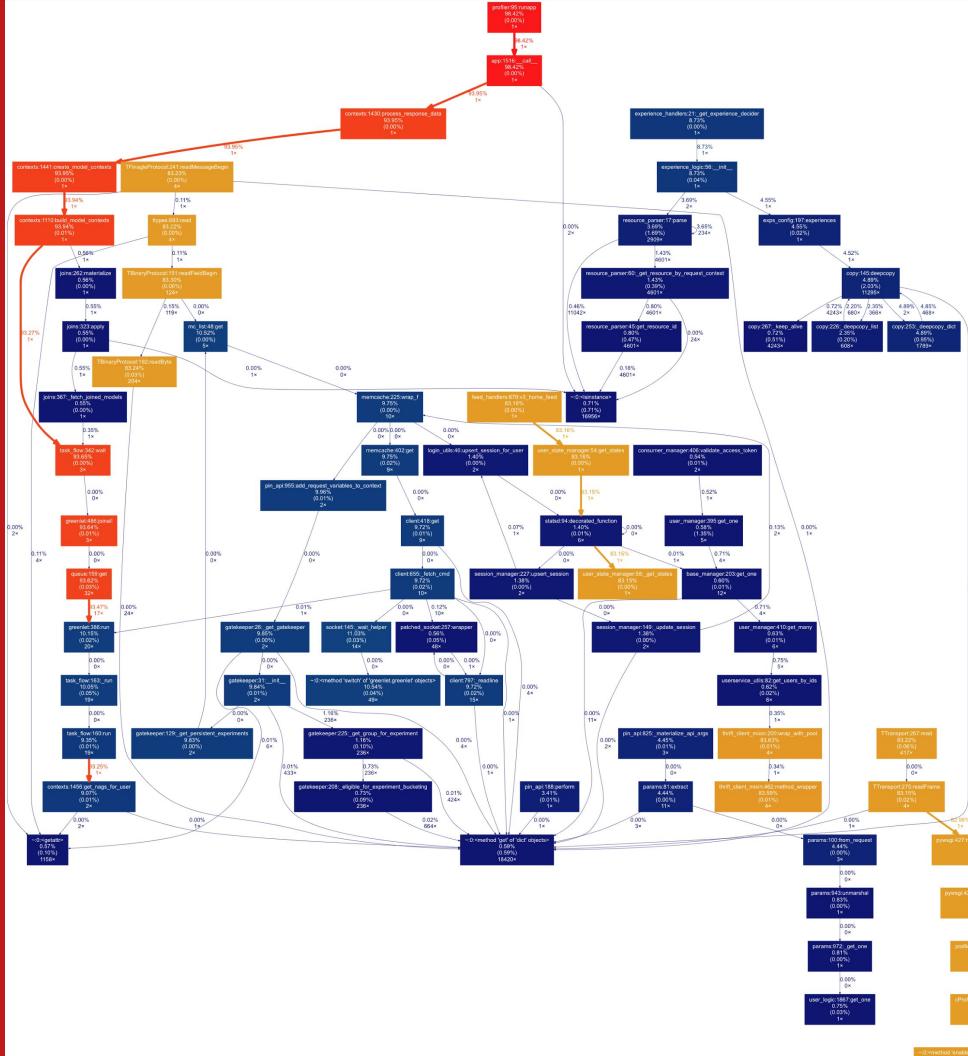
Profiling

Impact

Python Profilers

1. API cProfiler middleware
2. Line profiler
3. Sampling profiler and flamegraphs

API Profiler



API Profiler

- **Profiles function calls**
- **Save each request to a separate file**
- **Large overhead**
 - can't be used in production

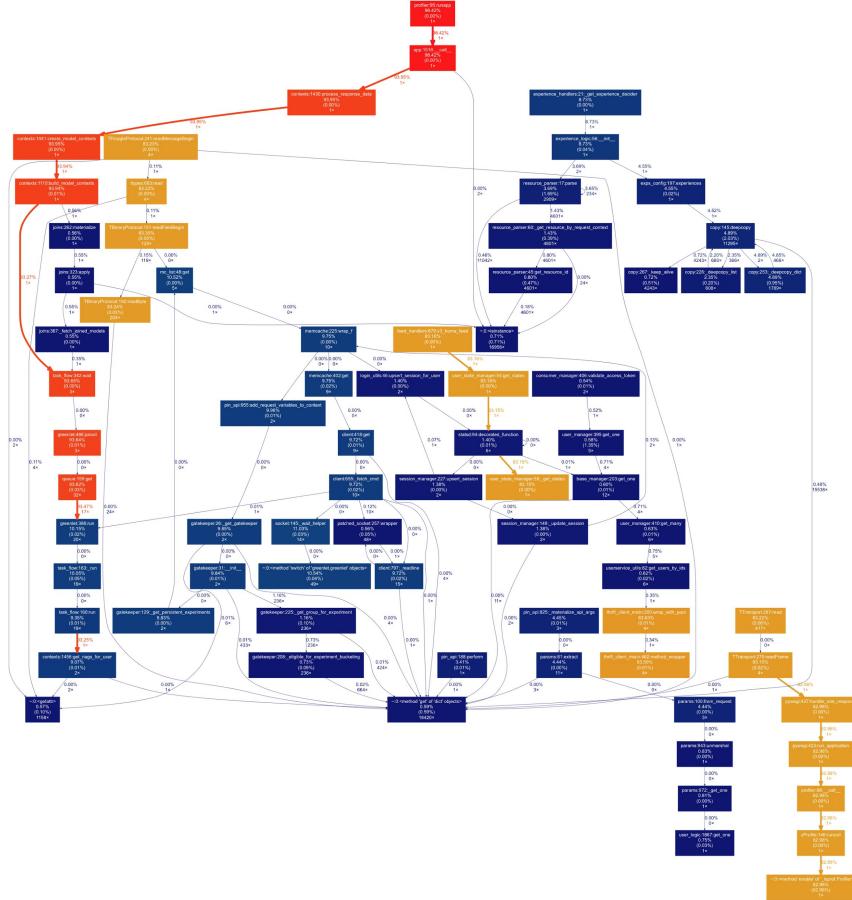
```
diff --git a/api/start.py b/api/start.py
index e910aa9..b0f58f0 100755
--- a/api/start.py
+++ b/api/start.py
@@ -5,6 +5,9 @@ import sys
sys.path.insert(0, '..')
sys.path.insert(1, '../common')

+from werkzeug.contrib.profiler import ProfilerMiddleware
+
+
import common.utils.patch_for_gevent
common.utils.patch_for_gevent.patch()

@@ -45,7 +48,7 @@ def main(unused_argv):
    if settings.API_WSGI_SERVER_LOGS_ENABLED is not None:
        log = 'default' if settings.API_WSGI_SERVER_LOGS_ENABLED else None

-    pywsgi.WSGIServer(("0.0.0.0", port), app, log=log).serve_forever()
+    pywsgi.WSGIServer(("0.0.0.0", port), ProfilerMiddleware(app, profile_dir='/home/jogo/tmp'), log=log).serve_forever()
```

GET /v3/feeds/home



```
python -mpstats GET.v3.search.pins.000908ms.1452734245.prof
```

python pstats module

Module for profiler data analysis

pstats module

stats

```
GET.v3.search.pins.000908ms.1452734245.prof% stats 10
```

```
Wed Jan 13 17:26:41 2016    GET.v3.search.pins.000908ms.1452734245.prof
```

```
34464 function calls (34150 primitive calls) in 0.956 seconds
```

```
Ordered by: internal time
```

```
List reduced from 1247 to 10 due to restriction <10>
```

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.809	0.809	0.809	0.809	{method 'enable' of '_lsprof.Profiler' objects}
12	0.056	0.005	0.057	0.005	/home/jogo/code/.venv/local/lib/python2.7/site-packages/gevent/socket.py:153(wait_read)
3	0.048	0.016	0.860	0.287	/home/jogo/code/.venv/local/lib/python2.7/site-packages/thrift/transport/TTransport.py:275(readFrame)
189	0.001	0.000	0.001	0.000	../core/gatekeeper/utils.py:6(pick_key_by_mod)
872	0.001	0.000	0.001	0.000	{method 'search' of '_sre.SRE_Pattern' objects}
414	0.001	0.000	0.012	0.000	../core/gatekeeper/experiment.py:206(get_group)
189	0.001	0.000	0.002	0.000	../core/gatekeeper/experiment.py:291(get_mod_group)
414	0.001	0.000	0.013	0.000	../core/gatekeeper/user_experiments.py:80(_convert_experiment_object_to_dict)
201	0.001	0.000	0.001	0.000	/usr/lib/python2.7/threading.py:120(acquire)
2195	0.001	0.000	0.001	0.000	{method 'get' of 'dict' objects}

```
GET.v3.search.pins.000908ms.1452734245.prof% stats _convert_experiment_object_to_dict
```

```
Wed Jan 13 17:26:41 2016    GET.v3.search.pins.000908ms.1452734245.prof
```

```
34464 function calls (34150 primitive calls) in 0.956 seconds
```

```
Ordered by: internal time
```

```
List reduced from 1247 to 1 due to restriction '<_convert_experiment_object_to_dict>'
```

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
414	0.001	0.000	0.013	0.000	../core/gatekeeper/user_experiments.py:80(_convert_experiment_object_to_dict)

pstats module

callees & callers

```
GET.v3.search.pins.000908ms.1452734245.prof% callees _convert_experiment_object_to_dict
Random listing order was used
List reduced from 1247 to 1 due to restriction <'_convert_experiment_object_to_dict'>
```

```
Function                      called...
                                ncalls  tottime  cumtime
..../core/gatekeeper/user_experiments.py:80(_convert_experiment_object_to_dict)  ->    414    0.001    0.012  ..../core/gatekeeper/experiment.py:206(get_group)
                                                414    0.000    0.000  {isinstance}
```

```
GET.v3.search.pins.000908ms.1452734245.prof% callers _convert_experiment_object_to_dict
Random listing order was used
List reduced from 1247 to 1 due to restriction <'_convert_experiment_object_to_dict'>
```

```
Function                      was called by...
                                ncalls  tottime  cumtime
..../core/gatekeeper/user_experiments.py:80(_convert_experiment_object_to_dict)  <-    414    0.001    0.013  ..../core/gatekeeper/user_experiments.py:53(get_limited_user_experiments)
```

pstats module

Find CPU bound function calls

```
jogo@lappy:~/scratchpad/profiling/Jan-13-2015$ python walk.py GET.v3.search.pins.000908ms.1452734245.prof
cumtime 0.8088 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/gevent/pywsgi.py:427(handle_one_response)
cumtime 0.8088 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/gevent/pywsgi.py:423(run_application)
cumtime 0.8088 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/werkzeug/contrib/profiler.py:88(__call__)
cumtime 0.8088 - /usr/lib/python2.7/profile.py:146(runcall)
cumtime 0.8088 - ~:0<method 'enable' of '_lsprof.Profiler' objects>
cumtime 0.0040 - .../core/experiments/gatekeeper.py:225(_get_group_for_experiment)
cumtime 0.0038 - .../core/models/user_model.py:306(is_employee)
cumtime 0.0036 - .../common/utils/employee_utils.py:121(is_employee_user)
cumtime 0.0032 - .../common/utils/employee_utils.py:95(is_employee_account)
cumtime 0.0026 - .../core/experiments/gatekeeper.py:208(_eligible_for_experiment_bucketing)
cumtime 0.0024 - ~:0<getattr>()
cumtime 0.0023 - .../common/utils/decorators.py:438(__call__)
cumtime 0.0017 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/werkzeug/local.py:340(__getattr__)
cumtime 0.0016 - .../api/serializers/context.py:1341(from_request)
cumtime 0.0015 - .../common/utils/url_utils.py:434(__init__)
cumtime 0.0015 - .../core/gatekeeper/utils.py:6(pick_key_by_mod)
cumtime 0.0014 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/ua_parser/user_agent_parser.py:175(Parse)
cumtime 0.0013 - .../core/utils/ratelimit.py:103(keys_to_check)
cumtime 0.0012 - .../core/utils/ratelimit.py:94(generate_key)
cumtime 0.0011 - .../core/gatekeeper/filters/filters.py:549(test)
cumtime 0.0011 - .../config/experiments.py:2558(<lambda>)
cumtime 0.0010 - /usr/lib/python2.7/threading.py:120(acquire)
cumtime 0.0010 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/thrift/protocol/TBase.py:56(write)
cumtime 0.0010 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/thrift/protocol/TProtocol.py:373(writeStruct)
cumtime 0.0010 - .../common/utils/version.py:71(__init__)
cumtime 0.0010 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/thrift/protocol/TProtocol.py:393(writeFieldByTType)
```

Idea: Automatically figure out which functions take a long time and don't use the network

Due to not collecting full stacktraces, we can only find out callees and go that way, so would work if everything is only called once.

```
"""
import pstats
import sys

p = pstats.Stats(sys.argv[1])
p.calc_callees()

def get_children(func, visited=None):
    """Return set of all functions called by func (recursive)"""
    if visited is None:
        visited = set()
    visited.add(func)
    callees = p.all_callees[func]
    for callee in callees:
        if callee not in visited:
            visited.update(get_children(callee, visited))
    return visited

def uses_network(children):
    """Checks a list if children to see if the network is called"""
    return "socket.py" in str(children)

def get_network_free_functions():
    functions = []
    for func in p.all_callees:
        children = get_children(func)
        if not uses_network(children):
            functions.append(func)

    # sort by cumtime
    functions = sorted(functions, key=lambda x: p.stats[x][3], reverse=True)
    return functions

def format_function_name(func):
    (path, line, name) = func
    return ("%s:%s(%s)" % (path, line, name))

# print top 50 slowest non network bound functions
for func in get_network_free_functions()[:50]:
    ncalls, ncalles, tottime, cumtime, callers = p.stats[func]
    print "cumtime %.4f - %s" % (cumtime, format_function_name(func))
```

pstats module

Find CPU bound function calls

```
jogo@lappy:~/scratchpad/profiling/Jan-13-2015$ python walk.py GET.v3.search.pins.000908ms.1452734245.prof
cumtime 0.8088 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/gevent/pywsgi.py:427(handle_one_response)
cumtime 0.8088 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/gevent/pywsgi.py:423(run_application)
cumtime 0.8088 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/werkzeug/contrib/profiler.py:88(__call__)
cumtime 0.8088 - /usr/lib/python2.7/cProfile.py:146(runcall)
cumtime 0.8088 - ~:0(<method 'enable' of '_lsprof_Profiler' objects>)
cumtime 0.0040 - ../core/experiments/gatekeeper.py:225(__get_group_for_experiment)
cumtime 0.0038 - ../core/models/user_model.py:306(is_employee)
cumtime 0.0036 - ../common/utils/employee_utils.py:121(is_employee_user)
cumtime 0.0032 - ../common/utils/employee_utils.py:95(is_employee_account)
cumtime 0.0026 - ../core/experiments/gatekeeper.py:280(_eligible_for_experiment_bucketing)
cumtime 0.0024 - ~:0(<getattr>)
cumtime 0.0023 - ../common/utils/decorators.py:438(__call__)
cumtime 0.0017 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/werkzeug/local.py:340(__getattr__)
cumtime 0.0016 - ../api/serializers/context.py:1341(from_request)
cumtime 0.0015 - ../common/utils/url_utils.py:434(__init__)
cumtime 0.0015 - ../core/gatekeeper/utils.py:6(pick_key_by_mod)
cumtime 0.0014 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/ua_parser/user_agent_parser.py:175(Parse)
cumtime 0.0013 - ../core/utils/ratelimit.py:103(keys_to_check)
cumtime 0.0012 - ../core/utils/ratelimit.py:94(generate_key)
cumtime 0.0011 - ../core/gatekeeper/filters/filters.py:549(test)
cumtime 0.0011 - ../config/experiments.py:2558(<lambda>)
cumtime 0.0010 - /usr/lib/python2.7/threading.py:120(acquire)
cumtime 0.0010 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/thrift/protocol/TBase.py:56(write)
cumtime 0.0010 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/thrift/protocol/TProtocol.py:373(writeStruct)
cumtime 0.0010 - ../common/utils/version.py:71(__init__)
cumtime 0.0010 - /home/jogo/code/.venv/local/lib/python2.7/site-packages/thrift/protocol/TProtocol.py:393(writeFieldByTType)
```

Idea: Automatically figure out which functions take a long time and don't use the network

Due to not collecting full stacktraces, we can only find out callees and go that way, so would work if everything is only called once.

```
"""
import pstats
import sys

p = pstats.Stats(sys.argv[1])
p.cal_callees()

def get_children(func, visited=None):
    """Return set of all functions called by func (recursive)"""
    if visited is None:
        visited = set()
    visited.add(func)
    callees = p.all_callees[func]
    for callee in callees:
        if callee not in visited:
            visited.update(get_children(callee, visited))
    return visited

def uses_network(children):
    """Checks a list if children to see if the network is called"""
    return "socket.py" in str(children)

def get_network_free_functions():
    functions = []
    for func in p.all_callees:
        children = get_children(func)
        if not uses_network(children):
            functions.append(func)

    # sort by cumtime
    functions = sorted(functions, key=lambda x: p.stats[x][3], reverse=True)
    return functions

def format_function_name(func):
    (path, line, name) = func
    return ("%s:%s(%s)" % (path, line, name))

# print top 50 slowest non network bound functions
for func in get_network_free_functions()[:50]:
    ncalls, ncalles, tottime, cumtime, callers = p.stats[func]
    print "cumtime %.4f - %s" % (cumtime, format_function_name(func))
```

Line Profiler

Function: get_url_structures at line 744

Line #	Hits	Time	Per Hit	% Time	Line Contents
744					@profile
745					def get_url_structures(url):
746					"""Parse and return domain/reverse-domain/host/l1 from a url.
747					Example:
748					"http://image.pinterest.com/ben/?a=1":
749					{"domain": "pinterest.com",
750					"reverse-domain": "com.pinterest/", # always have trailing
751					"host": "image.pinterest.com",
752					"reverse-host": "com.pinterest.image",
753					"l1": "image.pinterest.com/ben/" # always have trailing /
754					"reverse-l1": "com.pinterest.image/ben/" # always have trailing }
755					
756					Returns:
757					A dict with fields above. It is not necessary to contain all field
758					For example, "http://pinterest.com" doesn't contain l1 path.
759					"""
760	20001	17592	0.9	1.6	url_structure = {}
761	20001	16756	0.8	1.5	if url:
762	20001	140586	7.0	12.8	url_info = urlparse.urlparse(url)
763	20001	26608	1.3	2.4	if not url_info.scheme:
764					url = "http://%s" % url
765					url_info = urlparse.urlparse(url)
766	20001	73095	3.7	6.7	if url_info.hostname:
767					# host is case insensitive
768					# special treatment for "www."
769	20001	73464	3.7	6.7	normalized_netloc = netloc = url_info.hostname.lower().rstrip()
770	20001	23757	1.2	2.2	if netloc.startswith("www.":
771	10000	9668	1.0	0.9	netloc = netloc[4:]
772	20001	18024	0.9	1.6	url_structure[URL_HOST_STR] = netloc
773	20001	24235	1.2	2.2	host_segs = netloc.split('.')
774	20001	18156	0.9	1.7	host_segs.reverse()
775	20001	26056	1.3	2.4	url_structure[URL_REVERSE_HOST_STR] = '. '.join(host_segs) + "/
776	20001	26807	1.3	2.4	path = url_info.path.lstrip("//")
777	20001	16678	0.8	1.5	if path:
778	20001	22443	1.1	2.0	path_segs = path.split("/")
779	20001	16154	0.8	1.5	url_structure[URL_L1_STR] = "%s/%s/" % (
780	20001	26940	1.3	2.5	netloc, path_segs[0])
781	20001	16161	0.8	1.5	url_structure[URL_REVERSE_L1_STR] = "%s%s/" % (
782	20001	23272	1.2	2.1	url_structure[URL_REVERSE_HOST_STR], path_segs[0])
783					
784	20001	337160	16.9	30.7	domain = psl.get_public_suffix(normalized_netloc.lower())
785					
786	20001	19660	1.0	1.8	if domain:
787	20001	18633	0.9	1.7	url_structure[URL_DOMAIN_STR] = domain
788	20001	25030	1.3	2.3	domain_segs = domain.split('.')
789	20001	18582	0.9	1.7	domain_segs.reverse()
790	20001	16944	0.8	1.5	url_structure[URL_REVERSE_DOMAIN_STR] = '. '.join(
791	20001	28092	1.4	2.6	domain_segs) + "/"
792	20001	16154	0.8	1.5	return url_structure
~					
~					
~					
~					
~					

Line Profiler

https://github.com/rkern/line_profiler

```
Line Contents
=====
@profile
def _is_employee_test_user(email):
    """Return True if account has a pinterest email of a current employee

    This function does not check if the email is confirmed.
    """
    pattern = r"^(?P<user>[a-zA-Z0-9_.-]+)(?P<at>@)[a-zA-Z0-9_.-]+\\.[a-zA-Z0-9_.-]+$"
    if isinstance(email, str) or isinstance(email, unicode):
        match = re.match(pattern, email)
        if match:
            _, user, _, domain = match.groups()
            if domain == "pinterest.com":
                return LdapList().is_ldap(user)
    return False
```

Line Profiler

https://github.com/rkern/line_profiler

```
@profile  
def slow_function(a, b, c):  
    ...
```

```
$ kernprof -l script_to_profile.py
```

```
$ python -m line_profiler script_to_profile.py.lprof
```

```
Function: _is_employee_test_user at line 80  
  
Line #      Hits       Time  Per Hit   % Time  Line Contents  
=====  ======  ======  ======  ======  ======  
     80                      @profile  
     81                      def _is_employee_test_user(email):  
     82                          """Return True if account has a pinterest email of a current employee  
     83  
     84                          This function does not check if the email is confirmed.  
     85                          """  
     86     200          128      0.6      1.5  pattern = r"^(?P<user>[a-zA-Z0-9_.-]+)(?P<domain>[a-zA-Z0-9_.-]+@[a-zA-Z0-9_.-]+\.[a-zA-Z0-9_.-]+)$"  
     87     200          133      0.7      1.5  if isinstance(email, str) or isinstance(email, unicode):  
     88     200         4808     24.0     54.7  match = re.match(pattern, email)  
     89     200          123      0.6      1.4  if match:  
     90     200          194      1.0      2.2  _, user, _, domain = match.groups()  
     91     200          122      0.6      1.4  if domain == "pinterest.com":  
     92     100         3220     32.2     36.6  return LdapList().is_ldap(user)  
     93     100           58      0.6      0.7  return False  
  
Total time: 0.001775 s  
File: common/utils/employee_utils.py  
Function: _is_employee_account at line 95
```

Line Profiler

https://github.com/rkern/line_profiler

```
@profile  
def slow_function(a, b, c):  
    ...
```

```
$ kernprof -l script_to_profile.py
```

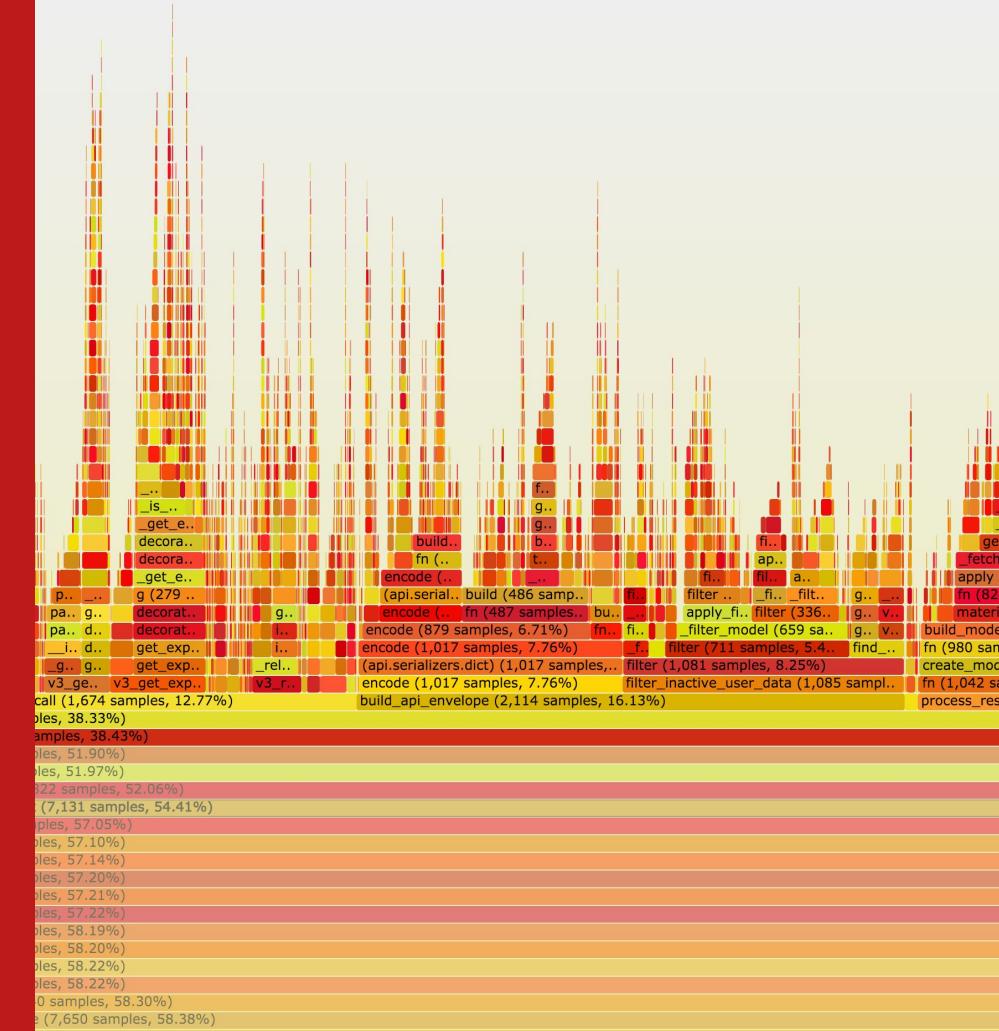
```
$ python -m line_profiler script_to_profile.py.lprof
```

```
Function: _is_employee_test_user at line 80  
  
Line #      Hits       Time  Per Hit   % Time  Line Contents  
=====  ======  ======  ======  ======  ======  
     80                      @profile  
     81                      def _is_employee_test_user(email):  
     82                          """Return True if account has a pinterest email of a current employee  
     83  
     84                          This function does not check if the email is confirmed.  
     85                          """  
     86     200          128      0.6      1.5  pattern = r"^(?P<user>[a-zA-Z0-9_.-]+)(?P<domain>[a-zA-Z0-9_.-]*@[a-zA-Z0-9-]+\.[a-zA-Z0-9-\.]+)$"  
     87     200          133      0.7      1.5  if isinstance(email, str) or isinstance(email, unicode):  
     88     200          4808     24.0     54.7  match = re.match(pattern, email)  
     89     200          123      0.6      1.4  if match:  
     90     200          194      1.0      2.2  _, user, _, domain = match.groups()  
     91     200          122      0.6      1.4  if domain == "pinterest.com":  
     92     100          3220     32.2     36.6  return LdapList().is_ldap(user)  
     93     100           58      0.6      0.7  return False  
  
Total time: 0.001775 s  
File: common/utils/employee_utils.py  
Function: _is_employee_account at line 95
```



Sampling Profiler

Flame Graphs

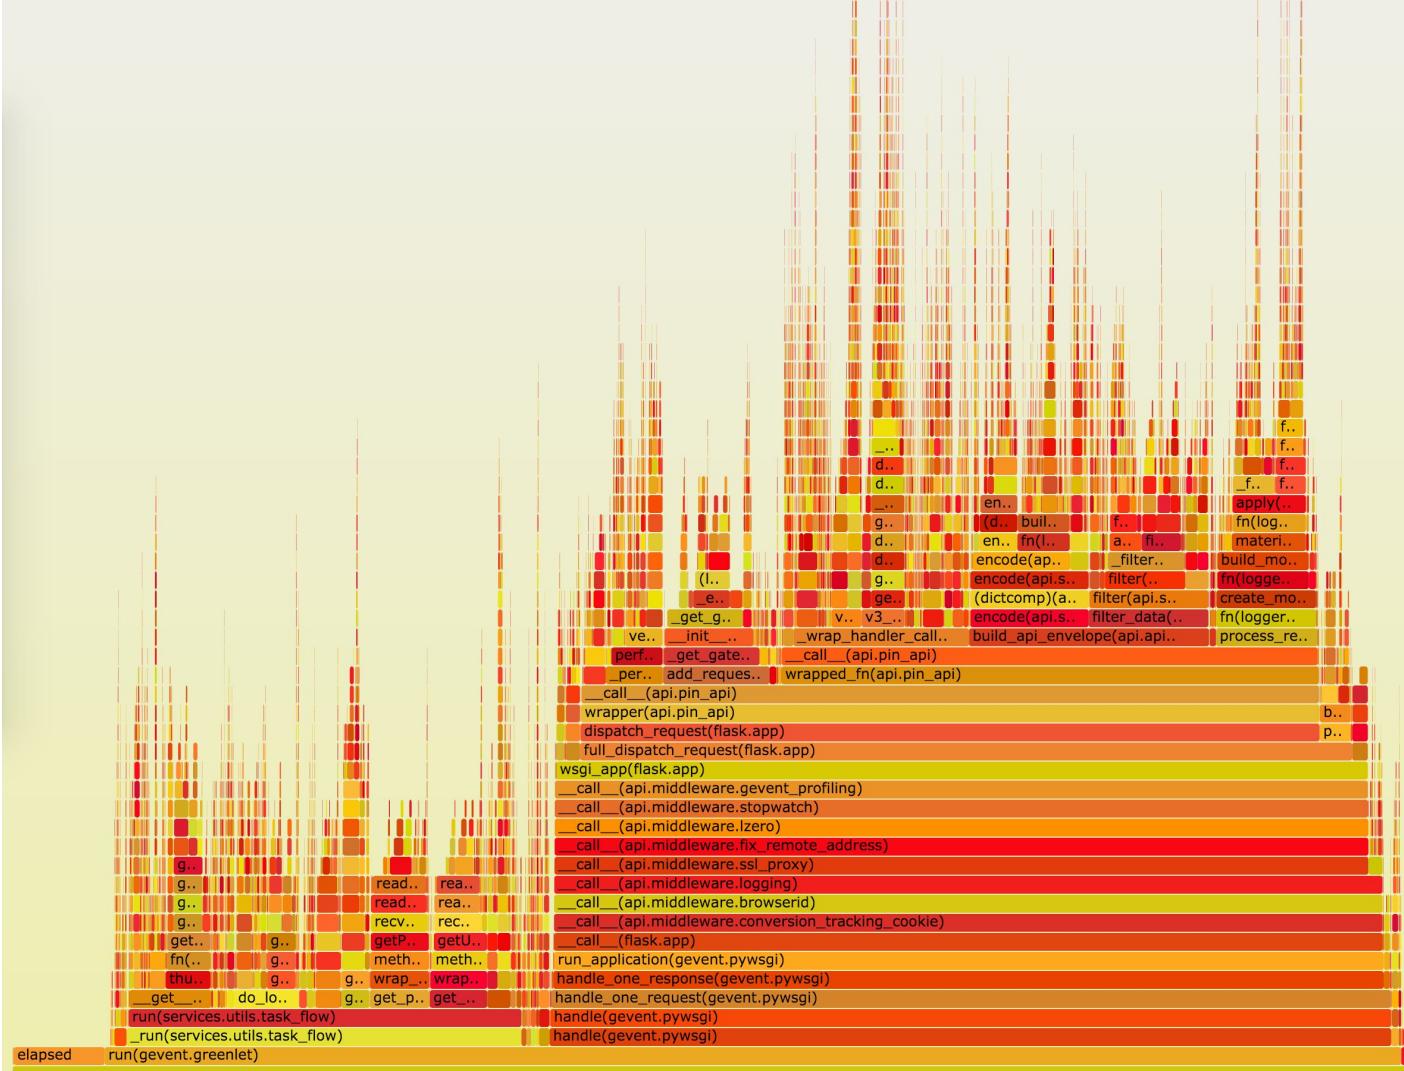


Sampling Profiler

<https://github.com/nylas/nylas-perf-tools>

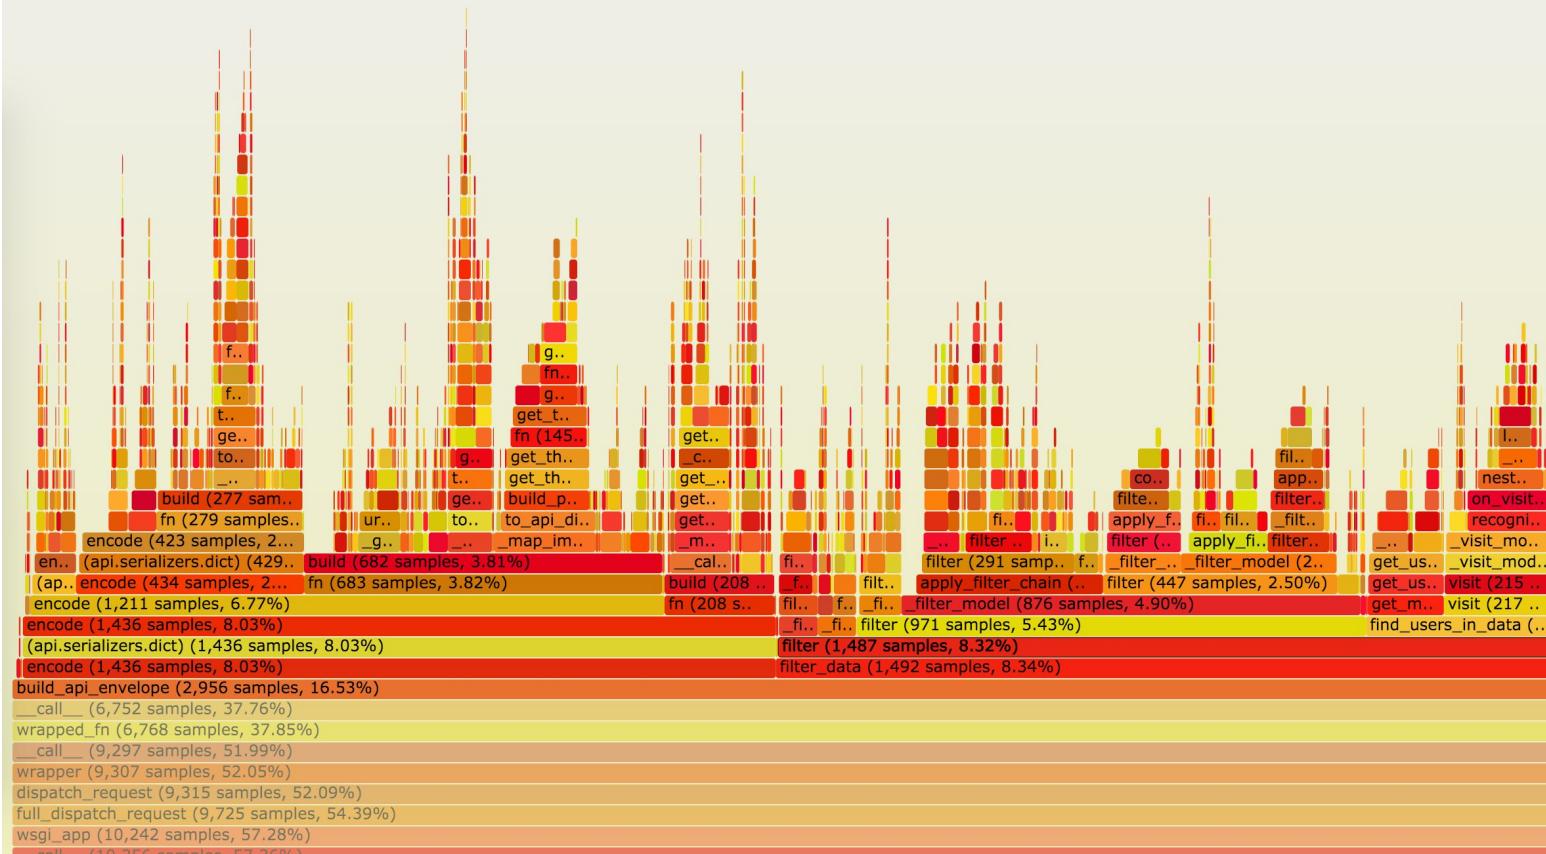
- **Low overhead**
 - can be used in production
- **Records stack every 5ms**
- **Emits data over**
http://localhost:16384
- **Format**
 - Stack: Count
- **Fuzzy**
- **Flame graphs**

```
24 class Sampler(object):
25     """
26     A simple stack sampler for low-overhead CPU profiling: samples the call
27     stack every `interval` seconds and keeps track of counts by frame. Because
28     this uses signals, it only works on the main thread.
29     """
30     def __init__(self, interval=0.005):
31         self.interval = interval
32         self._started = None
33         self._stack_counts = collections.defaultdict(int)
34
35     def start(self):
36         self._started = time.time()
37         try:
38             signal.signal(signal.SIGVTALRM, self._sample)
39         except ValueError:
40             raise ValueError('Can only sample on the main thread')
41
42         signal.setitimer(signal.ITIMER_VIRTUAL, self.interval, 0)
43
44     def _sample(self, signum, frame):
45         stack = []
46         while frame is not None:
47             stack.append(self._format_frame(frame))
48             frame = frame.f_back
49
50         stack = ';' .join(reversed(stack))
51         self._stack_counts[stack] += 1
52         signal.setitimer(signal.ITIMER_VIRTUAL, self.interval, 0)
```



[Reset Zoom](#)

Flame Graph

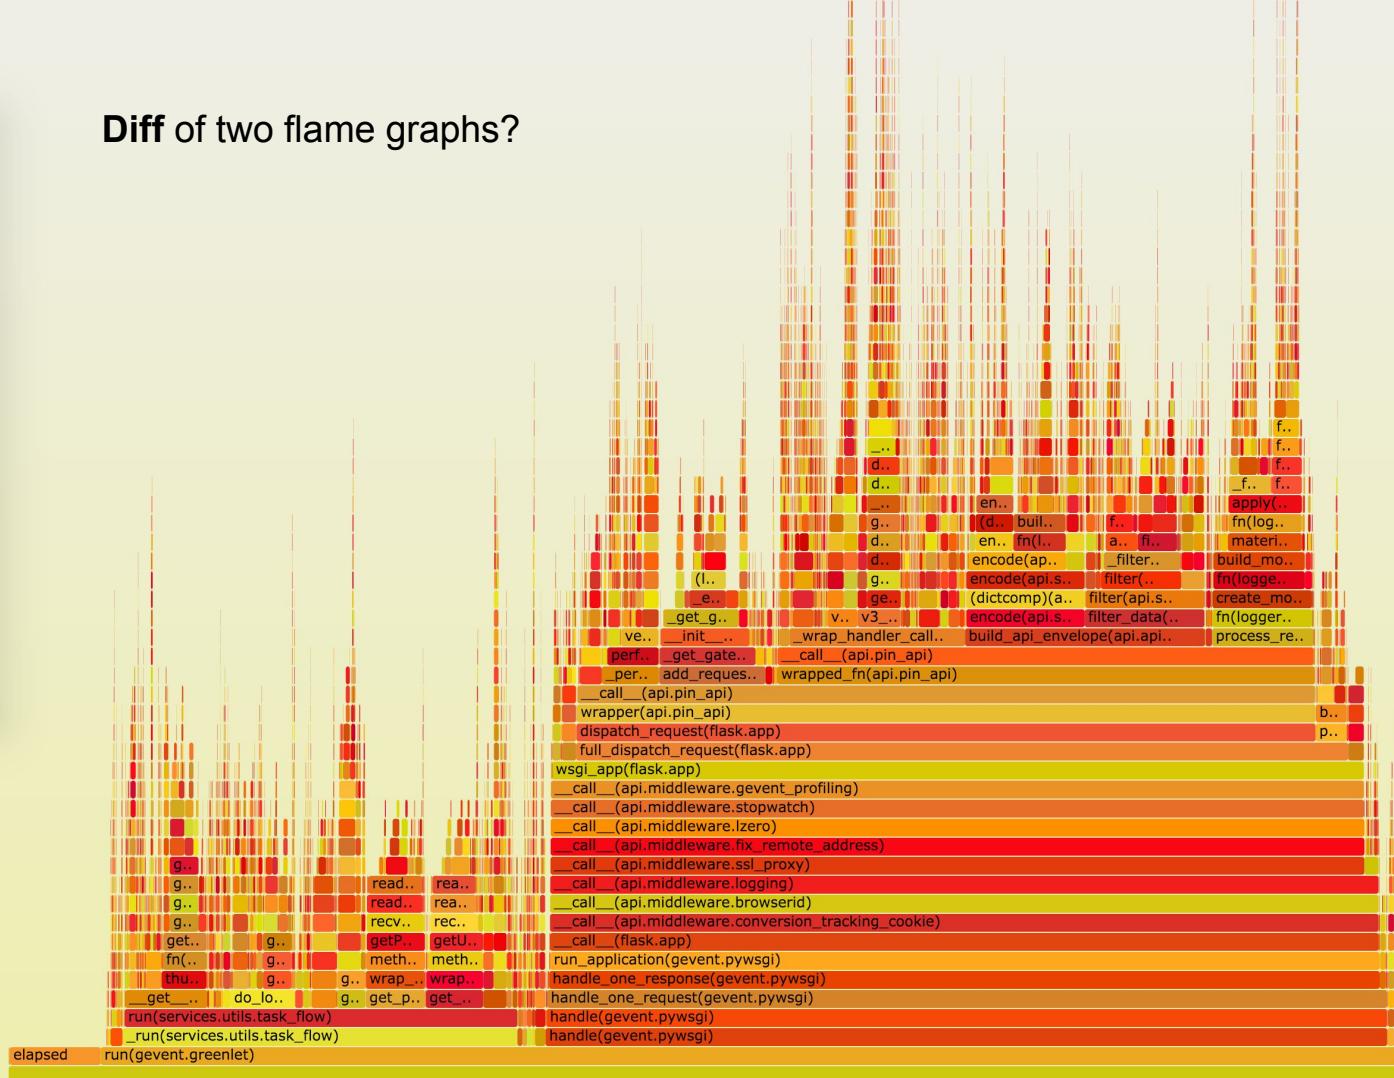
[Search](#)

Sampling Profiler

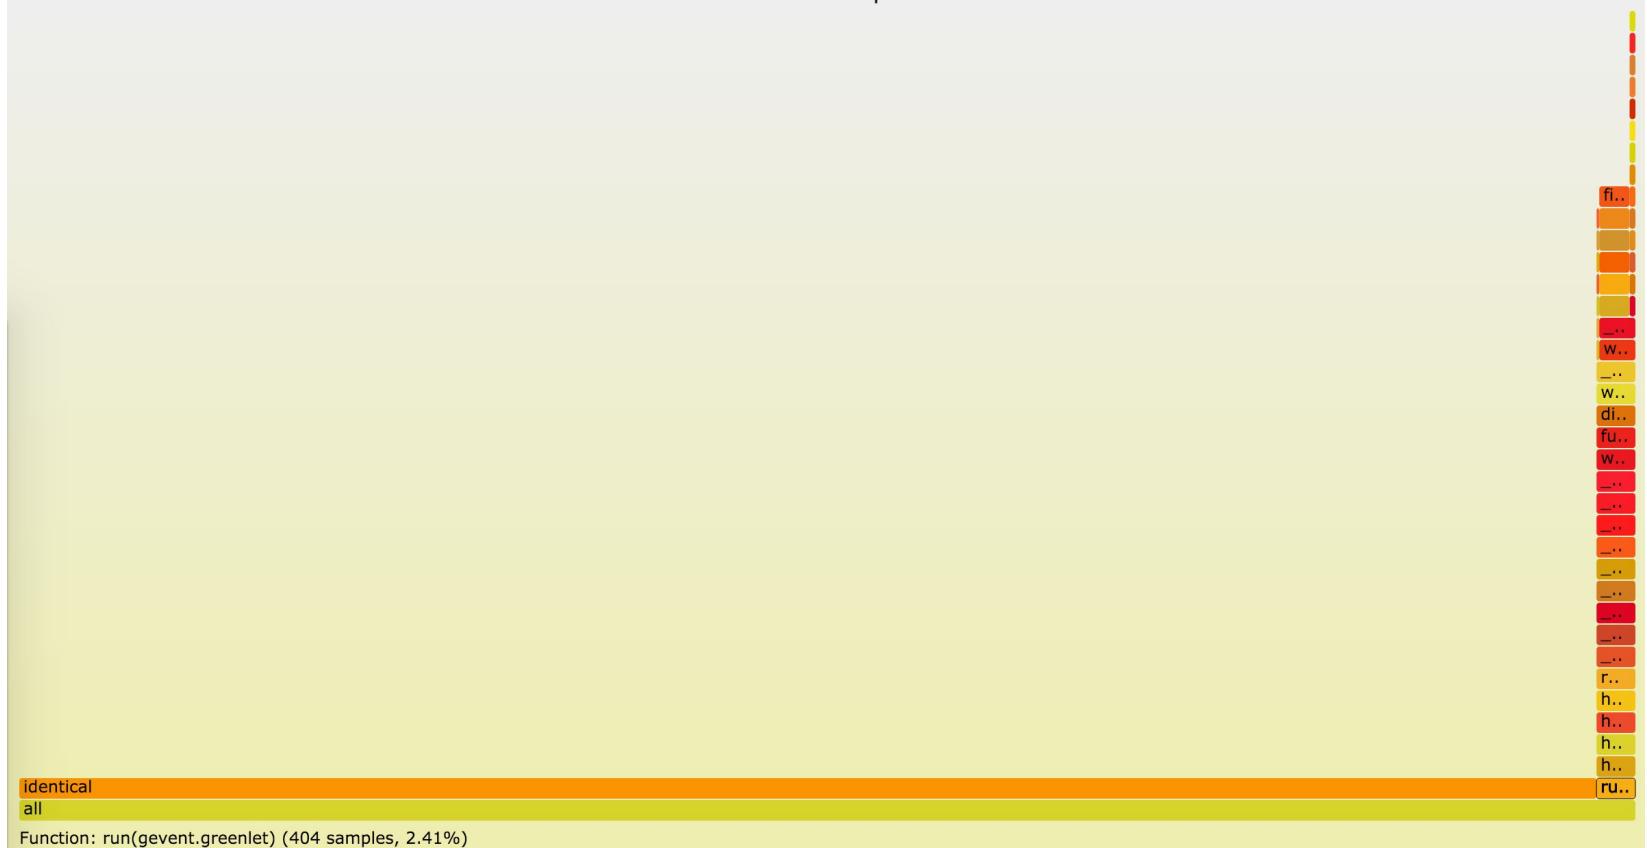
in production

- **\$ sudo kill -SIGUSR2 <pid>**
- **wait a few minutes to collect enough data**
- **\$ ssh host "curl localhost:16384" > profile.out**
- **\$ sudo kill -SIGUSR2 <pid>**

Diff of two flame graphs?



Diff Flamegraph



Diff Flamegraph

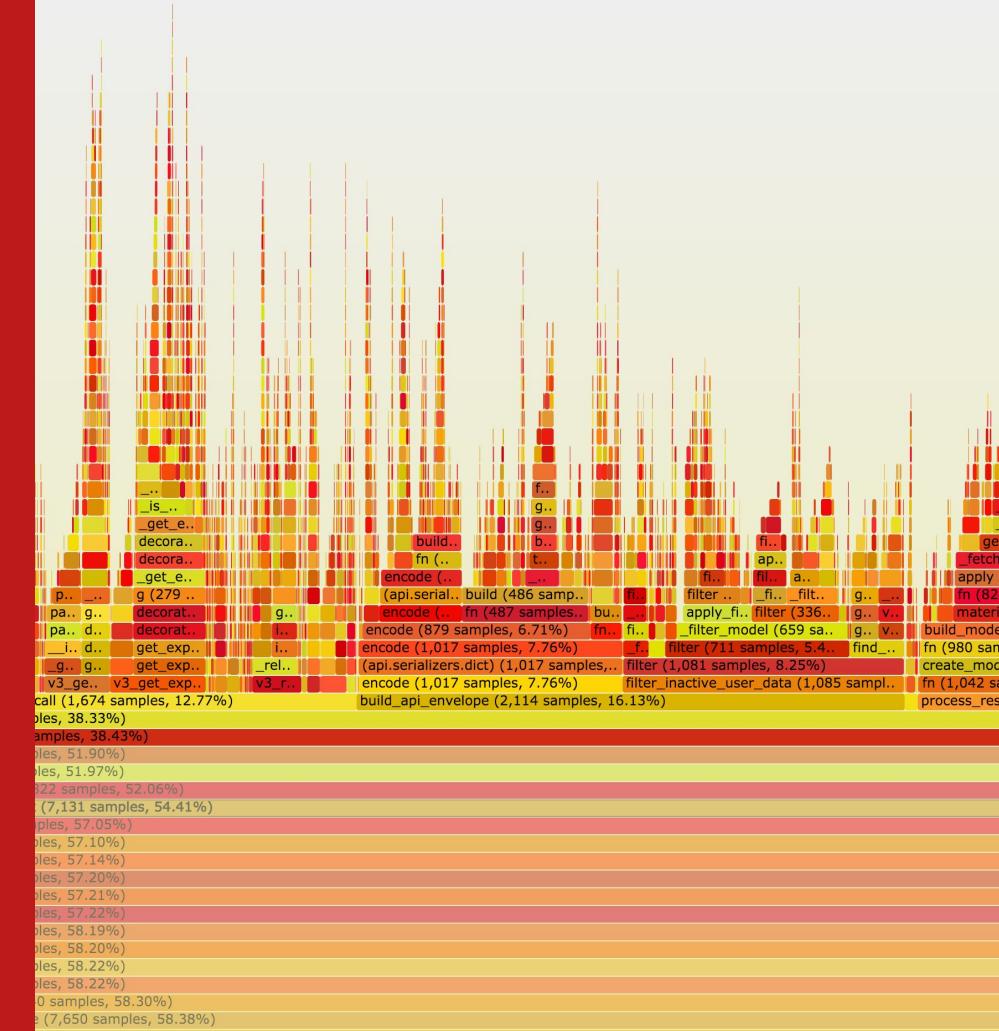
Reset Zoom

Flame Graph

Search



Examples



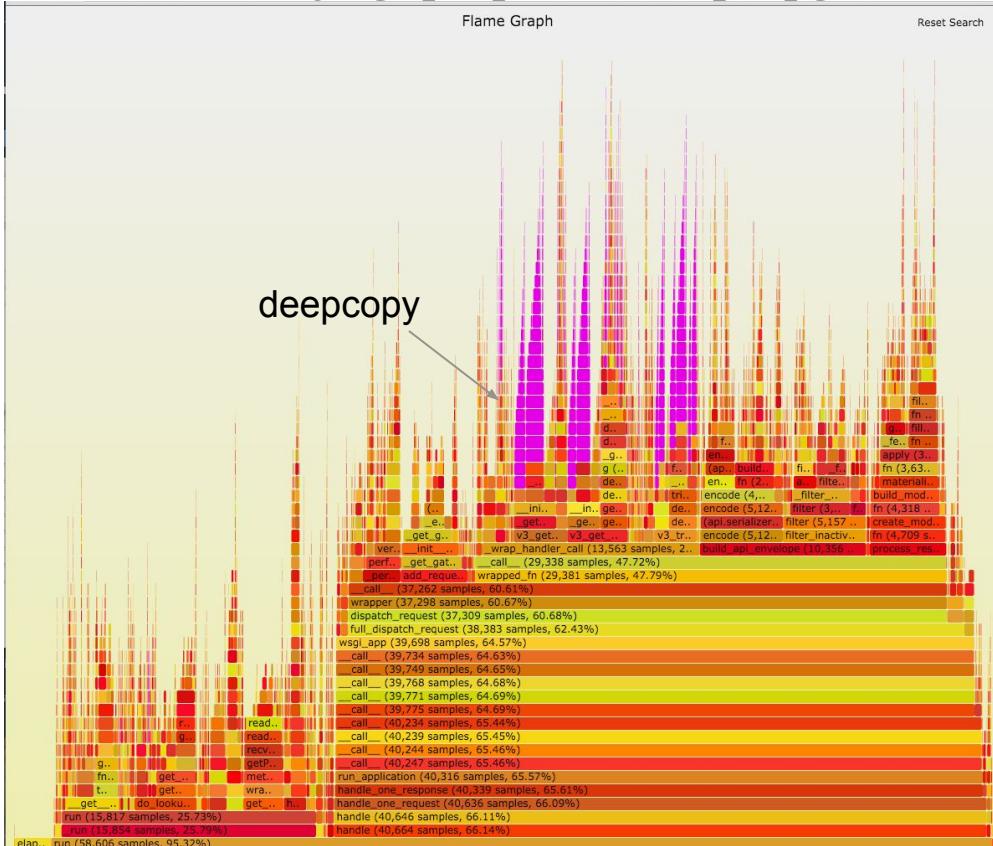
API Profiler

30ms of home feed spent in deepcopy

```
$ python walk.py GET.v3.feeds.home.000622ms.1452734743.prof
<snip>
cumtime 0.0553 - ../api/handlers/v3/experience_handlers.py:21(_get_experience_decider)
cumtime 0.0553 - ../core/logic/experiences/experience_logic.py:56(__init__)
cumtime 0.0309 - /usr/lib/python2.7/copy.py:145(deepcopy)
cumtime 0.0309 - /usr/lib/python2.7/copy.py:253(_deepcopy_dict)
cumtime 0.0288 - ../core/logic/experiences/exp_config/exp_config.py:197(experiences)
cumtime 0.0234 - ../core/logic/experiences/resource_parser.py:17(parse)
cumtime 0.0149 - /usr/lib/python2.7/copy.py:226(_deepcopy_list)
cumtime 0.0090 - ../core/logic/experiences/resource_parser.py:60(_get_resource_by_request_context)
...
```

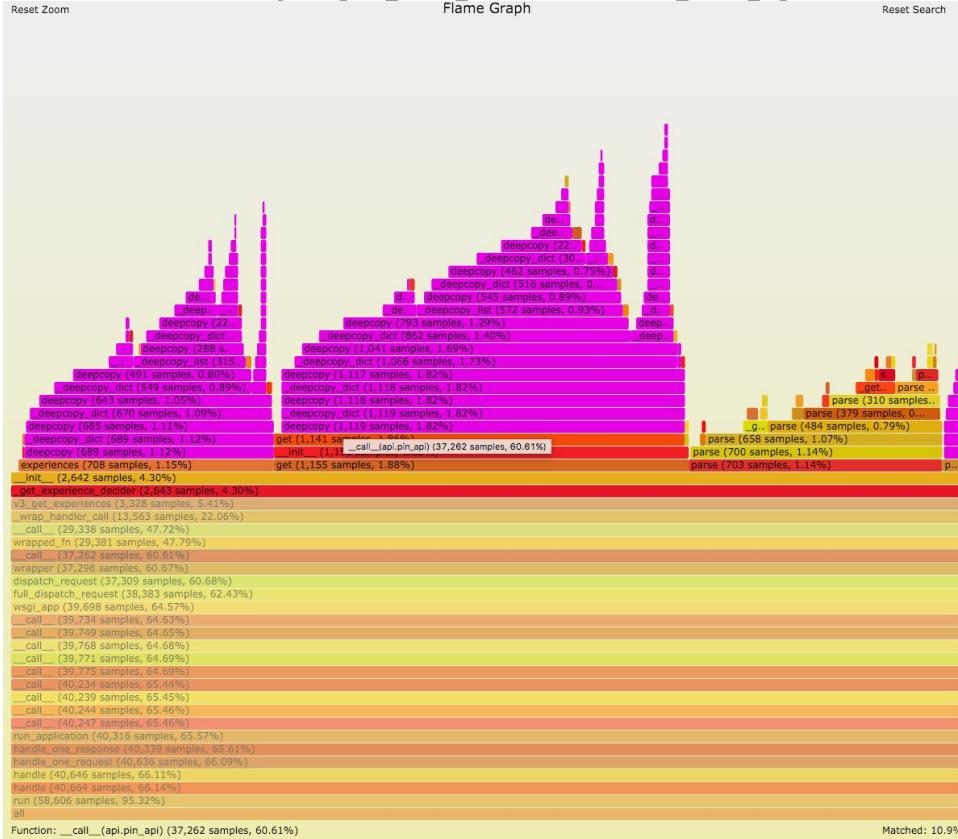
Sampling Profiler: Flame Graph

10% of ngapi spent in deepcopy



sampling profiler: flame graph

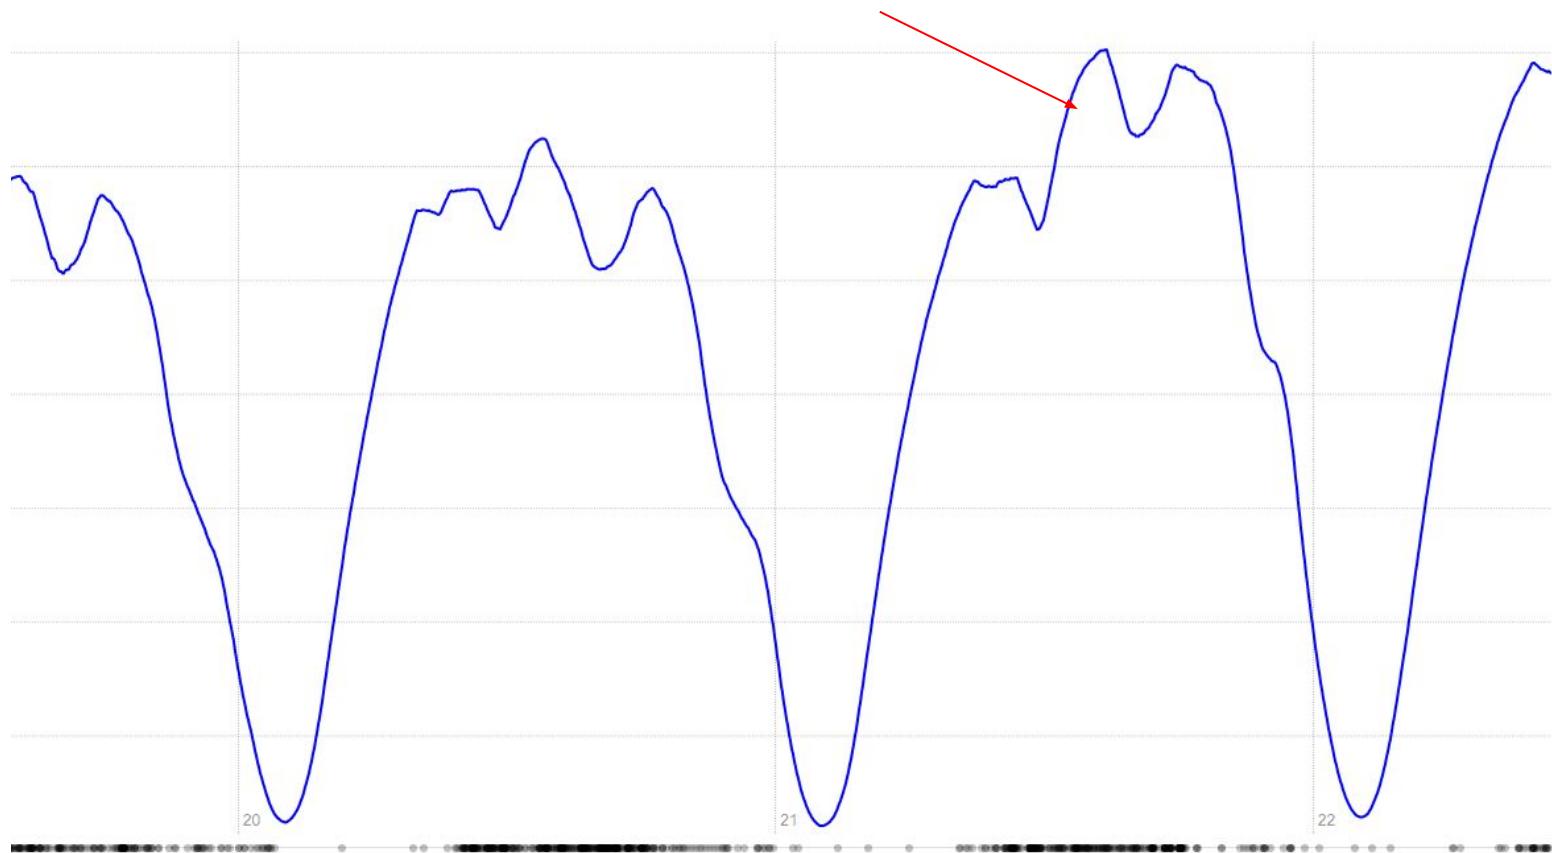
10% of ngapi spent in deepcopy



diff

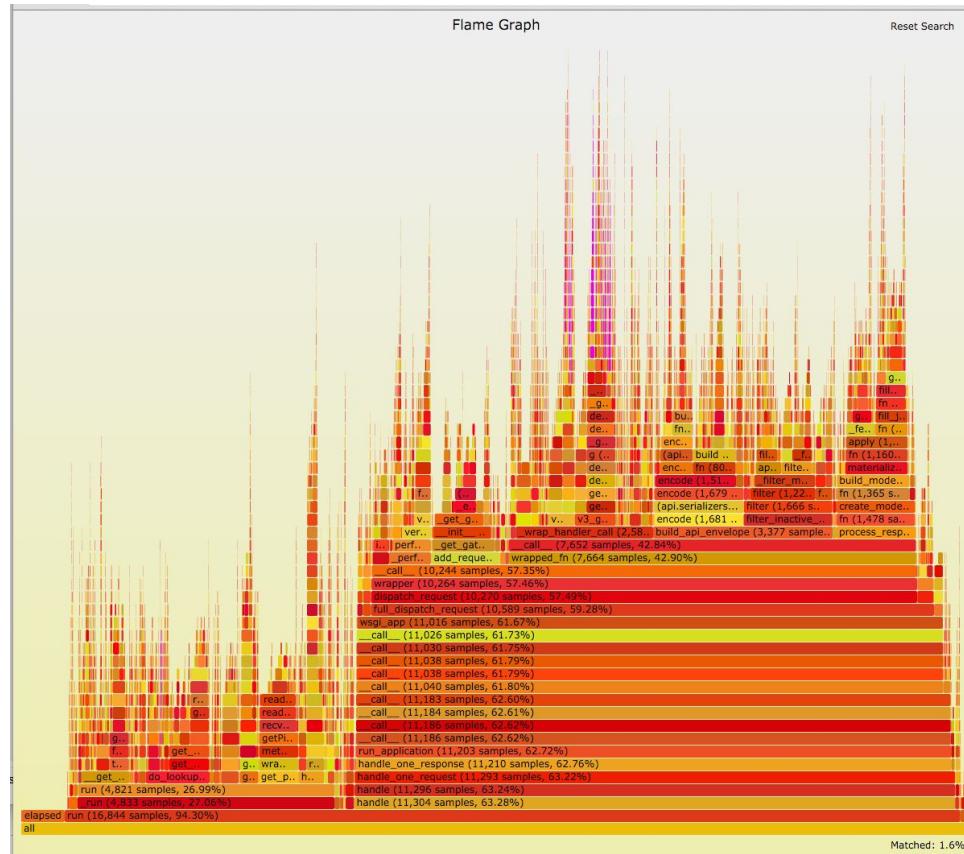
```
187
188     @property
189     def placements(self):
190         placement_configs = copy.deepcopy(self._static_placements)
191
192         if decider.decide_experiment(self.MANAGED_LIST_DARK_READ_DECIDER):
193             for placement_id, placement_config in self._whitelisted_managed_list_placements.iteritems():
194                 if placement_id not in placement_configs:
195                     self._incr_failure_stat('missing_static_config.placements.{}'.format(placement_id))
196                 elif placement_configs[placement_id] != placement_config:
197                     self._incr_failure_stat('mismatching_config.placements.{}'.format(placement_id))
198
199
200         ...
201
202
203
204         self._config = data
205
206     def get(self):
207         # Deep copy to avoid passing data around by reference.
208         return copy.deepcopy(self._config)
209
210
211
212     managed_list_config = ManagedListConfig()
213     static_config = StaticConfig()
214
215
216
217
218     @property
219     def placements(self):
220         placement_configs = copy.deepcopy(self._static_placements)
221
222         if decider.decide_experiment(self.MANAGED_LIST_DARK_READ_DECIDER):
223             for placement_id, placement_config in self._whitelisted_managed_list_placements.iteritems():
224                 if placement_id not in placement_configs:
225                     self._incr_failure_stat('missing_static_config.placements.{}'.format(placement_id))
226                 elif placement_configs[placement_id] != placement_config:
227                     self._incr_failure_stat('mismatching_config.placements.{}'.format(placement_id))
228
229
230
231
232         self._config = data
233
234     def get(self):
235         # Deep copy to avoid passing data around by reference.
236         return self._config
237
238
239
240     managed_list_config = ManagedListConfig()
241     static_config = StaticConfig()
242
243
244
245
246     @property
247     def placements(self):
248         placement_configs = self._static_placements.copy()
249
250         if decider.decide_experiment(self.MANAGED_LIST_DARK_READ_DECIDER):
251             for placement_id, placement_config in self._whitelisted_managed_list_placements.iteritems():
252                 if placement_id not in placement_configs:
253                     self._incr_failure_stat('missing_static_config.placements.{}'.format(placement_id))
254                 elif placement_configs[placement_id] != placement_config:
255                     self._incr_failure_stat('mismatching_config.placements.{}'.format(placement_id))
256
257
258
259
260         self._config = data
261
262     def get(self):
263         # Deep copy to avoid passing data around by reference.
264         return self._config
265
266
267
268     managed_list_config = ManagedListConfig()
269     static_config = StaticConfig()
270
271
272
273
274     @property
275     def placements(self):
276         placement_configs = self._static_placements.copy()
277
278         if decider.decide_experiment(self.MANAGED_LIST_DARK_READ_DECIDER):
279             for placement_id, placement_config in self._whitelisted_managed_list_placements.iteritems():
280                 if placement_id not in placement_configs:
281                     self._incr_failure_stat('missing_static_config.placements.{}'.format(placement_id))
282                 elif placement_configs[placement_id] != placement_config:
283                     self._incr_failure_stat('mismatching_config.placements.{}'.format(placement_id))
```

RPS / host



sampling profiler: flame graph

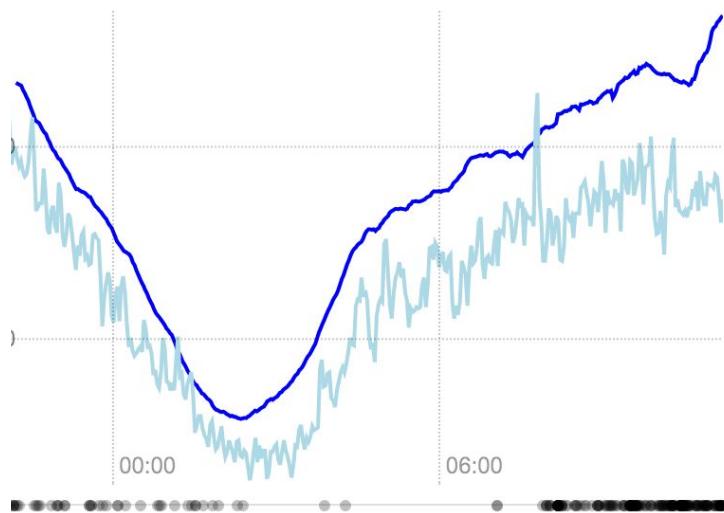
1.6% of ngapi spent in deepcopy



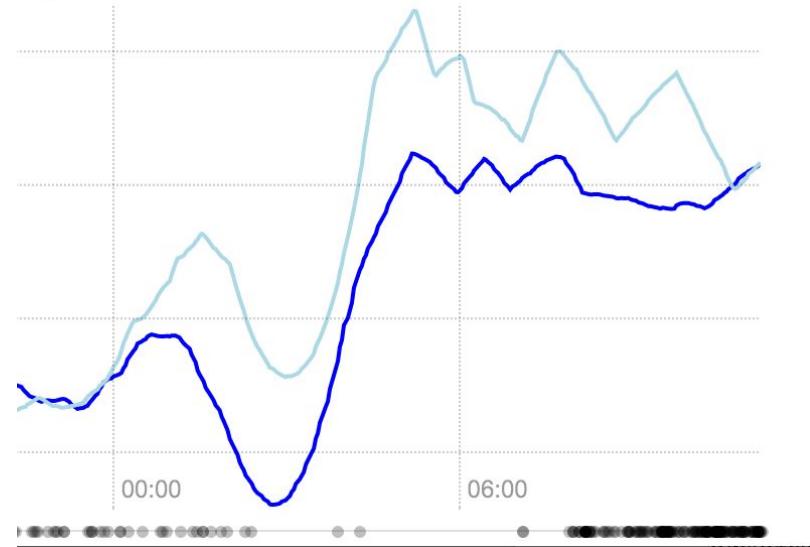
Example: Color Difference

API latency increase, RPS/host decrease

Latency



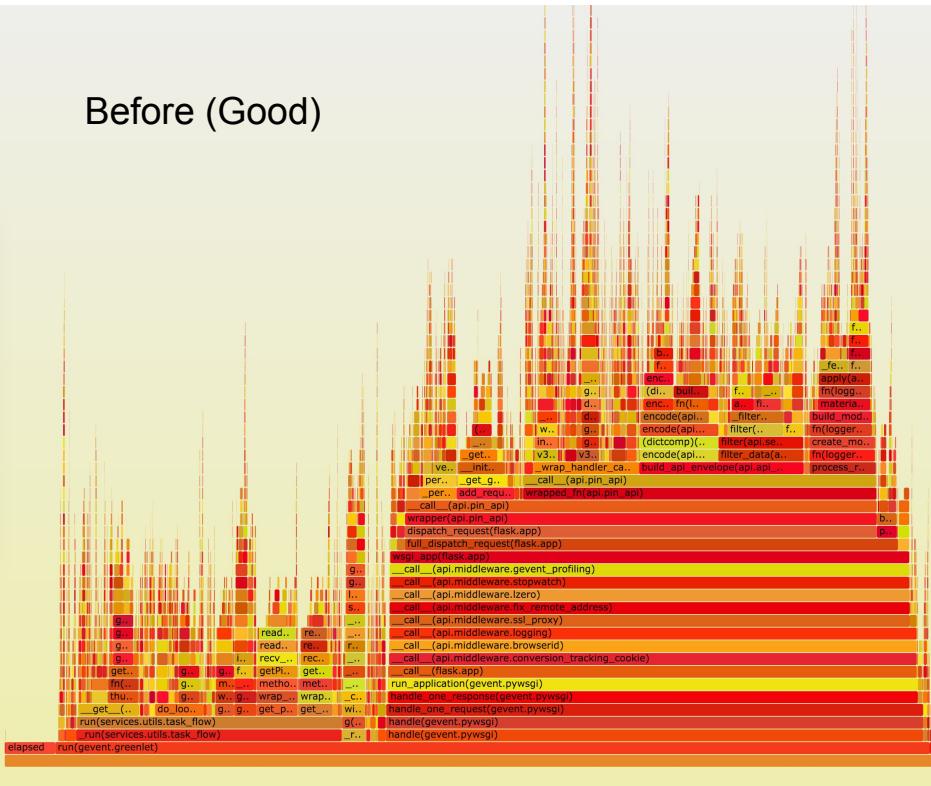
RPS/Host



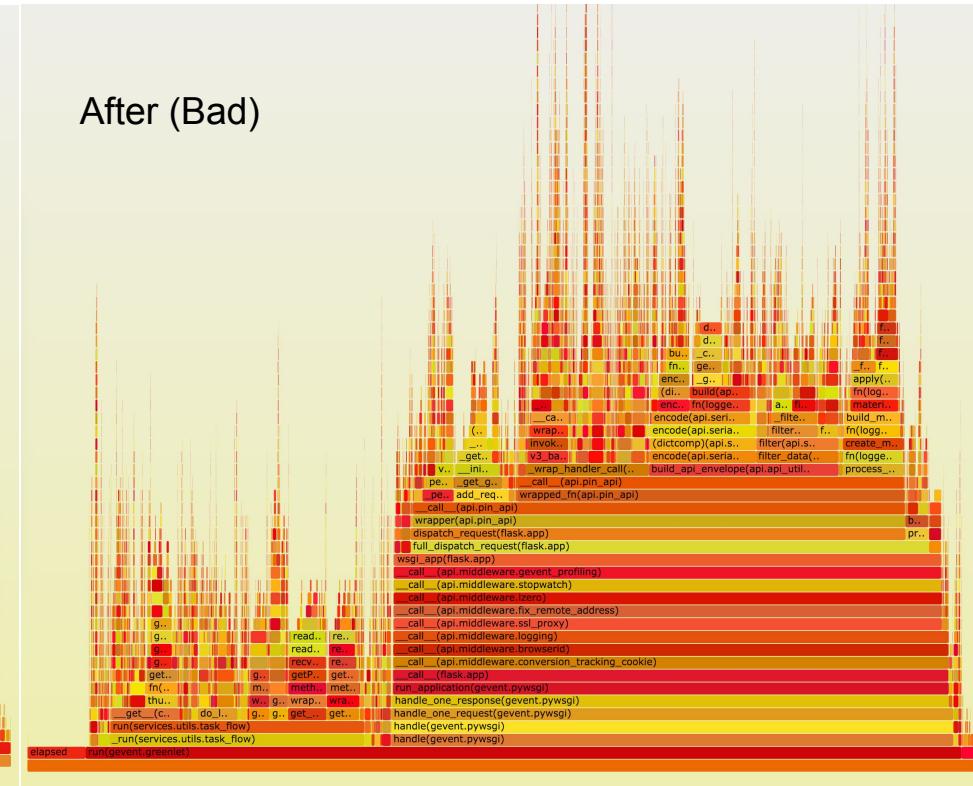
Example: Color Difference

Flamegraphs

Before (Good)

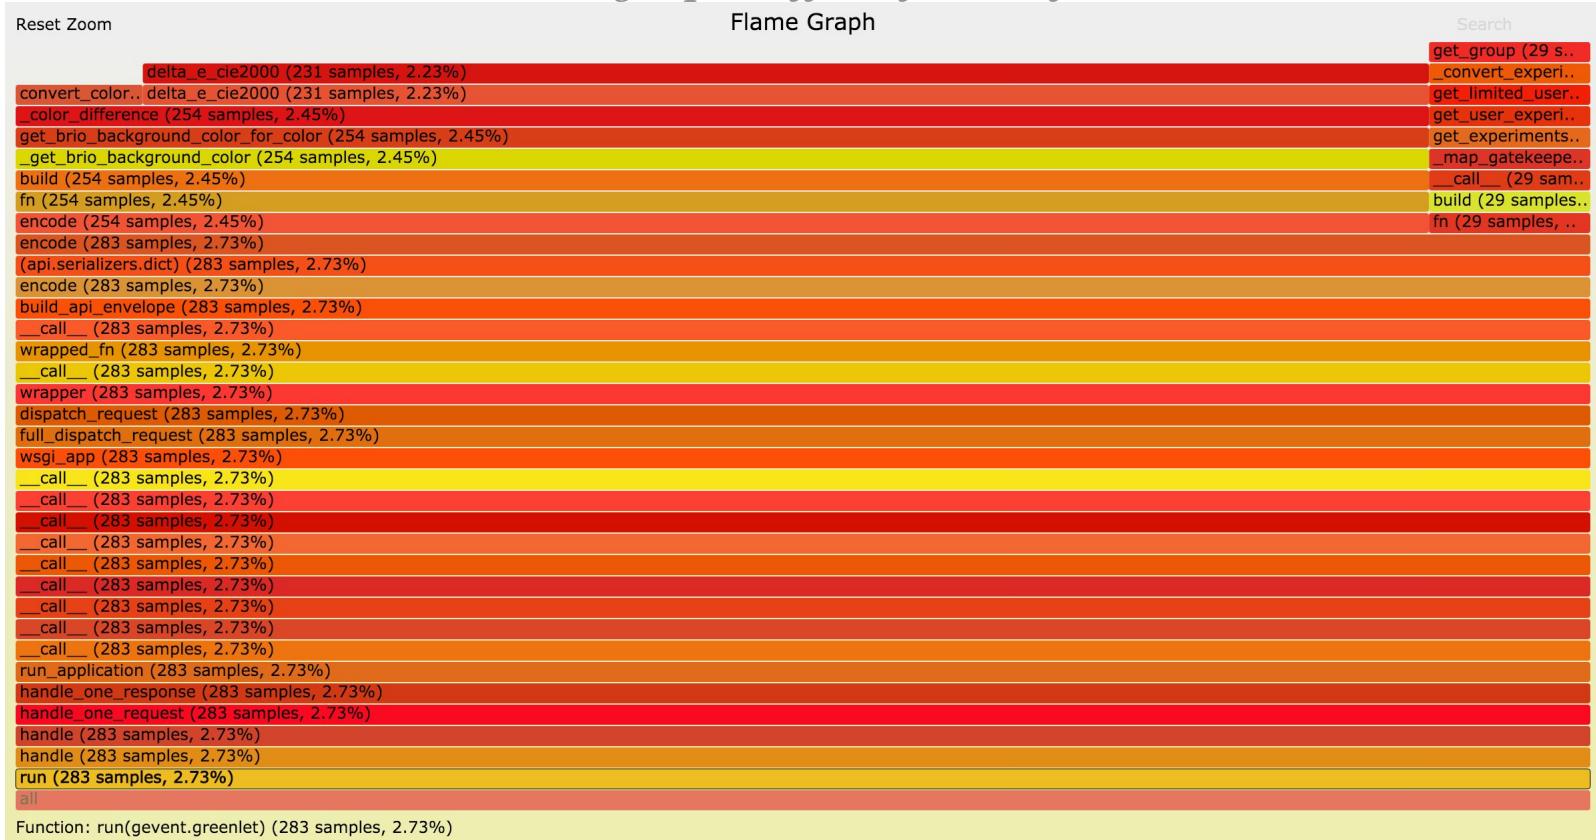


After (Bad)



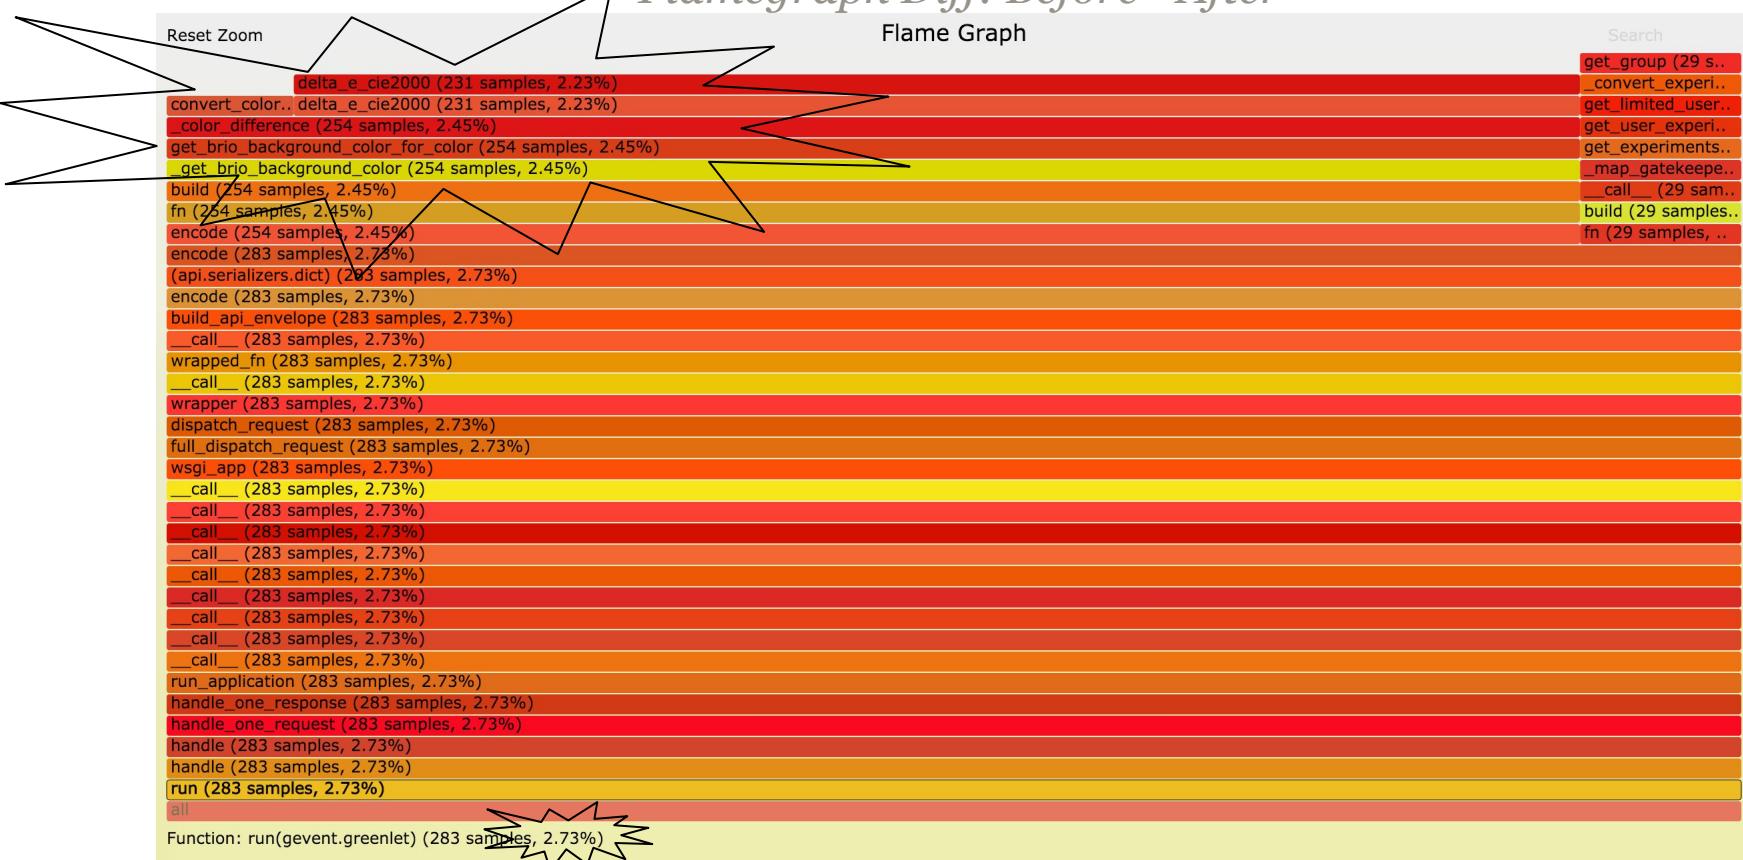
Example: Color Difference

Flamegraph Diff: Before - After



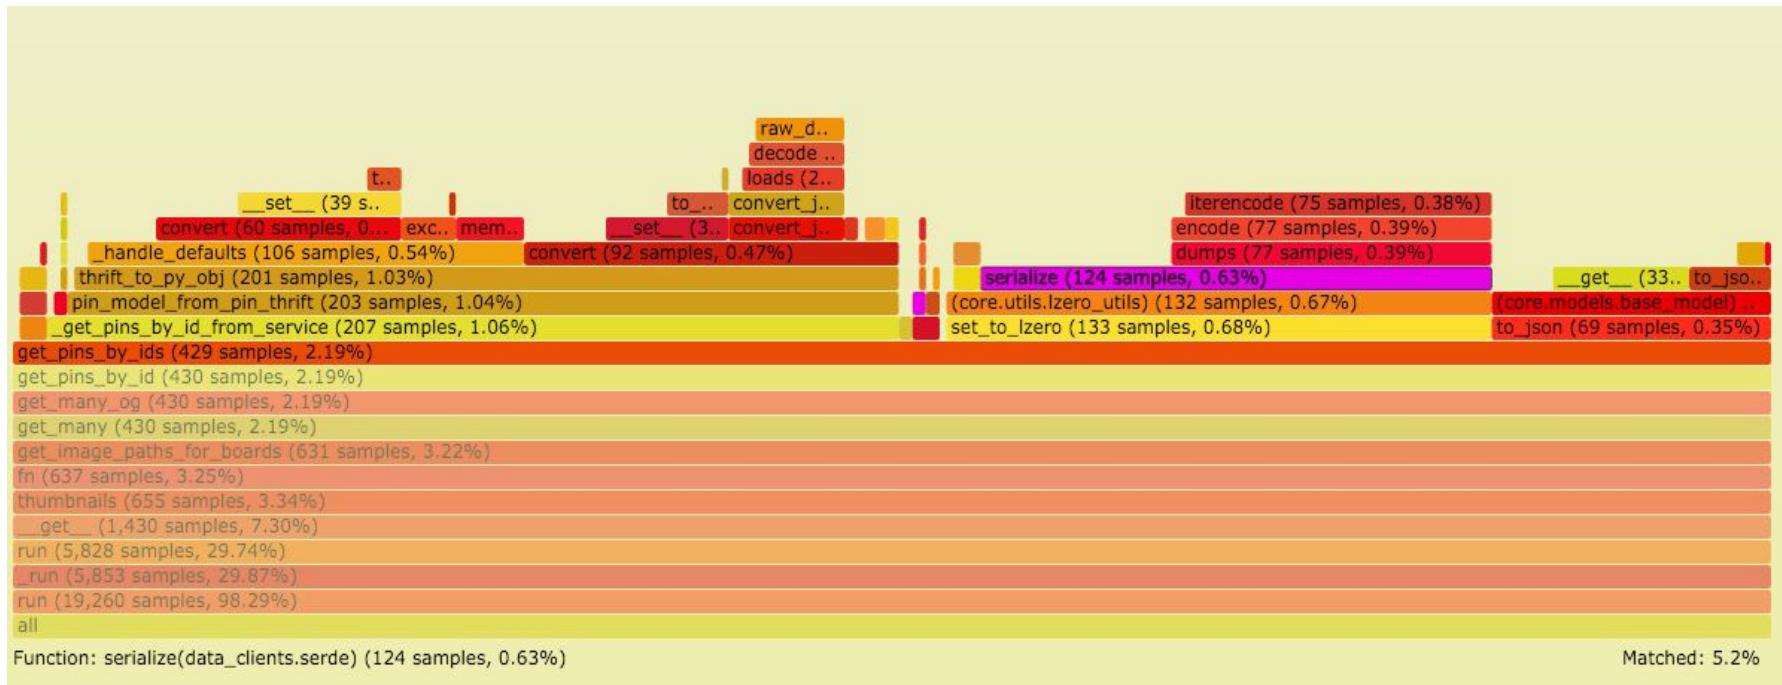
Example: Color Difference

Flamegraph Diff: Before - After



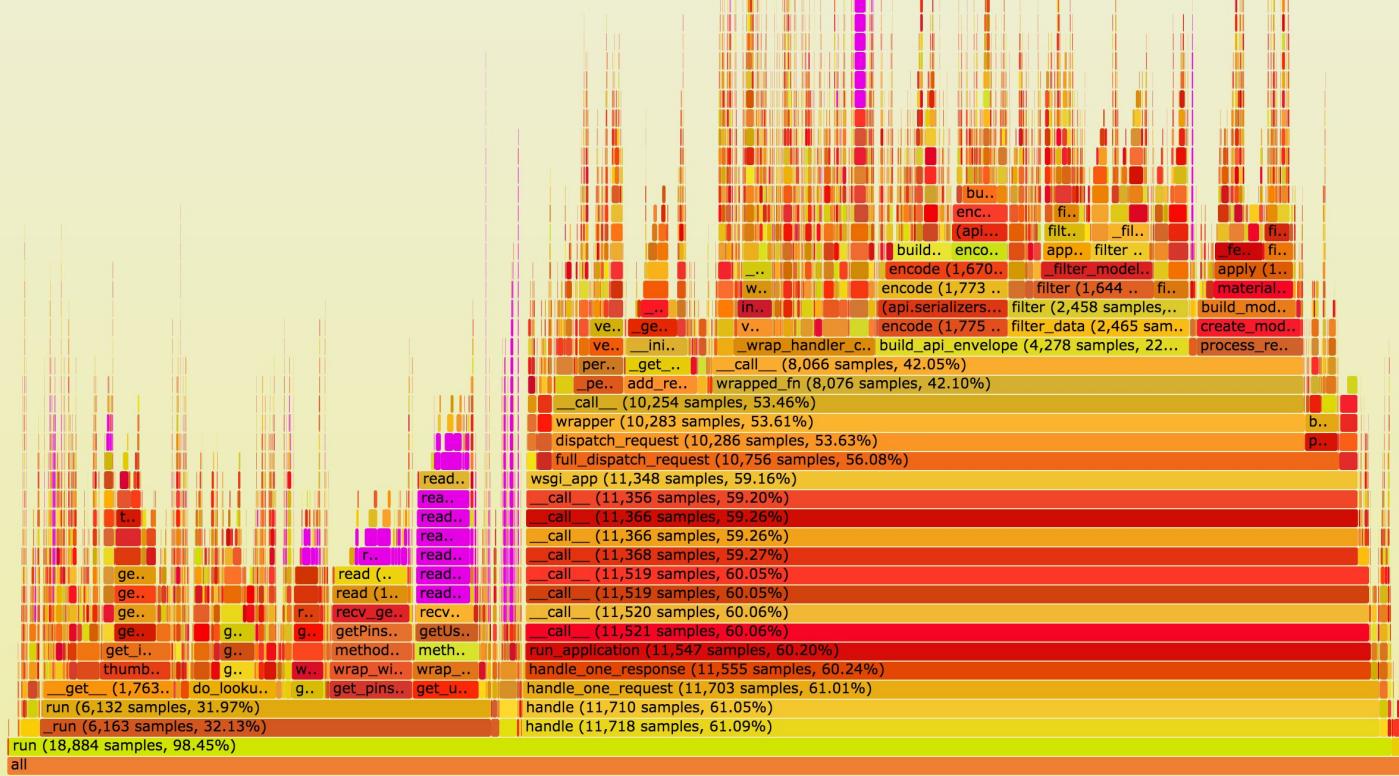
Example: Cache serialization

5%



Example: Thrift

RPC library



Matched: 13.6%

Example: Thrift

Before



Example: Thrift

fastbinary

```
class TBinaryProtocolAccelerated(TBinaryProtocol):  
    """C-Accelerated version of TBinaryProtocol.
```

This class does not override any of TBinaryProtocol's methods, but the generated code recognizes it directly and will call into our C module to do the encoding, bypassing this object entirely. We inherit from TBinaryProtocol so that the normal TBinaryProtocol encoding can happen if the fastbinary module doesn't work for some reason. ([TODO\(dreiss\)](#): Make this happen sanely in more cases.) To disable this behavior, pass fallback=False constructor argument.

In order to take advantage of the C module, just use TBinaryProtocolAccelerated instead of TBinaryProtocol.

Example: Thrift

<https://github.com/pinterest/thrift/commits/0.9.3-pinterest>

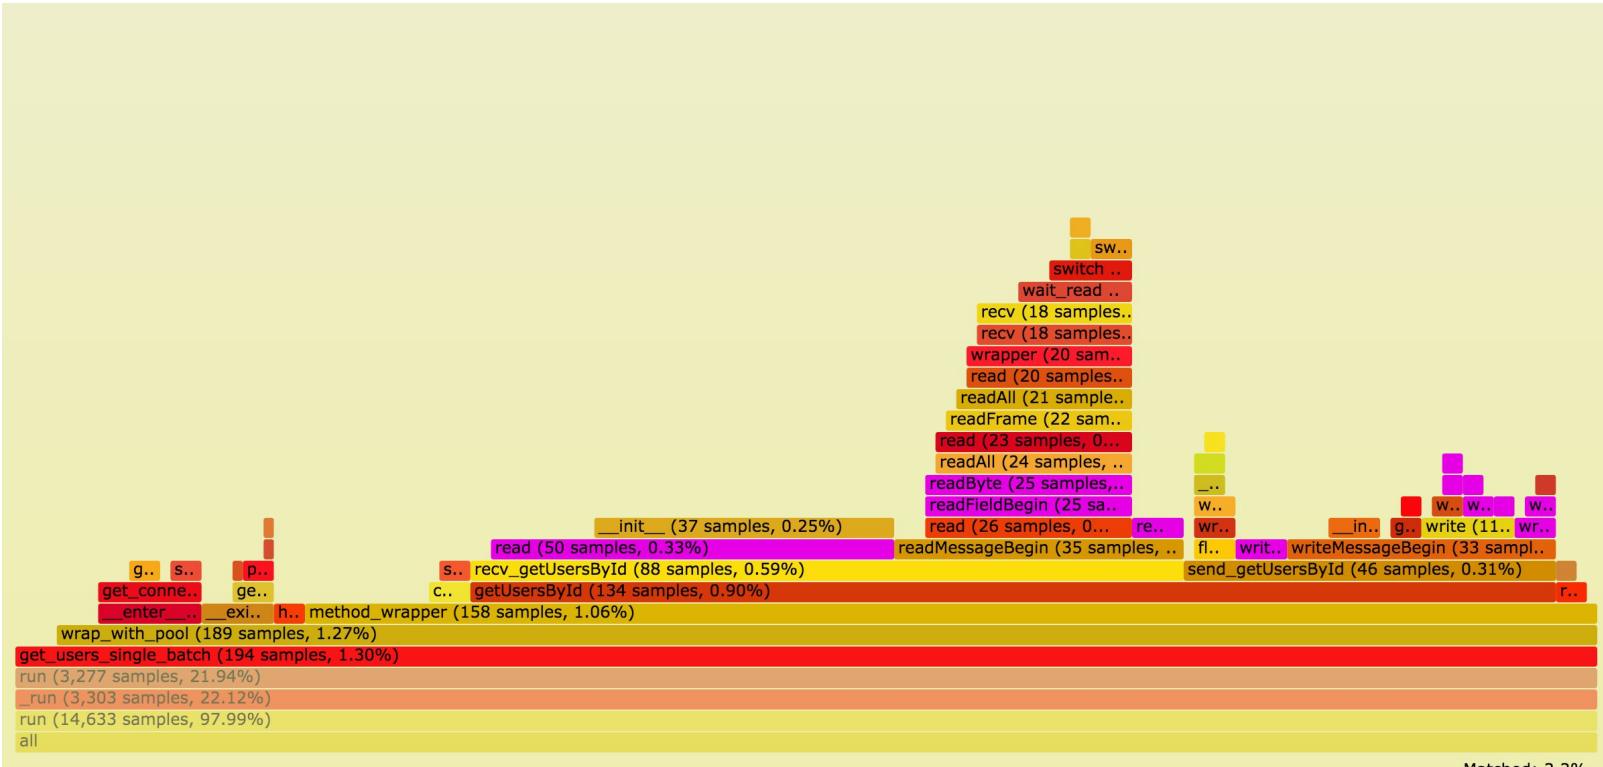
- Backport fixes for fastbinary from unreleased thrift 0.10.0
- Unicode support
- Support subclasses of TBinaryProtocolFactory
- Other bug fixes

```
indent_up();
```

```
- indent(out) << "if iprot.__class__ == TBinaryProtocol.TBinaryProtocolAccelerated "
+ indent(out) << "if isinstance(iprot, TBinaryProtocol.TBinaryProtocolAccelerated) "
    "and isinstance(iprot.trans, TTransport.CReadableTransport) "
    "and self.thrift_sproc is not None "
```

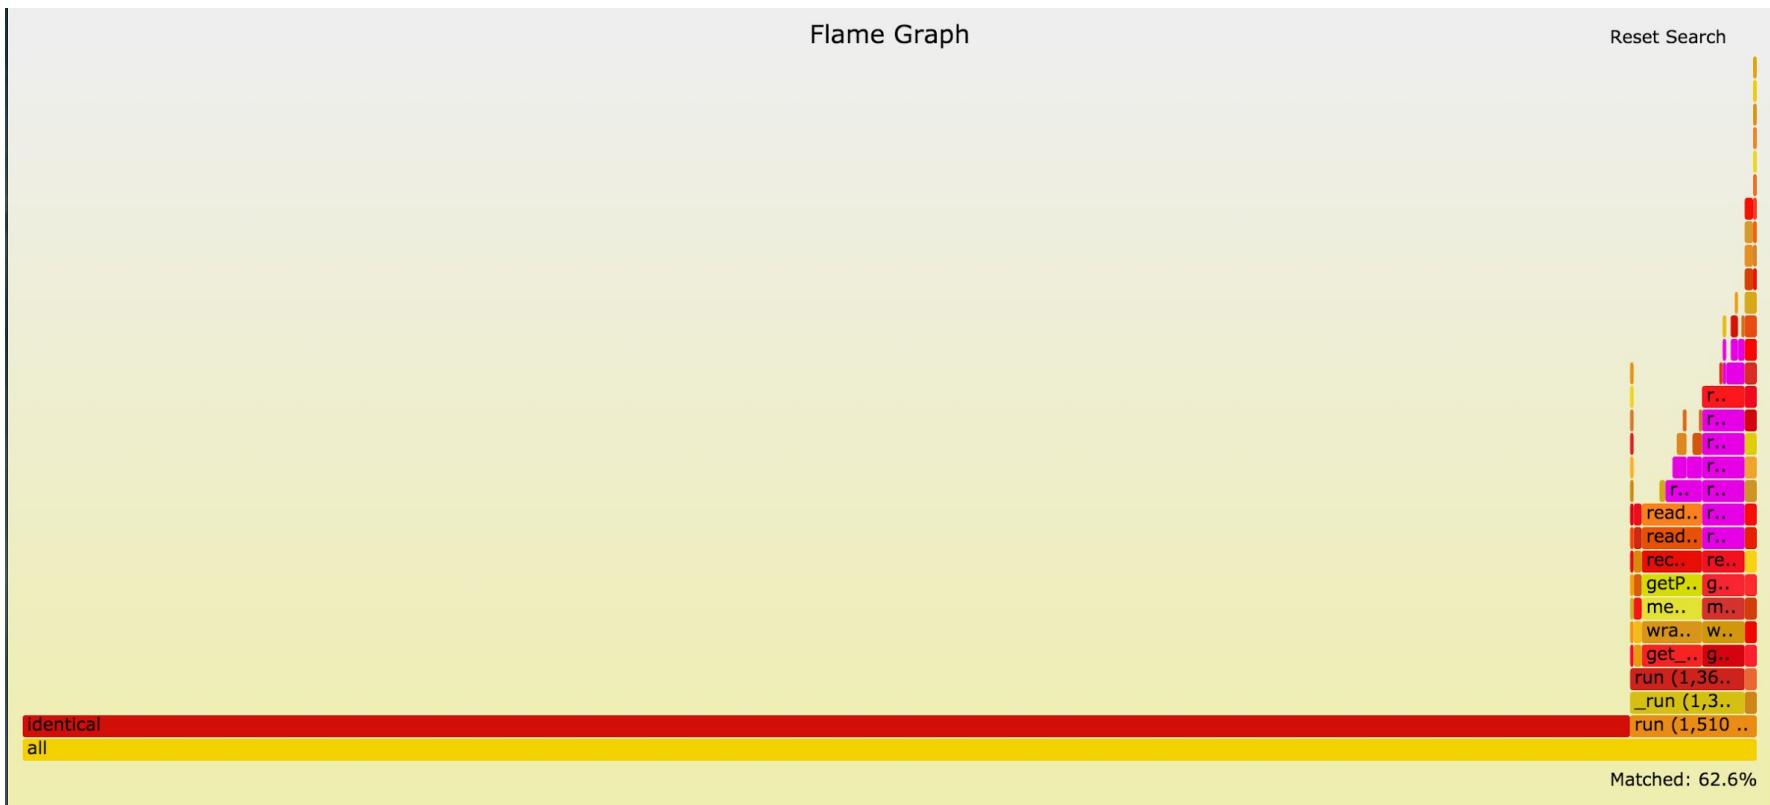
Example: Thrift

After



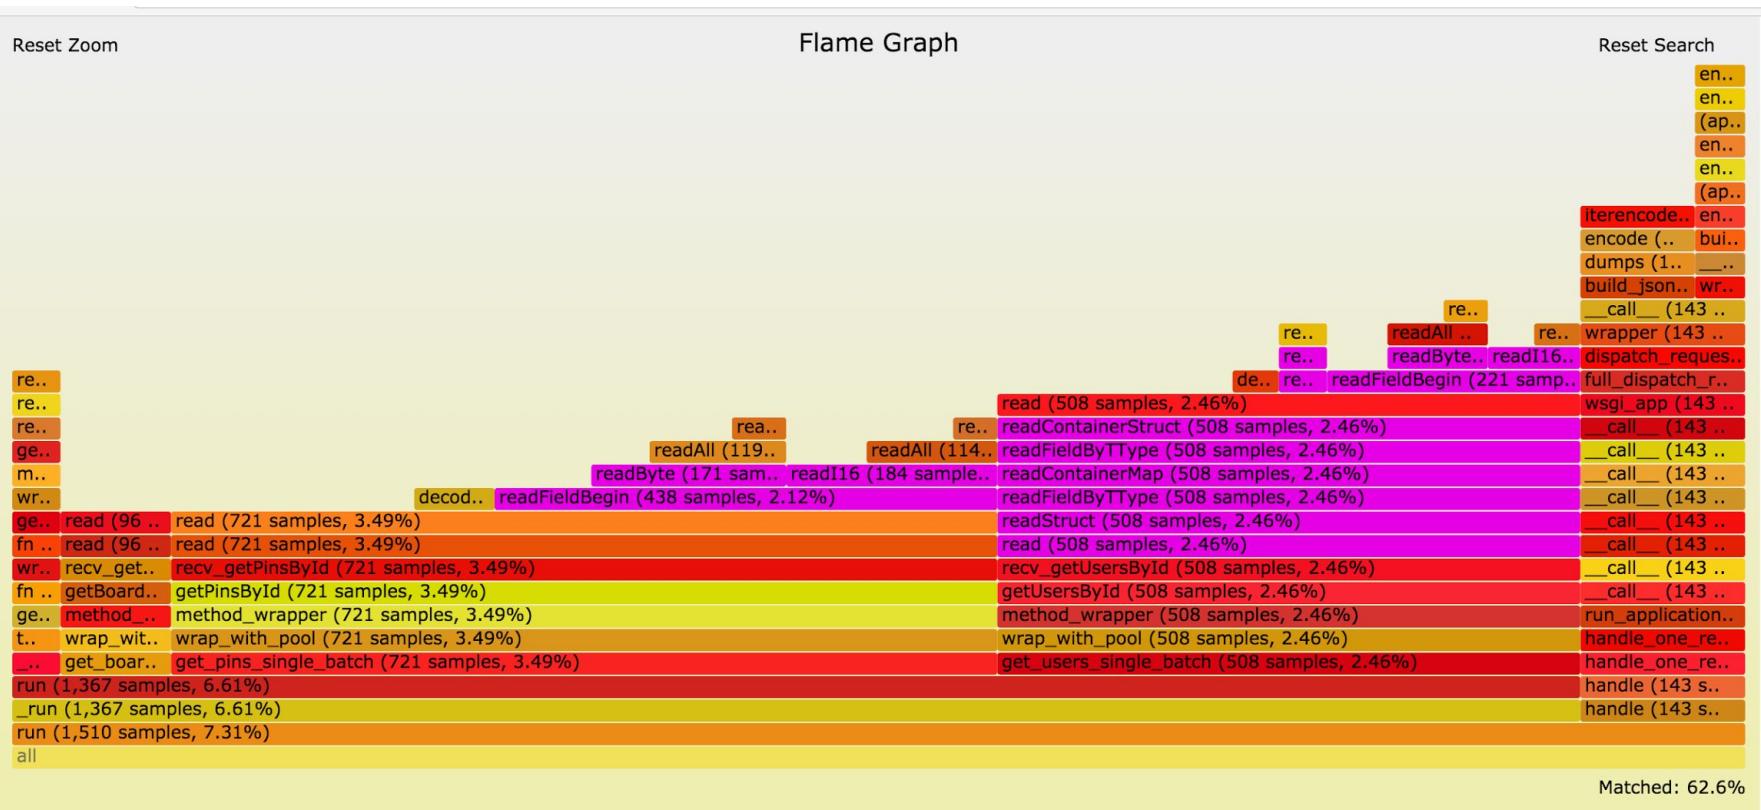
Example: Thrift

Difference: Before - After



Example: Thrift

Difference: Before - After





Problem

Gevent

Profiling

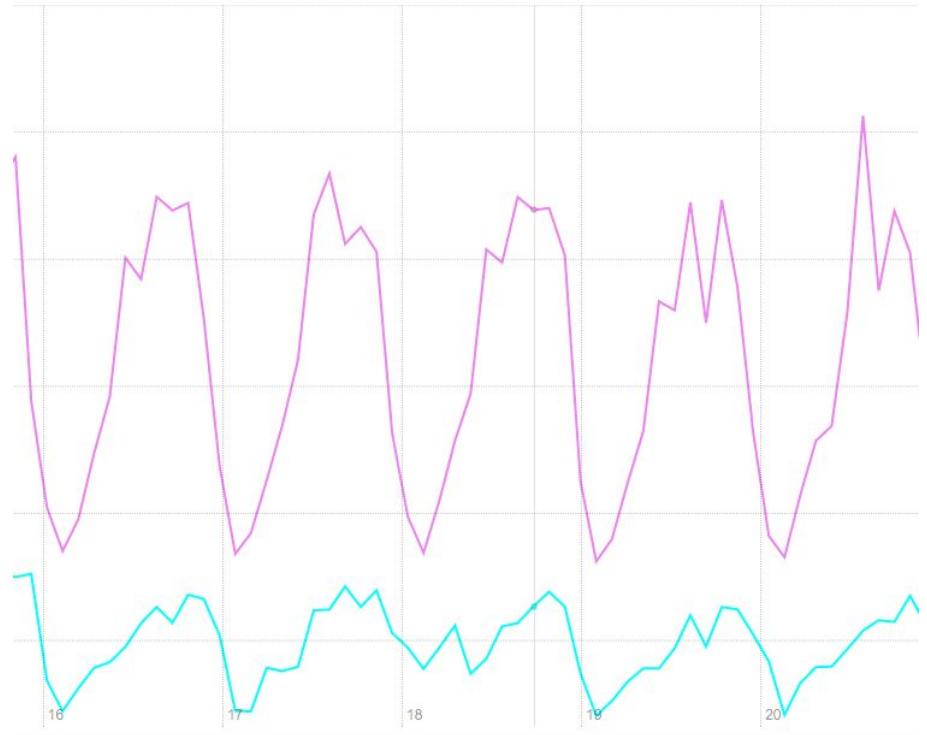


Impact

Pin Search P90 Latency

Q4 2015

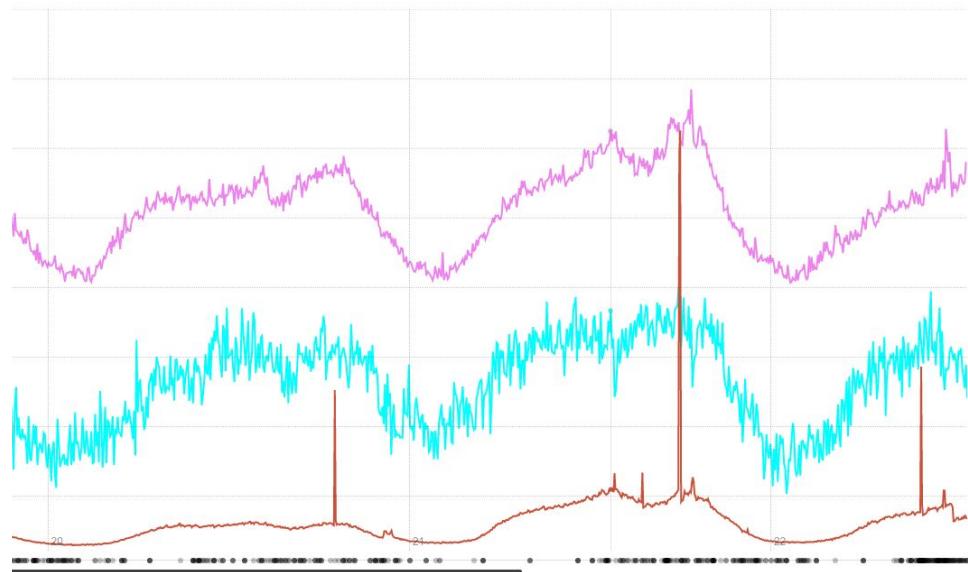
- **Latency**
 - Varnish: 2200ms
 - Search system: 640ms
- **70% of latency outside of Search system**
- **Height of Varnish latency >> height of asterix latency**



Pin Search P90 Latency

Today

- **Latency**
 - Varnish: 550ms
 - Search system: < 300ms
- **50% of search latency outside of Asterix**
- **Height of Varnish latency == height of asterix latency**

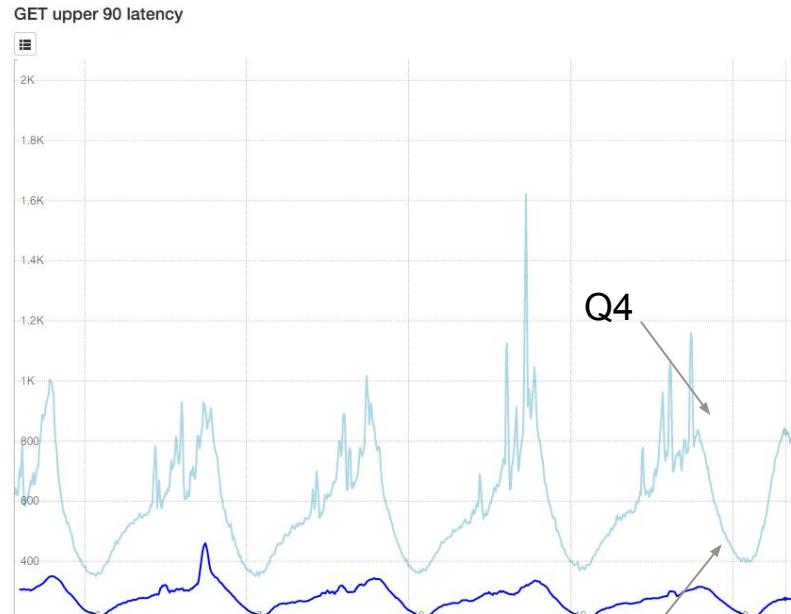


GET API Latency

Q4 vs today

- **Height (peak - valley)**

- Q4: 50%
- Now: 22%



Now



Conclusion

Conclusion

Lessons

- **Understand bottlenecks for latency and efficiency**
 - Network, CPU, Memory etc
 - Cooperative multitasking?
- **Profile in Production**
 - Synthetic benchmarks aren't enough
- **Visibility**
 - Profile code
 - Define efficiency metric
 - There are flame graph tools for most languages
- **Finding what to fix is harder than fixing it**
- **Start with low hanging fruit**
- **Detect and fix performance regressions**



What is Pinterest?



Pinterest is the world's catalog of ideas. Find and save recipes, parenting hacks, style inspiration and other ideas to try.