

Profiling the Unprofilable

PyCon APAC 2016, Seoul

Dmitry Trofimov

@DmitryTrofimov

14 August 2016



About me


- Dmitry Trofimov
- I work for JetBrains.
- I am the Team Lead and developer of PyCharm IDE
- I am interested in Python run-time: execution, debugging, profiling

*The best theory is inspired by practice. The
best practice is inspired by theory.*

Donald Knuth

<https://youtrack.jetbrains.com/issue/PY-14286>

<https://youtrack.jetbrains.com/issue/PY-14286>



Created by Yulia Zozulya 03 Nov 2014 17:39 Updated by Yulia Zozulya 03 Nov 2014 17:41

visible to: All Users ▾

★ PY-14286 Debugger: debugged code gets really slowed down

```
def main_work():  
    # do some busy work in parallel  
    print "Started main task"  
    x = 0  
    for i in xrange(100000000):  
        x += 1  
    print "Completed main task"
```

```
main_work()
```

With breakpoint in some part of the function execution get significantly slowed down (run ~3 secs, debug without breakpoints ~12 secs, debug with breakpoint on 4th line ~18 min).

PY-140.423

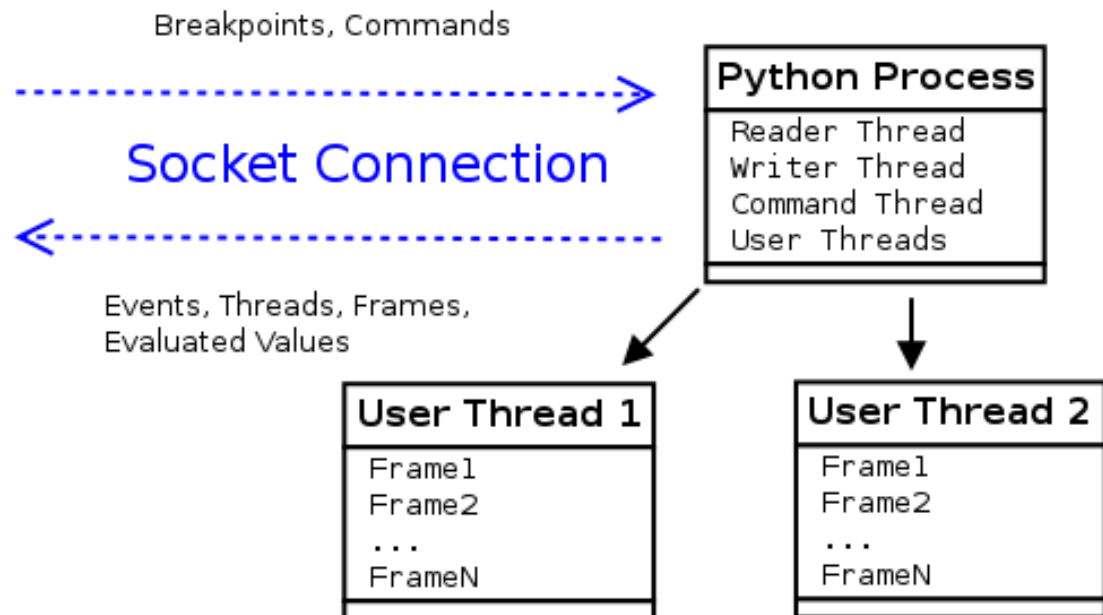
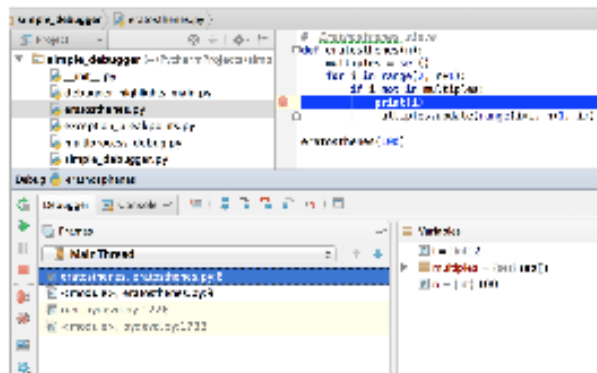
PyCharm debugger

<https://github.com/fabioz/PyDev.Debugger>

PyCon APAC 2015: Python Debugger Uncovered

<https://www.youtube.com/watch?v=DHf-6gW3-qs>

PyCharm Debugger



Trace Function

`sys.settrace(tracefunc)` ¶

Set the system's trace function, which allows you to implement a Python source code debugger in Python. `settrace()` for each thread being debugged.

Trace functions should have three arguments: *frame*, *event*, and *arg*. *frame* is the current stack frame. *event* is a string: 'call', 'line', 'return', 'exception', 'c_call', 'c_return', or 'c_exception'. *arg* depends on the event type.

The trace function is invoked (with *event* set to 'call') whenever a new local scope is entered; it should return a reference to a local trace function to be used that scope, or `None` if the scope shouldn't be traced.

```
def trace_function(self, frame, event, arg):  
    filename = frame.f_code.co_filename  
    line = frame.f_lineno  
  
    breakpoints_for_file = debugger.breakpoints.get(filename)  
    if breakpoints_for_file:  
        breakpoint = breakpoints_for_file[line]  
        ...
```

Issue PY-14286

Type: Performance

Subsystem: Debugger

Run	3 seconds
Debug	12 seconds
Debug + breakpoint	~18 minutes

Sample code

```
def main_work():  
    # do some busy work in parallel  
    print("Started main task")  
    x = 0  
    for i in xrange(100000000):  
        x += 1  
    print("Completed main task")  
  
main_work()
```

1. Reproducing

2. Analysis

- Normal run
 - Debug without breakpoints
 - Debug + BP in the function
 - Debug + BP in the same file
 - ~~Debug + BP in some other file~~
- = 4 different cases

Debugger

Execution Mode	Speed
Run	fast
Debug: no BP	fast
Debug: BP in main_work	slow
Debug: BP in file	slow

You can't improve what you can't measure
W. Edwards Deming

3. Measurement

```
import time

def main_work():
    print("Started main task")
    x = 0
    t = time.time()
    for i in range(1000000):
        x += 1
    print("Completed main task in %0.3fs" % (time.time() - t))

main_work()
```

PyCharm debugger

Execution Mode	Duration
Run	0.09s
Debug: no BP	0.11s
Debug: BP in main_work	9s
Debug: BP in file	9s

3. Something to compare with pdb

- Although less functional, but does almost the same
- Written in Python
- Is in standard library

4. Benchmarking

	PyCharm	pdb
Run	0.09s	0.09s
Debug: no breakpoints	0.11s	0.10s
Debug: BP in main_work	9s	5s
Debug: BP in the file	9s	5s

5. Find a bottleneck Profile

What is Profile?

Profile

A profile is a set of statistics that describes how often and for how long various parts of the program executed.

Let's use a Python profiler.

Python profilers

- cProfile
- yappi
- line_profiler

cProfile

- Part of the standard library
- Written in C
- The default choice for Python profiling

yappi

- Written in C
- Does the same as cProfile
- Can profile separate threads

line_profiler

- Written in Cython
- Provides statistics about lines

cProfile

Let's use it.

Unprofilable?

cProfile

cProfile provides deterministic profiling of Python programs.

What does 'deterministic profiling' means?

There are 2 major types of profilers

- Tracing(deterministic) profilers
- Sampling(Statistical) profilers

Tracing profilers

aka deterministic profilers aka Event-based profilers

Events: events like `c_{call,return,exception}`,
`python_{call,return,exception}`.

Documentation » The Python Standard Library

`sys.setprofile(profilefunc)`¶

Set the system's profile function, which allows you to implement a Python source code profiler in Python. See chapter [The Python Profilers](#) for more information on the Python profiler. The system's profile function is called similarly to the system's trace function (see [settrace\(\)](#)), but it isn't called for each executed line of code (only on call and return, but the return event is reported even when an exception has been set).

cProfile	setprofile
----------	------------

yappi	setprofile
-------	------------

line_profiler	settrace
---------------	----------

Unprofilable?

Sampling(Statistical) profilers

- Collect call stack samples regularly
- Less accurate and specific
- Add almost no overhead

Python statistical profilers

- statprof
- plop
- Intel Vtune Amplifier
- vmprof

statprof

- Written in Python
- Last update: 2012

plop

- Written in Python
- "a work in progress"

Intel Vtune Amplifier

- Has low overhead
- Profiles Python and native code
- Proprietary and not open-source
- Doesn't work on MacOSX

vmprof

- Written in C -- has low overhead
- Works on Linux, Mac, Windows
- Supports Python 2.7, 3.5, and PyPy
- Free and open-source

Use vmprof

Bottleneck found, what next?

6. Optimization

- Design
- Algorithms and data structures

Why debug without breakpoints
works so much faster?


```
def trace_function(self, frame, event, arg):
    filename = frame.f_code.co_filename
    line = frame.f_lineno

    breakpoints_for_file = debugger.breakpoints.get(filename)
    if breakpoints_for_file:
        breakpoint = breakpoints_for_file[line]
        if breakpoint is not None:
            ... # handle breakpoint
        return trace_function
    else:
        return None
```

Documentation » The Python Standard Library

`sys.settrace(tracefunc)`

The trace function is invoked (with *event* set to 'call') whenever a new local scope is entered; it should return a reference to a local trace function to be used that scope, or `None` if the scope shouldn't be traced.

The local trace function should return a reference to itself (or to another function for further tracing in that scope), or `None` to turn off tracing in that scope.

```
def trace_function(self, frame, event, arg):
    filename = frame.f_code.co_filename
    line = frame.f_lineno

    breakpoints_for_file = debugger.breakpoints.get(filename)

    for breakpoint in breakpoints_for_file.itervalues():
        if breakpoint.func_name == frame.f_code.co_name:
            break
    else:
        return None

    ...
```

A. Algorithmic optimization

	Baseline	Optimization A
Run	0.09s	0.09s
Debug: no breakpoints	0.11s	0.11s
Debug: BP in main_work	9s	9s
Debug: BP in the file	9s	0.11s

6. Optimization

- Design
- Algorithms and data structures
- Source code

Line profiling

B. Source optimization

	Opt A	Opt A+B
Run	0.09s	0.09s
Debug: no breakpoints	0.11s	0.11s
Debug: BP in main_work	9s	8s
Debug: BP in the file	0.11s	0.11s

6. Optimization

- Design
- Algorithms and data structures
- Source code
- Build
- Compile
- Assembly
- Run time

Is optimization reached it's limit?

Go beyond Python!

Rewrite everything in C?

- Compatibility with Jython, IronPython, PyPy etc
- Avoid code duplication
- Python is much better language then C

Cython

- Static compiler for Python
- Gives the combined power of Python and C

```
def primes(int kmax):  
    cdef int n, k, i  
    cdef int p[1000]  
    result = []  
    if kmax > 1000:  
        kmax = 1000  
    k = 0  
    n = 2  
    while k < kmax:  
        i = 0  
        while i < k and n % p[i] != 0:  
            i = i + 1  
        if i == k:  
            p[k] = n  
            result.append(n)  
        n = n + 1  
    return result
```

```
def trace_dispatch(self, frame, str event, arg):  
    cdef str filename;  
    cdef bint is_exception_event;  
    cdef bint has_exception_breakpoints;  
    cdef bint can_skip;  
    cdef PyDBAdditionalThreadInfo info;  
    cdef int step_cmd;  
    cdef int line;  
    cdef str curr_func_name;  
    cdef bint exist_result;
```

C. Cython compile optimization


	Opt A+B	Opt A+B+C
Run	0.09s	0.09s
Debug: no breakpoints	0.11s	0.11s
Debug: BP in main_work	8s	4s
Debug: BP in the file	0.11s	0.11s

7. Compare results

	Baseline	pdb	A+B+C
Run	0.09s	0.09s	0.09s
Debug: no breakpoints	0.11s	0.10s	0.11s
Debug: BP in main_work	9s	5s	4s
Debug: BP in the file	9s	5s	0.11s

Unprofilable again?

<https://youtrack.jetbrains.com/issue/PY-14286>



Created by Yulia Zozulya 03 Nov 2014 17:39 Updated by Yulia Zozulya 03 Nov 2014 17:41

visible to: All Users ▾

★ PY-14286 Debugger: debugged code gets really slowed down

```
def main_work():  
    # do some busy work in parallel  
    print "Started main task"  
    x = 0  
    for i in xrange(1000000000):  
        x += 1  
    print "Completed main task"
```

```
main_work()
```

With breakpoint in some part of the function execution get significantly slowed down (run ~3 secs, debug without breakpoints ~12 secs, debug with breakpoint on 4th line ~18 min).

PY-140.423

<https://www.python.org/dev/peps/pep-0523>

PEP 523 -- Adding a frame evaluation API to CPython

PEP:	523
Title:	Adding a frame evaluation API to CPython
Author:	Brett Cannon <brett at python.org>, Dino Viehland <dinov at microsoft.com>
Status:	Draft
Type:	Standards Track
Created:	16-May-2016
Post-History:	16-May-2016

Conclusion

- Use profilers to find bottlenecks in your code
- There are different profilers: each has own upsides
- Start to optimize things from higher level to lower
- To optimize Python on a lower level use Cython

Links

- vmprof.readthedocs.io
- cython.org
- github.com/fabioz/PyDev.Debugger

Q&A